

Advanced Network Technologies

Applications

Assignment Project Exam Help

<https://eduassistpro.github.io/>

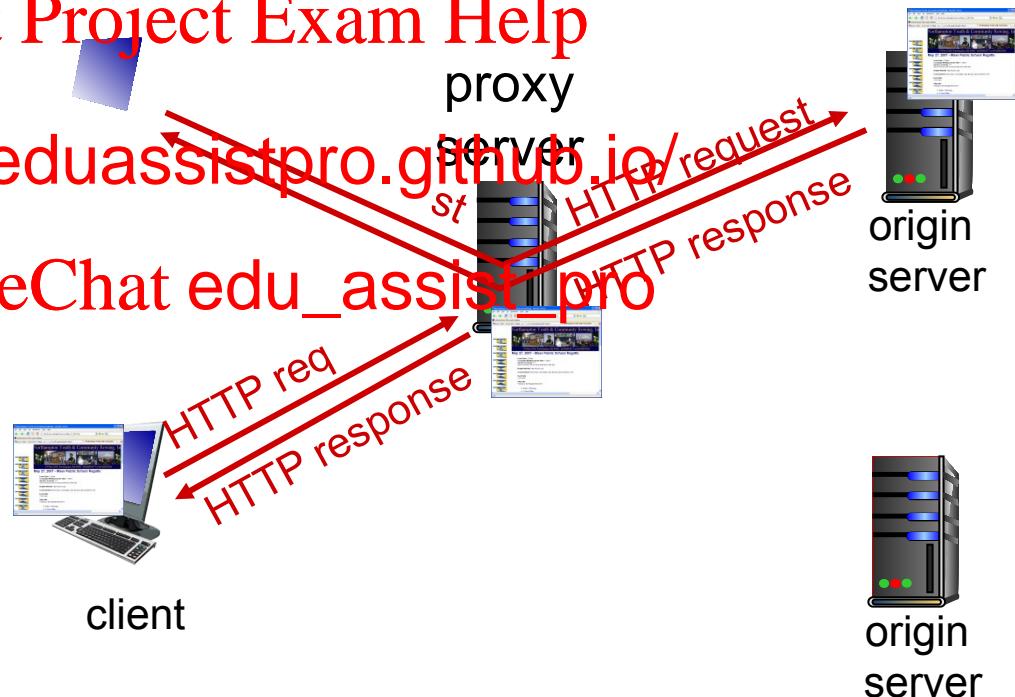
Add WeChat edu_assist_pro

Dr. Wei Bao | Lecturer
School of Computer Science



goal: satisfy client request without involving origin server

- › user sets browser: Web
- accesses via cache
- › browser sends all H requests to cache
- › if object in cache:
 - then cache returns object
 - else cache requests object from origin server, then returns object to client



› Q: Does the cache act as a client or a server?

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

- › R: cache acts as both client and server

- server for original requesting client
- client to origin serv

why Web caching?

- › reduce response time for client request

<https://eduassistpro.github.io/>

- › typically cache is installed by ISP (university, company, residential ISP)

Add WeChat [edu_assist_pro](#)

assumptions:

- ❖ avg object size: 100K bits
- ❖ avg request rate from browsers to origin servers: 15/sec (1.5 Mbps service)
- ❖ RTT from institutional router to any origin server: 2 sec
- ❖ access link rate: 1.54

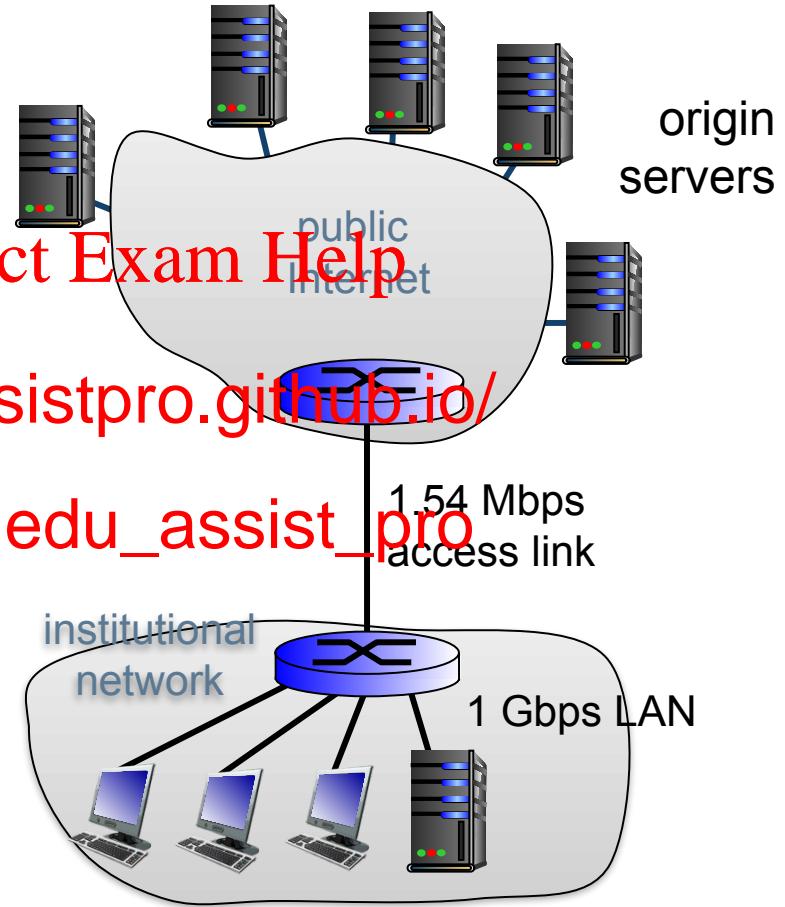
Assignment Project Exam Help

<https://eduassistpro.github.io/>

consequences:

- ❖ LAN utilization: 0.15%
- ❖ LANU = avg req rate * size / link bandwidth
- ❖ access link utilization = **99% problem!**
- ❖ ALU = avg req rate * size / link bandwidth
- ❖ total delay = 2 sec + seconds + usecs

Add WeChat edu_assist_pro



Q: what happens with fatter access link?

Caching example: fatter access link

assumptions:

- ❖ avg object size: 100K bits
- ❖ avg request rate from browsers to origin servers: 15/sec
- ❖ RTT from institutional router to any origin server: 2 sec
- ❖ access link rate: 1.54

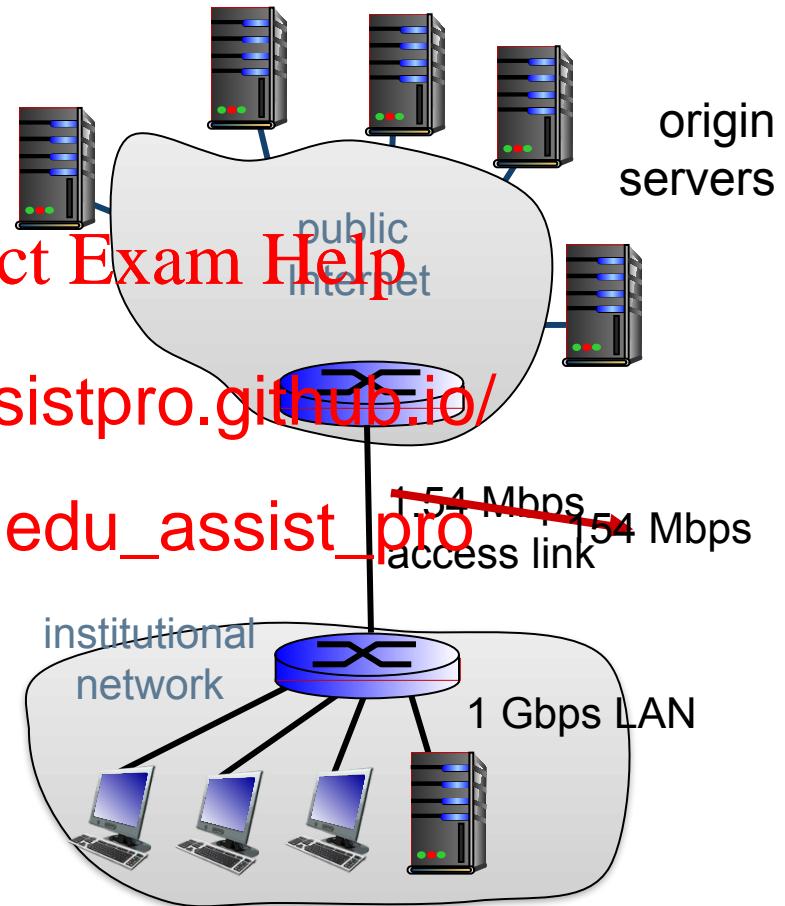
Assignment Project Exam Help

<https://eduassistpro.github.io/>

consequences:

- ❖ LAN utilization: 0.15%
- ❖ access link utilization = ~~99%~~ ^{0.99%}
- ❖ total delay = 2 sec + seconds + usecs

msecs



Cost: increased access link speed (not cheap!)

Caching example: install local cache

assumptions:

- ❖ avg object size: 100K bits
- ❖ avg request rate from browsers to origin servers: 15/sec
- ❖ RTT from institutional router to any origin server: 2 sec
- ❖ access link rate: 1.54

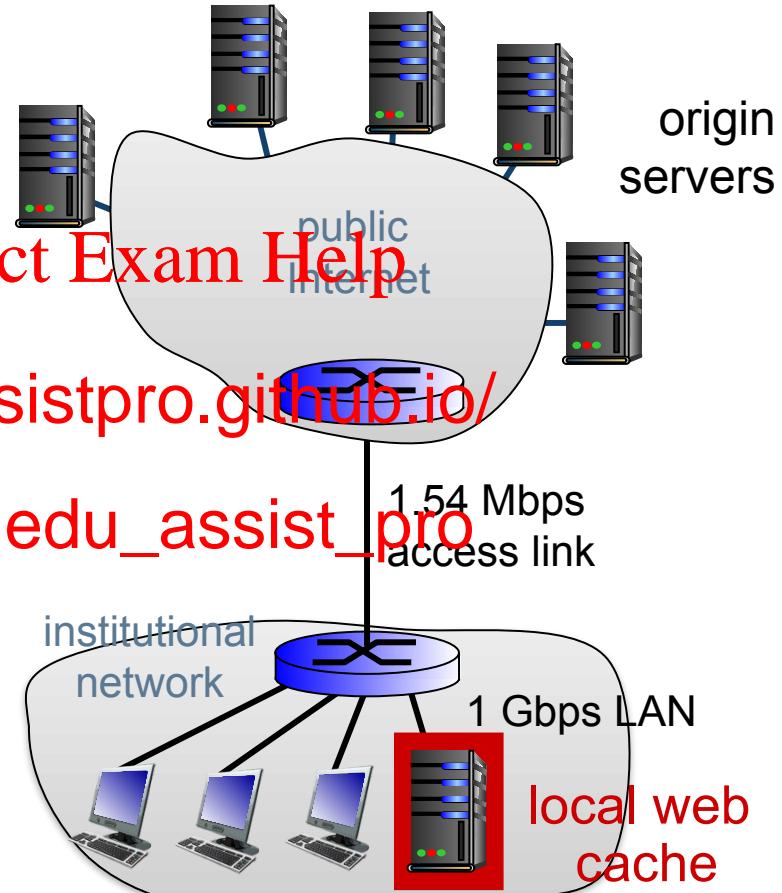
Assignment Project Exam Help

<https://eduassistpro.github.io/>

consequences:

- ❖ LAN utilization: 0.15%
- ❖ access link utilization = 0%
- ❖ total delay = usecs

Add WeChat edu_assist_pro



Cost: web cache (cheap!)

Caching example: install local cache

Calculating access link utilization, delay with cache:

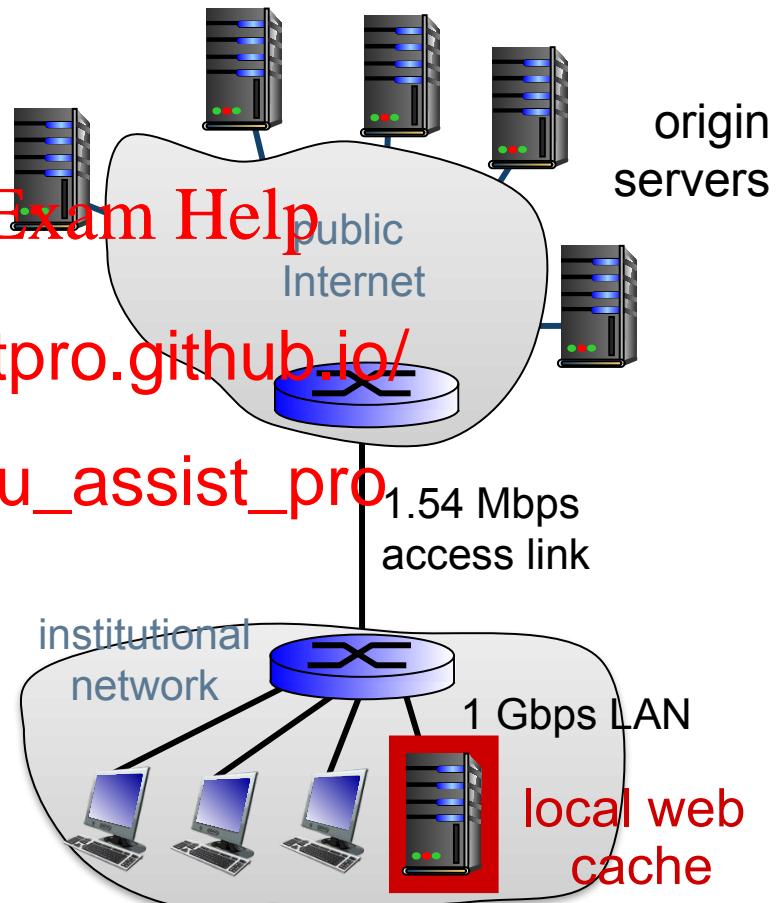
- › suppose cache hit rate is 0.4
 - 40% requests satisfied at cache,
 - 60% requests satisfied at origin
- › access link utilization <https://eduassistpro.github.io/>
 - 60% of requests use access link
- › average total delay

$$= 0.6 * (\text{delay from origin servers}) + 0.4 * (\text{delay when satisfied at cache})$$

Link utilization is around 60%, queueing delay is small enough

$$= 0.6 (\sim 2.x \text{ second}) + 0.4 (\sim \text{usecs})$$

less than with 154 Mbps link (and cheaper too!)



- › **Goal:** don't send object if client **client** has up-to-date cached version

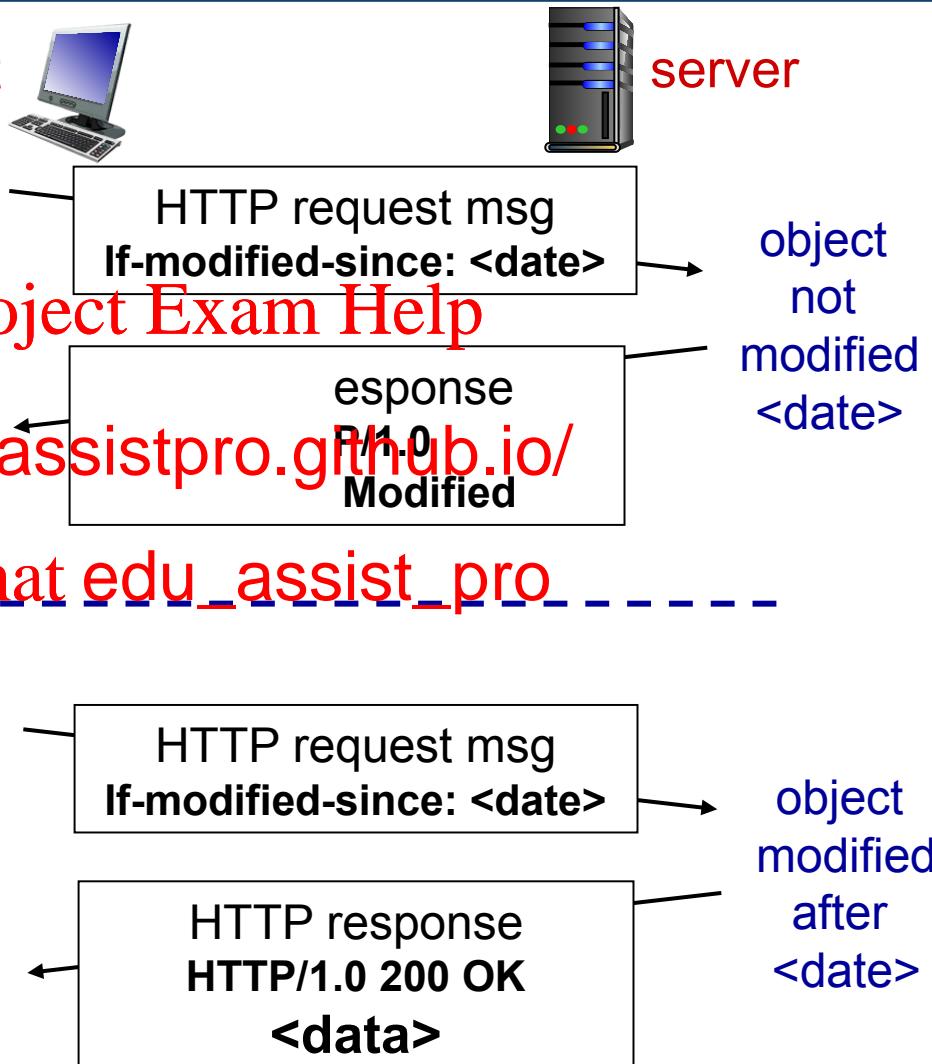
- no object transmission delay
- lower link utilization

- › **client:** specify date of copy in HTTP request

If-modified-since: <https://eduassistpro.github.io/>

- › **server:** response contains no object if cached copy is up-to-date:

HTTP/1.0 304 Not Modified



Key goal: decreased delay in multi-object HTTP requests

HTTP1.1: introduced multiple, pipelined GETs over single TCP connection

Assignment Project Exam Help

- server responds (e-first-served scheduling) to <https://eduassistpro.github.io/>
- with FCFS, small object may have transmission blocking (head-of-line (HOL) blocking) by large object(s)
- loss recovery (retransmitting lost TCP segments) stalls object transmission

Key goal: decreased delay in multi-object HTTP requests

HTTP/2: [RFC 7540, 2015] increased flexibility at server in sending object

- methods, status <https://eduassistpro.github.io/> unchanged from HTTP 1.1
- transmission order of requested objects based on client-specified object priority (not necessarily FCFS)
- *push* unrequested objects to client
- divide objects into frames, schedule frames to mitigate Head-of-line (HOL) blocking

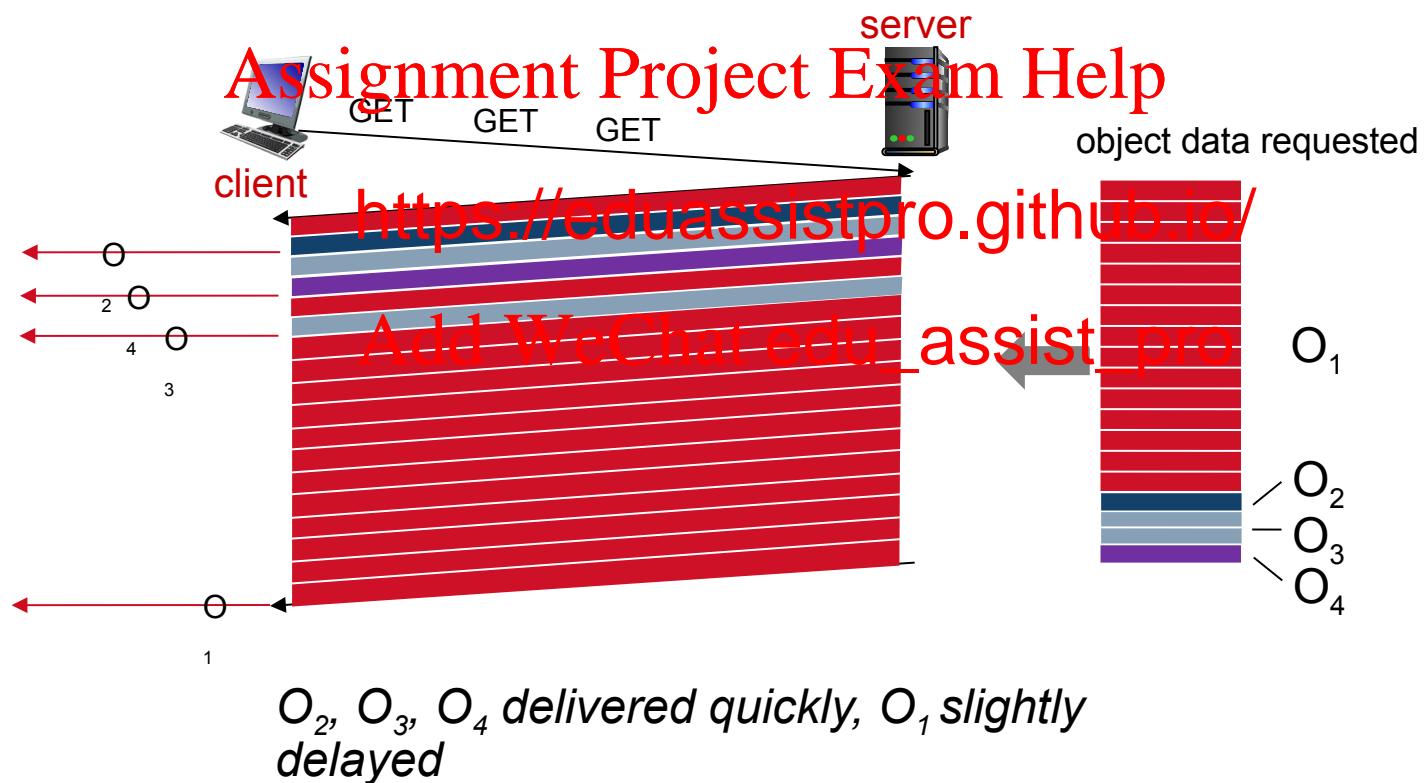
HTTP/2: mitigating HOL blocking

HTTP 1.1: client requests 1 large object (e.g., video file, and 3 smaller objects)



HTTP/2: mitigating HOL blocking

HTTP/2: objects divided into frames, frame transmission interleaved





Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Client

Server

Frames:

- Basic HTTP/2 data unit, replacing HTTP/1.1 header and body format.
 - HTTP/2 frame
 - Header frames, Data frames
 - Bidirectional channel where frames are transmitted
 - Replacing HTTP/1.1 Request-Response mode
- A single TCP connection to carry multiple streams
- Assignment Project Exam Help
https://eduassistpro.github.io/
Add WeChat edu_assist_pro

The HTTP/2 Server Push mechanism allows the server to send resources proactively without waiting for a request, when it believes the client will need them.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat `edu_assist_pro`

<https://blog.golang.org/h2push>

- › Web and HTTP (Done)

Assignment Project Exam Help

- › FTP

<https://eduassistpro.github.io/>

- › Email

Add WeChat edu_assist_pro

- › DNS

- › P2P



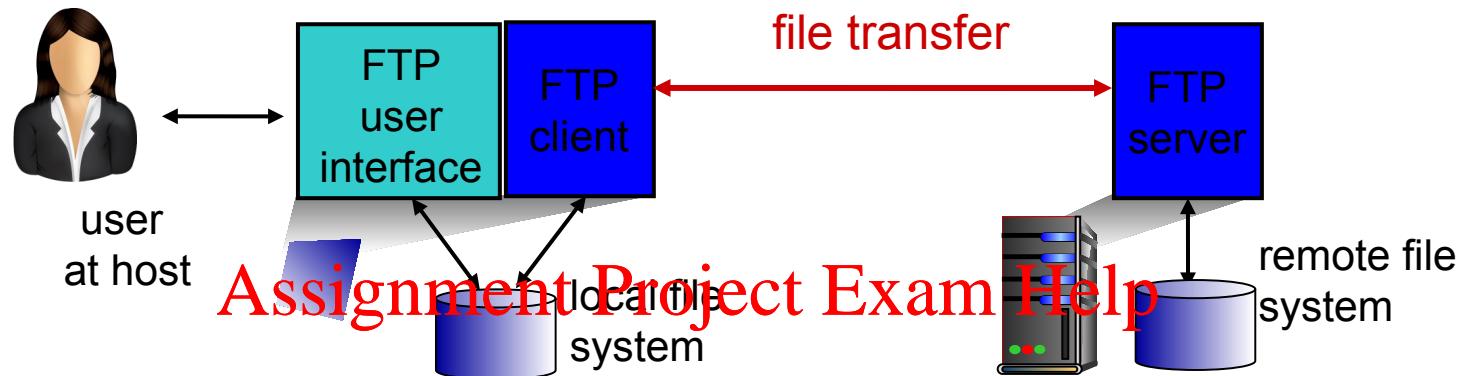
THE UNIVERSITY OF
SYDNEY

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

FTP: the file transfer protocol



<https://eduassistpro.github.io/>

- ❖ transfer file ~~Add from rem~~ **Add WeChat** **edu_assist_pro**
- ❖ client/server model
 - **client:** side that initiates transfer (either to/from remote)
 - **server:** remote host
- ❖ ftp: RFC 959
- ❖ ftp server: port 21, 20

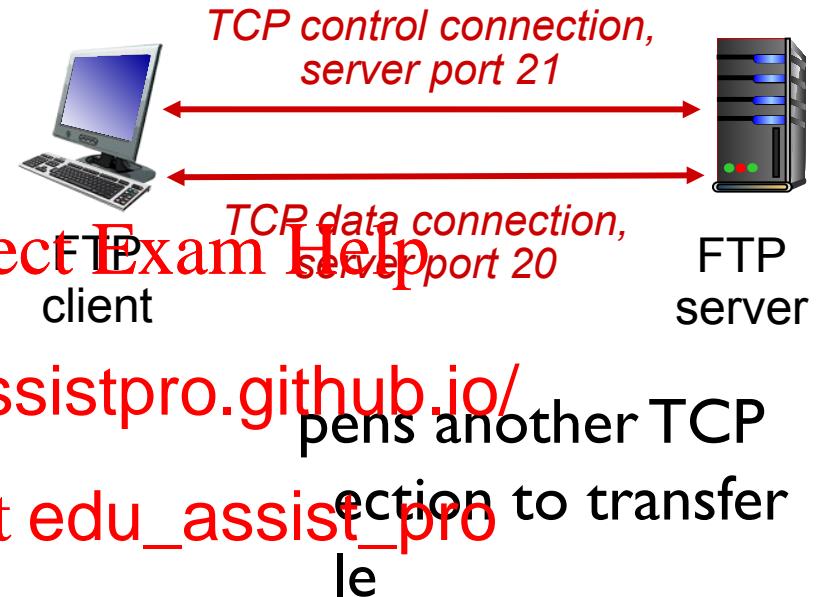
FTP: separate control, data connections

- › FTP client contacts FTP server at port 21, using TCP

- › client authorized over control connection

- › client browses remote directory, sends commands over control connection

- › when server receives file transfer command, **server** opens 2nd TCP data connection (for file) to client
- › after transferring one file, server closes data connection



- › control connection: “*out of band*”
- › FTP server maintains “state”: current directory, earlier authentication

sample commands:

- › sent as ASCII text over control channel
- › **USER** *username*
- › **PASS** *password*
- › **LIST** return list of file in current directory
- › **RETR** *filename* retrieves (gets) file
- › **STOR** *filename* stores (puts) file onto remote host

sample return codes

- › status code and phrase (as in HTTP)
331 Username OK,
d required
a conn
- › **open;**
starting
- › **425** Can't open data connection
- › **452** Error writing file



THE UNIVERSITY OF
SYDNEY

Assignment Project Exam Help
Email

<https://eduassistpro.github.io/>

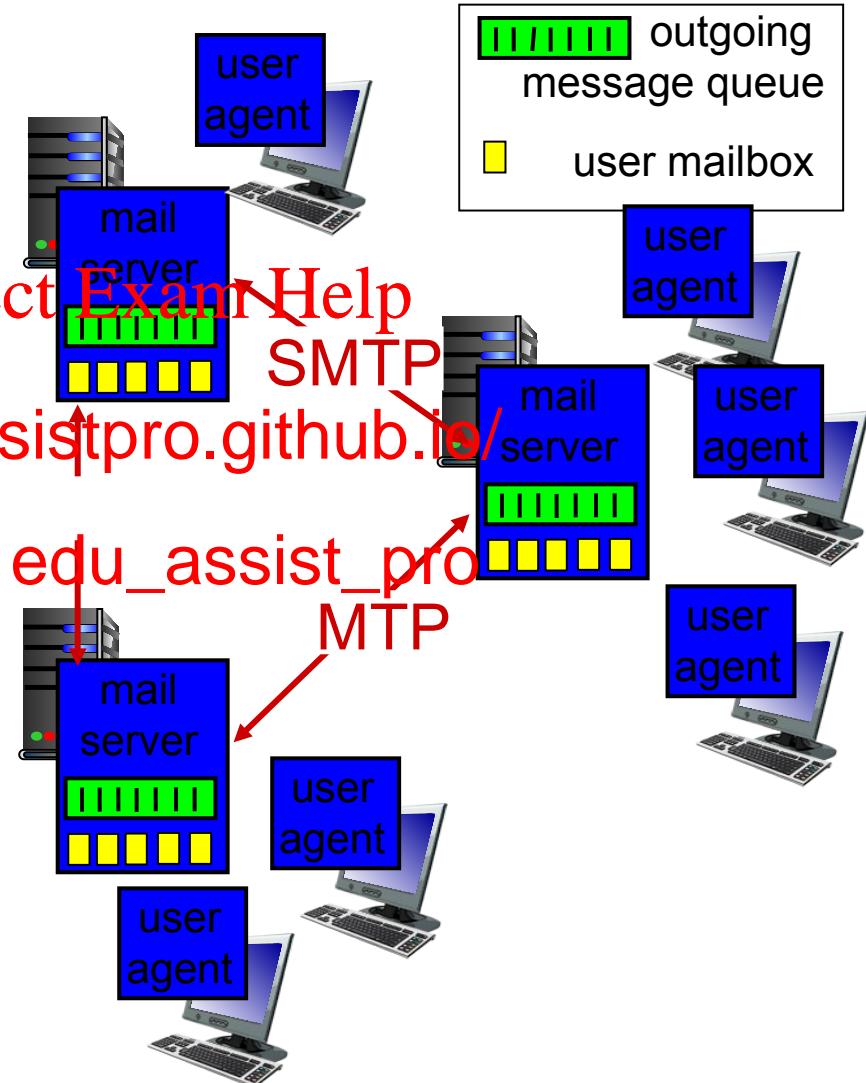
Add WeChat **edu_assist_pro**
SMTP: Simple Mail Transfer Protocol

IMAP: Internet Message Access Protocol

POP3: Post Office Protocol 3

Three major components:

- › user agents (clients)
- › mail servers
- › simple mail transfer protocol:
SMTP



User Agent

- › a.k.a. “mail reader”
- › composing, editing, reading mail messages
- › e.g., Outlook, Thunderbird, iPhone mail client

mail servers:

- › *mailbox* contains incoming messages for user
- › *message queue* of (to be sent) mail <https://eduassistpro.github.io/>
- › *SMTP protocol* to send email messages between mail servers
 - client: sending mail to server
 - “server”: receiving mail from server



- › uses TCP to reliably transfer email message from client to server, port 25
- › direct transfer: sending server to receiving server
- › three phases of transfer

- handshaking (greeting)

- transfer of messages
- closure

<https://eduassistpro.github.io/>

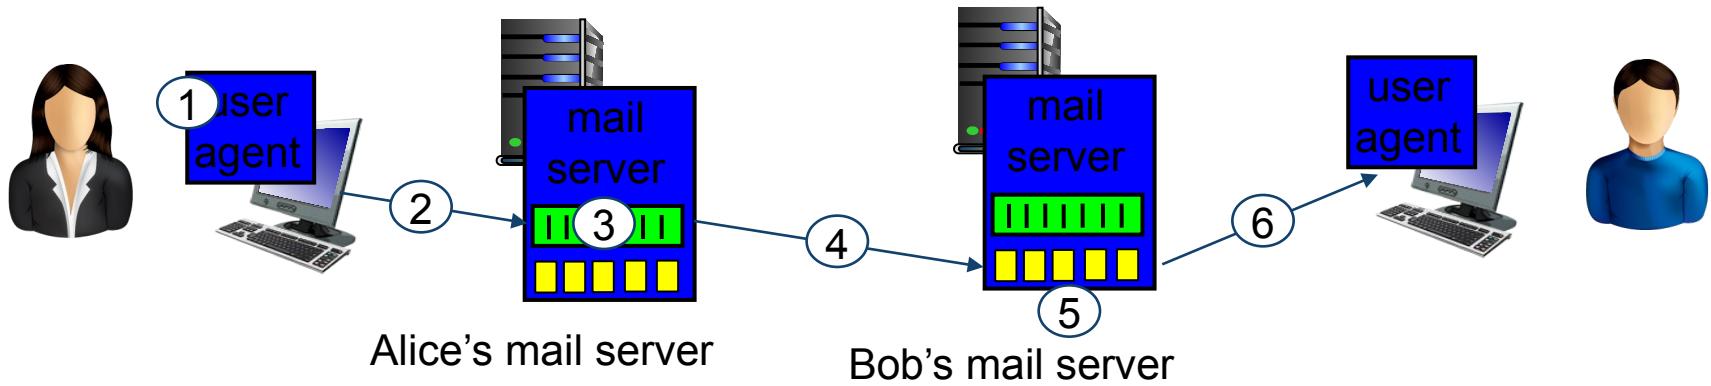
- › command/response interaction (like HTTP)
 - **commands:** ASCII text
 - **response:** status code and phrase

- › messages must be in 7-bit ASCII

- › **Q:** is SMTP stateful or stateless?
 - Stateful

Scenario: Alice sends message to Bob

- 1) Alice uses UA to compose message “to” bob@someschool.edu
 - 2) Alice’s UA sends message to her mail server; message in message queue
 - 3) client side of SMTP opens TCP connection with Bob’s mail server
 - 4) SMTP client sends Alice’s message over the TCP connection
 - 5) Bob’s mail server places the message in Bob’s mailbox
 - 6) Bob’s user agent retrieves his user agent to read the message
- Assignment Project Exam Help**
<https://eduassistpro.github.io/>
Add WeChat **edu_assist_pro**



Sample SMTP interaction

S: 220 hamburger.edu
C: HELO crepes.fr
S: 250 Hello crepes.fr, pleased to meet you
C: MAIL FROM: <alice@crepes.fr>
S: 250 alice@crepes.fr... Sender ok
C: RCPT TO: <bob@hamburger.edu>
t ok S: 250 bob@h
C: DATA <https://eduassistpro.github.io/>
S: 354 Enter mail, end with a blank line by itself
C: Do you like ketchup?
C: Add WeChat edu_assist_pro
C: How about pickles?
C: .
S: 250 Message accepted for delivery
C: QUIT
S: 221 hamburger.edu closing connection

- › SMTP uses persistent connections

- › SMTP requires message (header & body) to be ASCII

- › SMTP server uses CRLF .CRLF to determine end of message

- Carriage return
- Line feed

comparison with HTTP:

- › HTTP: pull

<https://eduassistpro.github.io/>

in SCII response

- › HTTP: each object encapsulated in its own response msg
- › SMTP: multiple objects sent in one msg

SMTP: protocol for exchanging
email msgs

RFC 822: standard for text
message format:

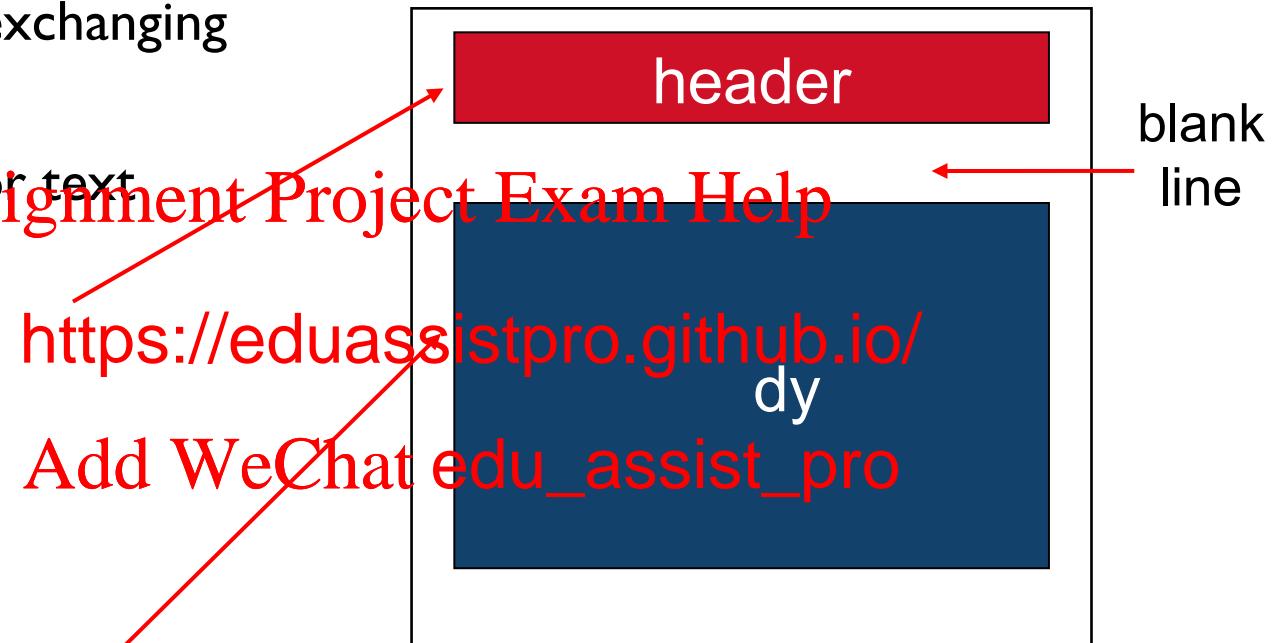
› header lines, e.g.,

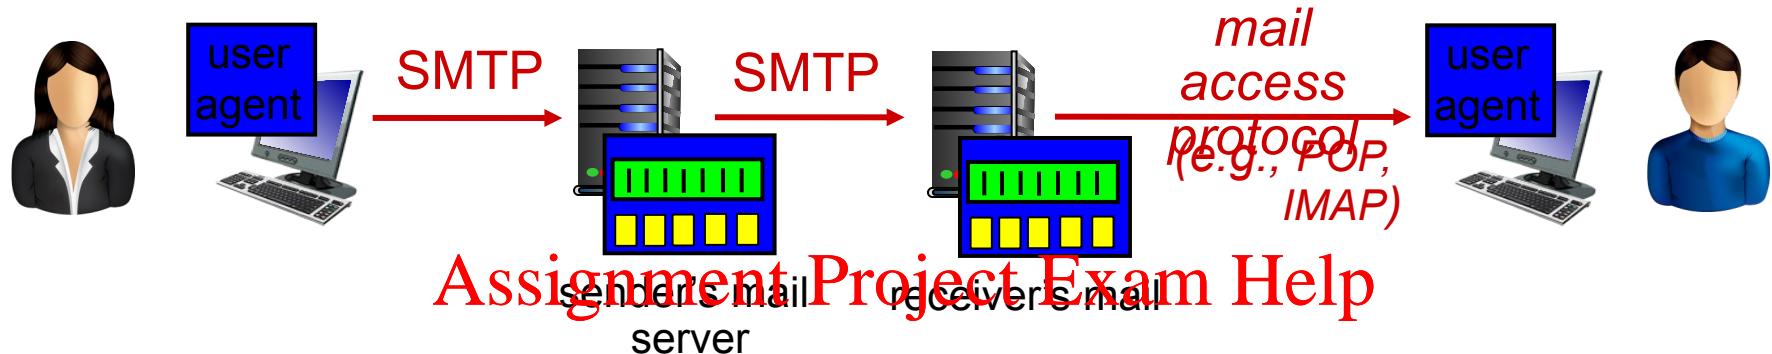
- To:
- From:
- Subject:

different from SMTP MAIL FROM,
RCPT TO: commands!

› Body: the “message”

- ASCII characters only





- › **SMTP:** delivery/storage <https://eduassistpro.github.io/>
- › mail access protocol: retrieval from server
 - **POP:** Post Office Protocol [RFC 1939]: download
 - **IMAP:** Internet Mail Access Protocol [RFC 1730]: more features, including manipulation of stored msgs on server
 - **HTTP:** Using a browser to access a webmail <https://webmail.sydney.edu.au>

Add WeChat **edu_assist_pro**

authorization phase

› client commands:

- **user**: declare username
- **pass**: password

› server responses

- +OK
- -ERR

transaction phase, client:

- › **list**: list message numbers
- › **retr**: retrieve message by number
- › **dele**: delete
- › **quit**

```
S: +OK POP3 server ready
C: user bob
S: +OK
C: pass hungry
S: +OK user successfully logged on
C: list
```

8
https://eduassistpro.github.io/

Add WeChat → edu_assist_pro

ge 1 contents >

```
S: .
C: dele 1
C: retr 2
S: <message 1 contents>
S: .
C: dele 2
C: quit
S: +OK POP3 server signing off
```

more about POP3

- › previous example uses POP3 “download and delete” mode
 - Bob cannot re-read e changes client
- › POP3 “download-and-keep”: copies of messages on different clients
- › POP3 is stateless across sessions

IMAP

- › keeps all messages in one place: at server
 - › to organize

<https://eduassistpro.github.io/>

- › kept state across sessions
 - names of folders and mappings between message IDs and folder name



THE UNIVERSITY OF
SYDNEY

Assignment Project Exam Help
DNS

<https://eduassistpro.github.io/>

Add WeChat **edu_assist_pro**

Internet hosts, routers:

- IP address (32 bit) - used for addressing datagrams
- “name”, e.g., www.yahoo.com - used by humans

people: many identi

- name, passport #

Q: how to map between IP address and name, and vice versa ?

Domain Name System:

› **distributed database**

implemented in hierarchy of many **name servers**

<https://eduassistpro.github.io/>

hyper protocol: hosts,

Add WeChat  edu_assist_pro

communicate to (address/name

translation)

DNS services

- › hostname to IP address translation
- › host aliasing
 - canonical, alias na <https://eduassistpro.github.io/>
- › mail server aliasing
- › load distribution
 - replicated Web servers: many IP addresses correspond to one name

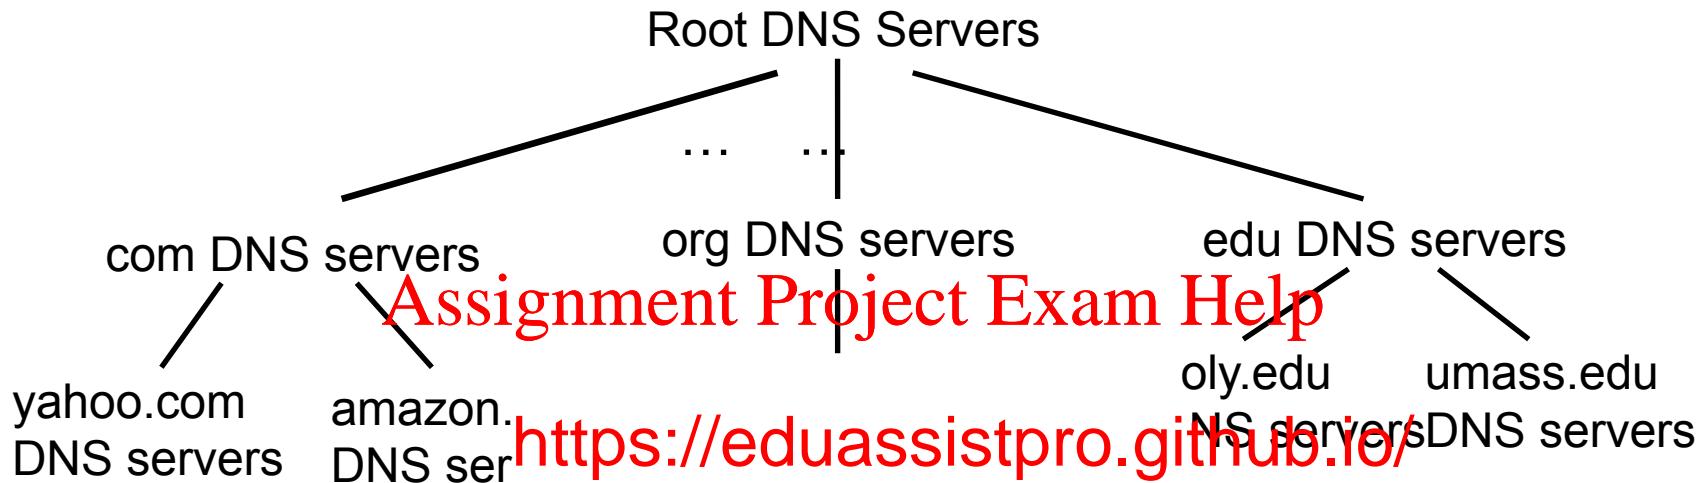
why not centralize DNS?

- › single point of failure
- › distant centralized database

Assignment Project Exam Help

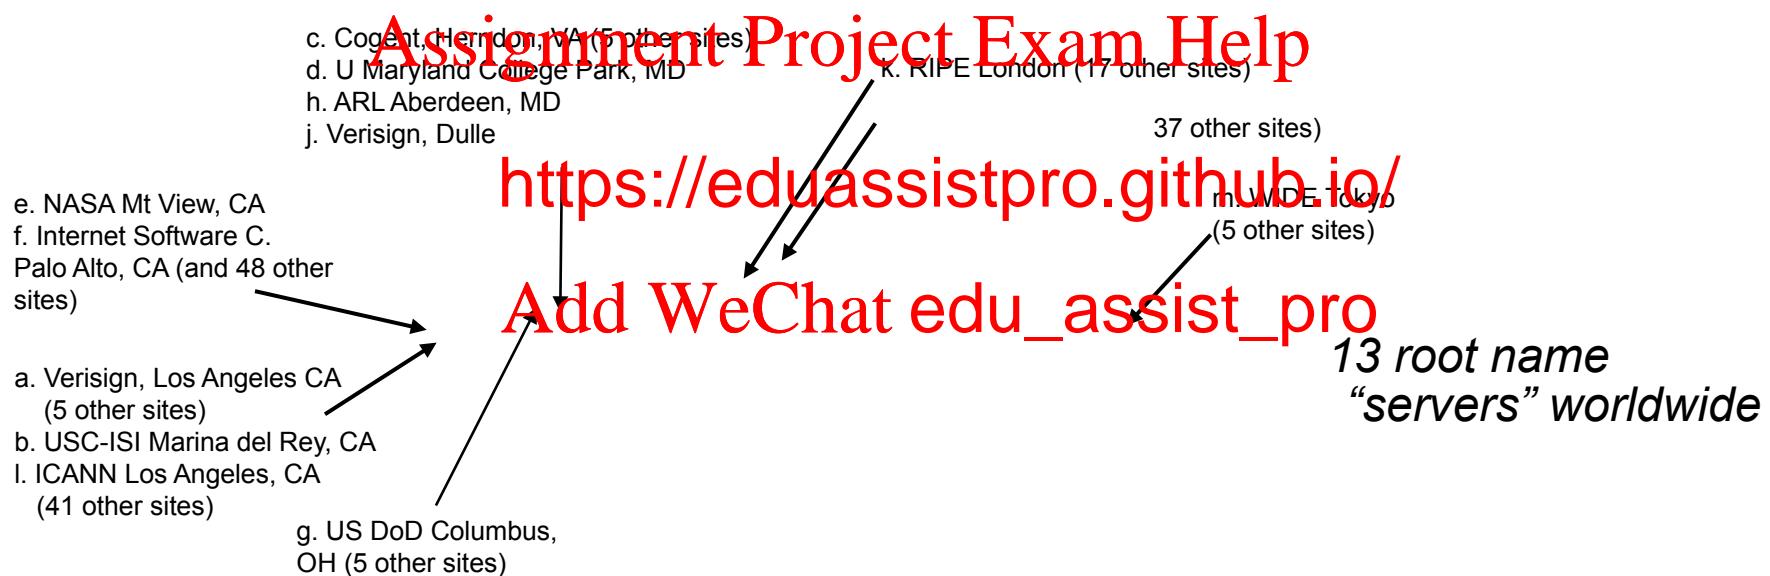
Add WeChat edu_assist_pro

DNS: a distributed, hierarchical database



Add WeChat **edu_assist_pro**
client wants IP for www.amazon.com;

- › client queries root server to find com DNS server
- › client queries .com DNS server to get amazon.com DNS server
- › client queries amazon.com DNS server to get IP address for www.amazon.com



top-level domain (TLD) servers:

- responsible for com, org, net, edu, aero, jobs, museums, and all top-level country domains, e.g.: uk, fr, ca, jp
- Network Solutions maintains servers for .com TLD
- Educause for .edu TL <https://eduassistpro.github.io/>

authoritative DNS servers:

- organization's own DNS server(s), providing hostname to IP mappings for organization's named hosts
- can be maintained by organization or service provider

- › does not strictly belong to hierarchy
- › each ISP (residential ISP, company, university) has one
 - also called “default name server”
- › when host makes DNS query, query is sent to its local DNS server
 - has local cache of recent <https://eduassistpro.github.io/> (but may be out of date!)
 - acts as proxy, forwards query into hierarchy

Assignment Project Exam Help

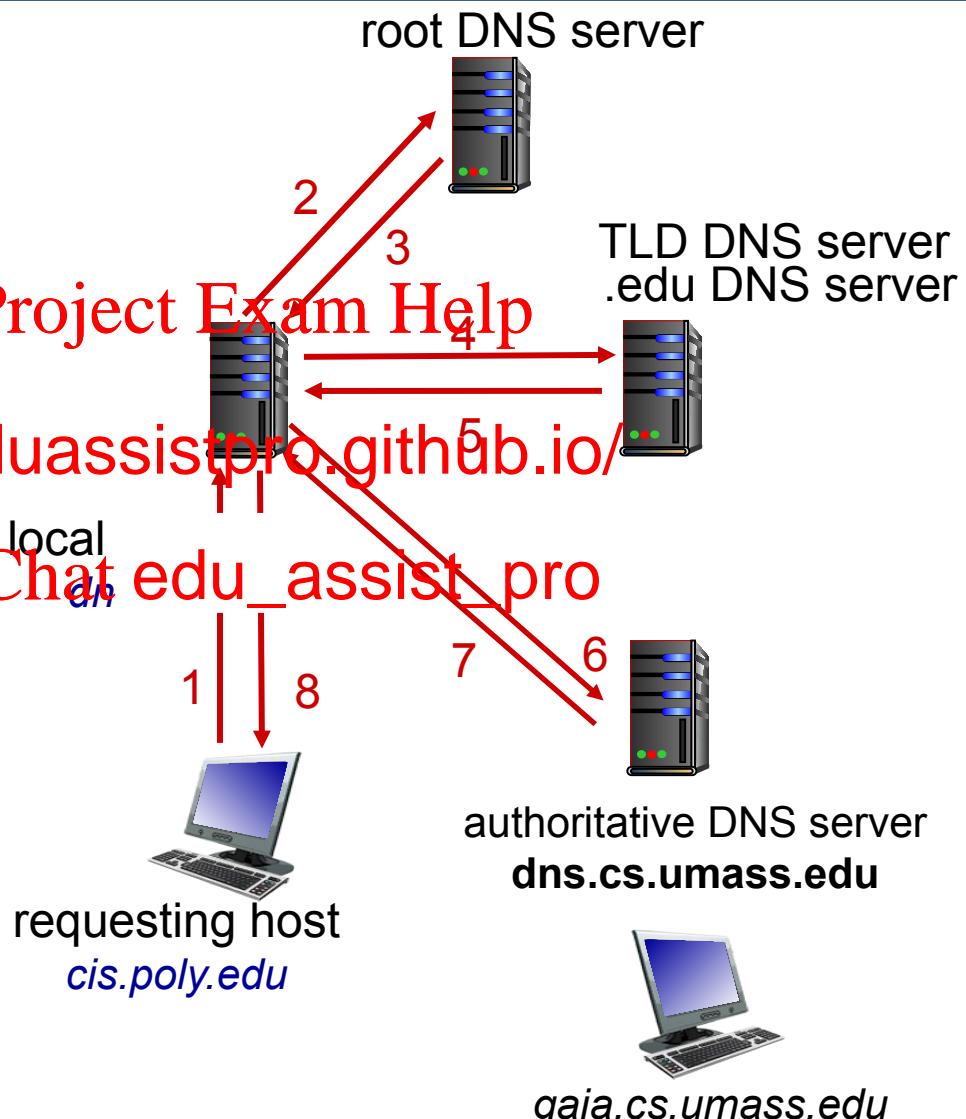
Add WeChat edu_assist_pro

DNS name resolution example

- › host at cis.poly.edu wants IP address for gaia.cs.umass.edu

iterated query: <https://eduassistpro.github.io/>

- ❖ contacted server replies with name of server to contact
- ❖ “I don’t know this name, but ask this server”

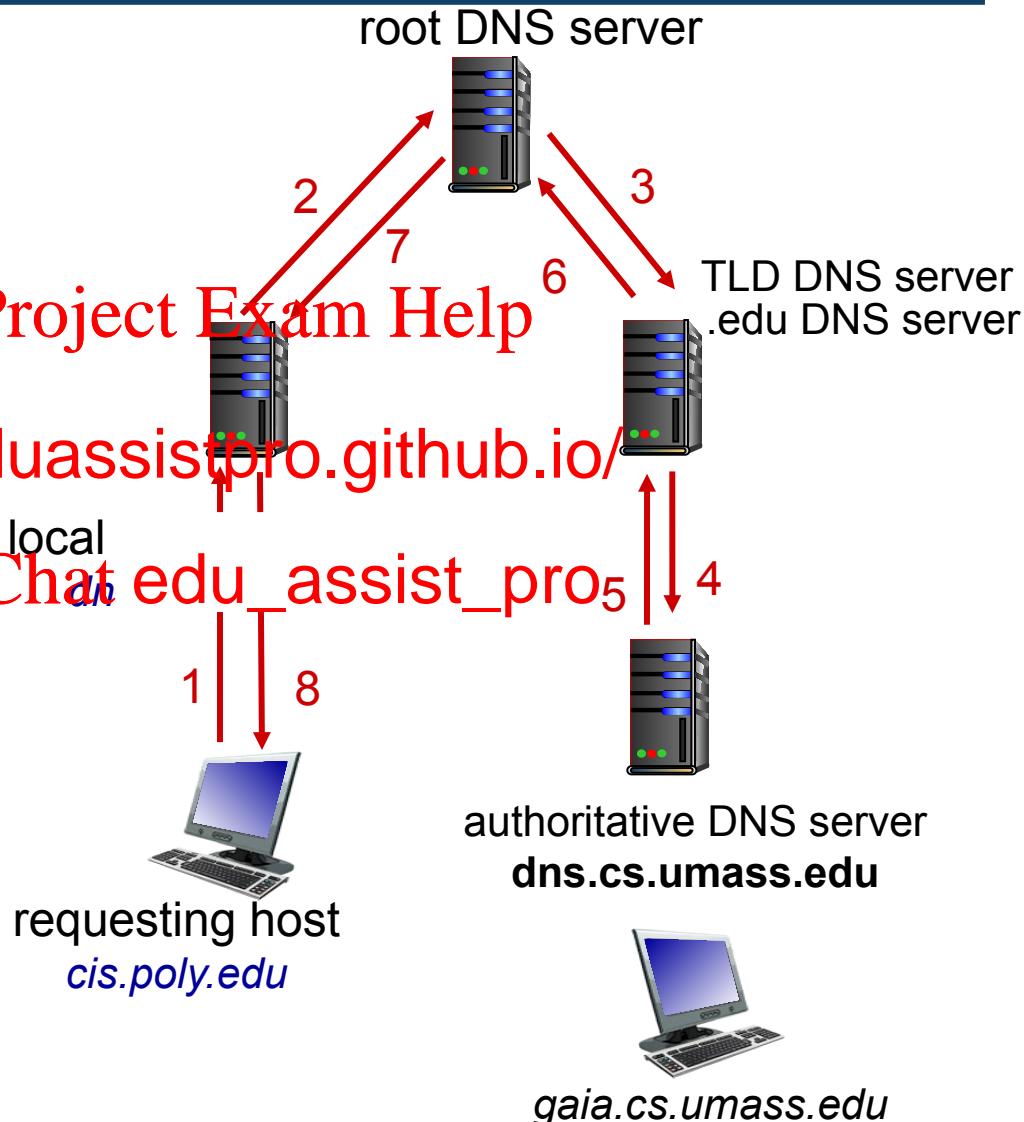


DNS name resolution example (cont'd)

recursive query:

- ❖ puts burden of name resolution on contacted name server
- ❖ heavy load at upper levels of hierarchy?

<https://eduassistpro.github.io/>



- › once (any) name server learns mapping, it *caches* mapping
 - cache entries timeout (disappear) after some time (TTL)
- › cached entries may be *out-of-date* (best effort name-to-address translation!) **Assignment Project Exam Help**
 - if name host changes expire <https://eduassistpro.github.io/> ternet-wide until all TTLs
- › update/notify mechanisms proposed I^E
Add WeChat edu_assist_pro
 - RFC 2136

DNS: distributed db storing resource records (RR)

RR format: (**name**, **value**, **type**, **ttl**)

Assignment Project Exam Help

type=A

- **name** is hostname
- **value** is IP address

<https://eduassistpro.github.io/>

http://namefor some
" (the r

type=NS

- **name** is domain (e.g., foo.com)
- **value** is hostname of authoritative name server for this domain

Add WeChat edu_assist_pro

.com is really
t.backup2.ibm.com

- **value** is canonical name

type=MX

- **value** is name of mailserver associated with **name**

- › example: new startup “Network Utopia”
- › register name networkuptopia.com at *DNS registrar* (e.g., Network Solutions)
 - provide names, IP addresses of authoritative name server
 - registrar inserts two RRs into .com TLD server:

Assignment Project Exam Help

NS) (networkutopia.co

(dns1.networkutop

<https://eduassistpro.github.io/>

- › create at authoritative server

Add WeChat edu_assist_pro

type A record for www.networkuptopia.com;

(www.networkutopia.com, 212.212.212.22, A)

(www.home.networkutopia.com, www.networkutopia.com, CNAME)



Assignment Project Exam Help Socket Programming

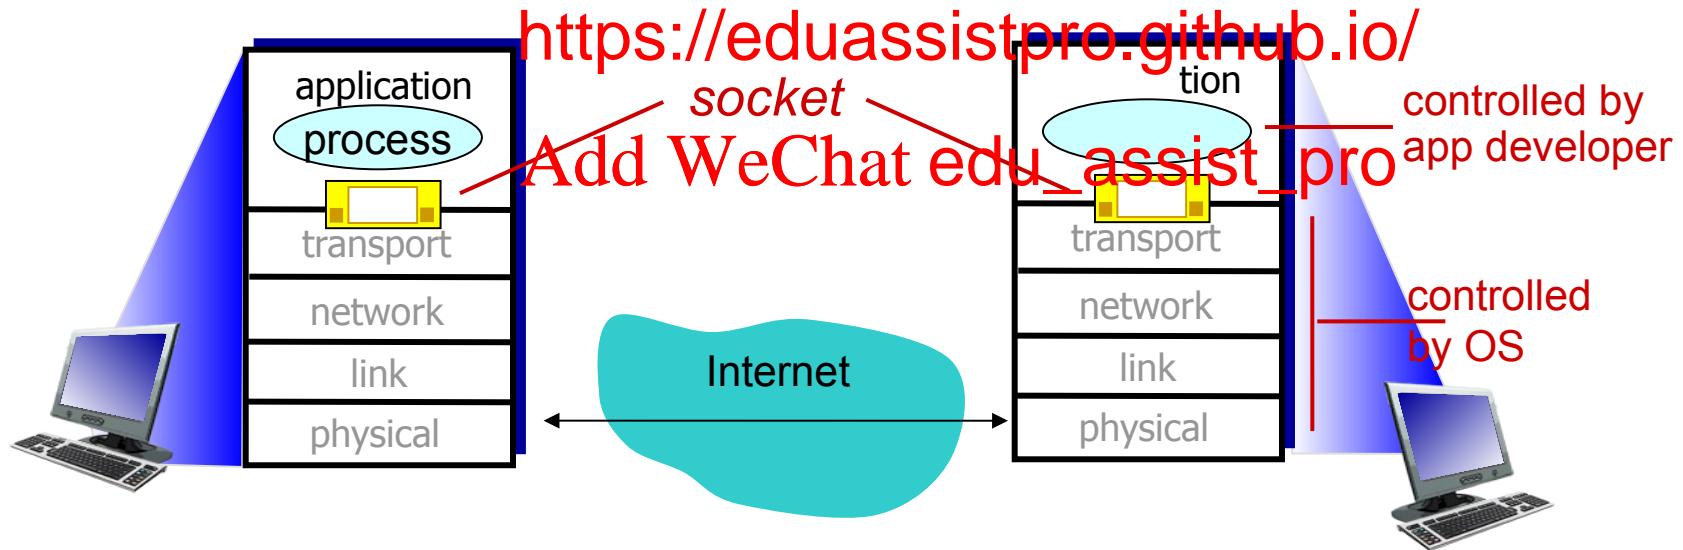
<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

goal: learn how to build client/server applications that communicate using sockets

socket: door between application process and end-end-transport protocol

Assignment Project Exam Help



Two socket types for two transport services:

- **UDP:** unreliable datagram
- **TCP:** reliable, stream-oriented

Assignment Project Exam Help <https://eduassistpro.github.io/>

1. Client reads a character (data) from its keyboard and sends the data to the server.
2. The server receives the data, converts characters to uppercase.
3. The server sends the modified data to the client.
4. The client receives the modified data and displays the line on its screen.

UDP: no “connection” between client & server

- › no handshaking before sending data
- › sender explicitly attaches IP destination address and port # to each packet
- › receiver extracts sender information from received packet

Assignment Project Exam Help
https://eduassistpro.github.io/
Add WeChat edu_assist_pro
UDP: transmitted data may be lost and out-of-order

Application viewpoint:

- › UDP provides *unreliable* transfer of groups of bytes (“datagrams”) between client and server

server (running on serverIP)

create socket, port= x:

serverSocket
~~Assignment Project Exam Help~~
socket(AF_INET,SOCK_DGRAM)

read datagram from
serverSocket

write reply to
serverSocket
specifying
client address,
port number

client

create socket:

clientSocket =

socket(AF_INET,SOCK_DGRAM)

create datagram with server IP and
x; send datagram via

clientSocket

read datagram from
clientSocket

close
clientSocket

<https://eduassistpro.github.io/>
Add WeChat edu_assist_pro

Python UDPCClient

```

include Python's socket library → from socket import *
serverName = 'hostname'
serverPort = 12000
.create UDP for server → .AF_INET,
.get user keyboard input → .SOCK_DGRAM)
message = input('se sentence:')
message = messa
Attach server name, port to message; send into socket → clientSocket.sendto(message,(serverName, serverPort))
.read reply characters from socket into string → modifiedMessage, serverAddress =
clientSocket.recvfrom(2048)
.print out received string and close socket → print (modifiedMessage.decode('utf-8'))
clientSocket.close()
convert from string to bytes
convert from bytes to string
New feature in Python 3

```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat **edu_assist_pro**

Python UDPServer

```
from socket import *
serverPort = 12000
create UDP socket → serverSocket = socket(AF_INET, SOCK_DGRAM)
bind socket to local port → bind(serverSocket, ("127.0.0.1", serverPort))
number 12000
loop forever → while 1:
Read from UDP socket into →   message, clientA = serverSocket.recvfrom(2048)
message, getting client's →   message=message.decode('utf-8')
address (client IP and port) →   modifiedMessage = message.upper()
send upper case string →   serverSocket.sendto(modifiedMessage.encode('utf-8'),
back to this client →     clientAddress)
```

Assignment Project Exam Help
<https://eduassistpro.github.io/>
Add WeChat edu_assist_pro

client must contact server

- › server process must first be running
- › server must have **Assignment Project Exam Help** (door) that welcomes contact

client contacts server by:

- › creating TCP socket, connecting server by specifying IP address, port number of server process
- › **client connects:** client TCP establishes connection to server TCP

- › when contacted by client, **server TCP creates new socket** for server process to communicate with that particular client

er to talk with

<https://eduassistpro.github.io/>

numbers used to
clients

application viewpoint:

TCP provides reliable, in-order byte-stream transfer (“pipe”) between client and server

server (running on hostid)

```

create socket,
port=x,
serverSocket = socket()
wait for incoming
connection request
serverSocket.listen
  
```

Accept client

```

connectionSocket =
serverSocket.accept()
  
```

read request from
connectionSocket

write reply to
connectionSocket

close
connectionSocket

client

```

create socket,
connect(hostid, port=x)
socket()
connect(hostid, x)
  
```

send request using
clientSocket

read reply from
clientSocket

close
clientSocket

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Python TCPClient

```
from socket import *
serverName = 'servername'
create TCP socket for
server, remote port 12000
serverPort = 12000
clientSocket = so
clientSocket.con https://eduassistpro.github.io/
sentence = input('Input lowercase
clientSocket.send(sentence.encode())
## Do not spe
No need to attach server
name, port
modifiedSentence = clientSocket.recv(1024)
print ('From Server:', modifiedSentence.decode('utf-8'))
clientSocket.close()
```

Assignment Project Exam Help
EAM)

Add WeChat **edu_assist_pro**

Python TCPServer

```
from socket import *
serverPort = 12000
serverSocket = socket(AF_INET, SOCK_STREAM)
serverSocket.bind(("", serverPort))
```

create TCP welcoming
socket → **Assignment Project Exam Help**

server begins listening for
incoming TCP requests → **https://eduassistpro.github.io/**

loop forever → **while 1:**

server waits on accept()
for incoming requests, new
socket created on return → **connectionSocket, serverSocket.accept()**

read bytes from socket (but
not address as in UDP) → **sentence = connectionSocket.recv(1024)**
capitalizedSentence = sentence.decode('utf-8').upper().encode('utf-8')
connectionSocket.send(capitalizedSentence)

close connection to this
client (but *not* welcoming
socket) → **connectionSocket.close()**



Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Pure peer-to-peer model architecture

- › no always-on server
- › arbitrary end systems directly communicate
- › peers are intermittently connected and change addresses

examples:

- file distribution (BitTorrent)
- Streaming (Zattoo, KanKan)
- VoIP (Skype)

Assignment Project Exam Help

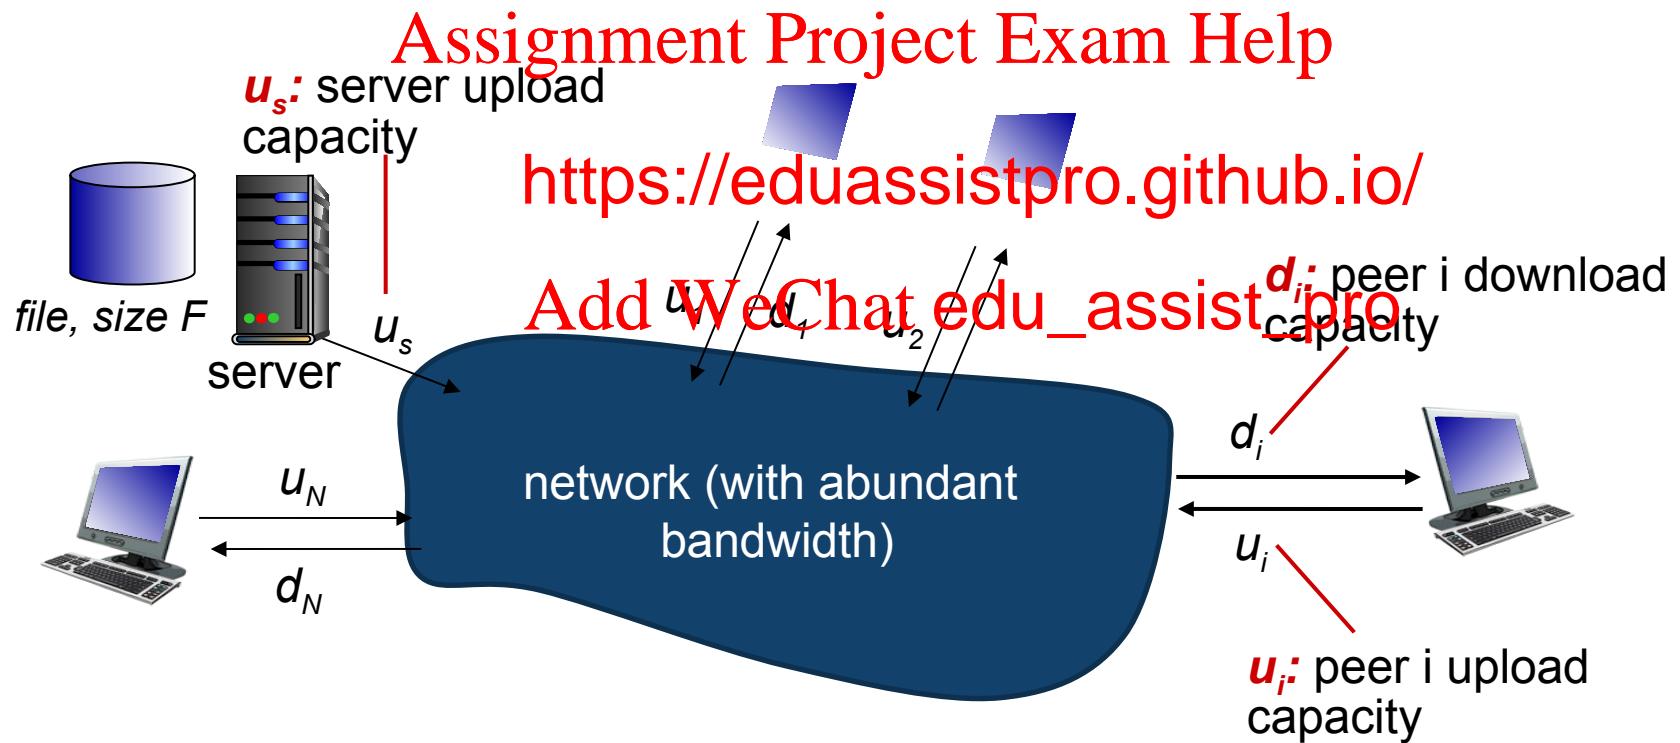
<https://eduassistpro.github.io/>
Add WeChat edu_assist_pro



File distribution: client-server vs. p2p

Question: how much time to distribute file (size F) from one server to N peers?

- peer upload/download capacity is limited resource



File distribution time: client-server

› **server transmission:** must sequentially send (upload) N file copies:

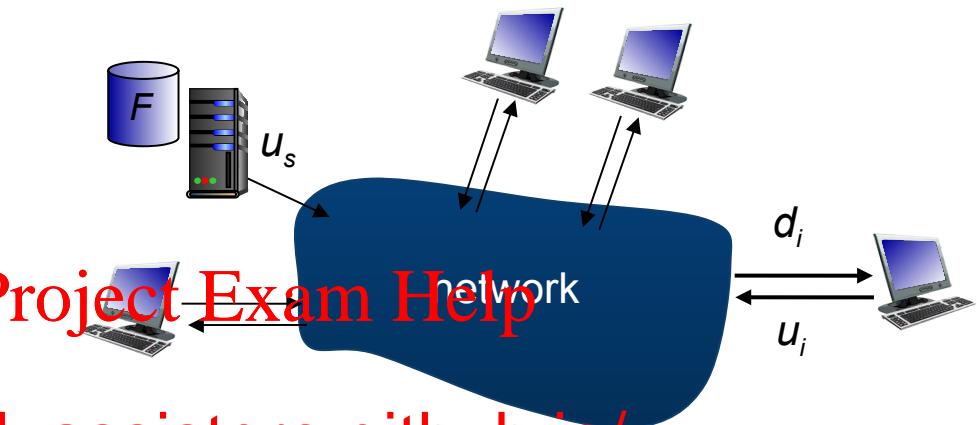
- time to send one copy: F/u_s
- time to send N copies:
❖ **client:** each client <https://eduassistpro.github.io/> download file co

- d_{\min} = min client download rate
- (worst case) client download time:
 F/d_{\min}

*time to distribute F
to N clients using
client-server approach*

$$D_s > \max\{NF/u_s, F/d_{\min}\}$$

increases linearly in N



› **server transmission:** must upload at least one copy

- time to send one copy: F/u_s

- ❖ **client:** each client must download file copy

- client download ti

- ❖ **clients:** as aggregate

- Max upload rate $u_s + \sum u_i$

- $NF/(u_s + \sum u_i)$

time to distribute F

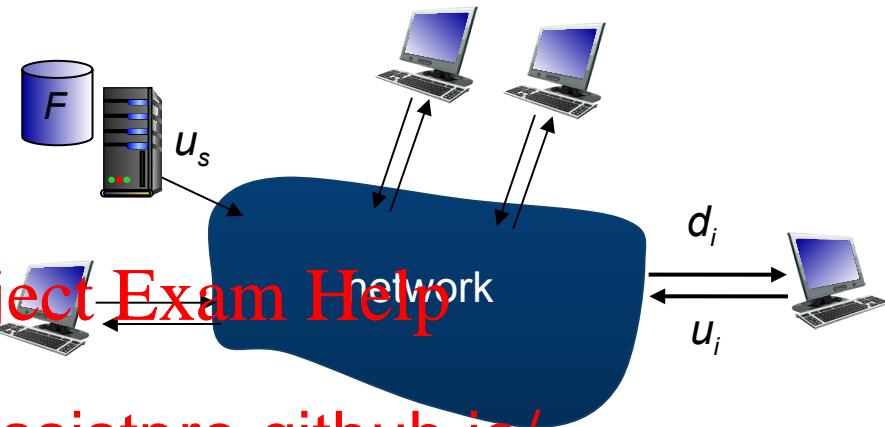
to N clients using

P2P approach

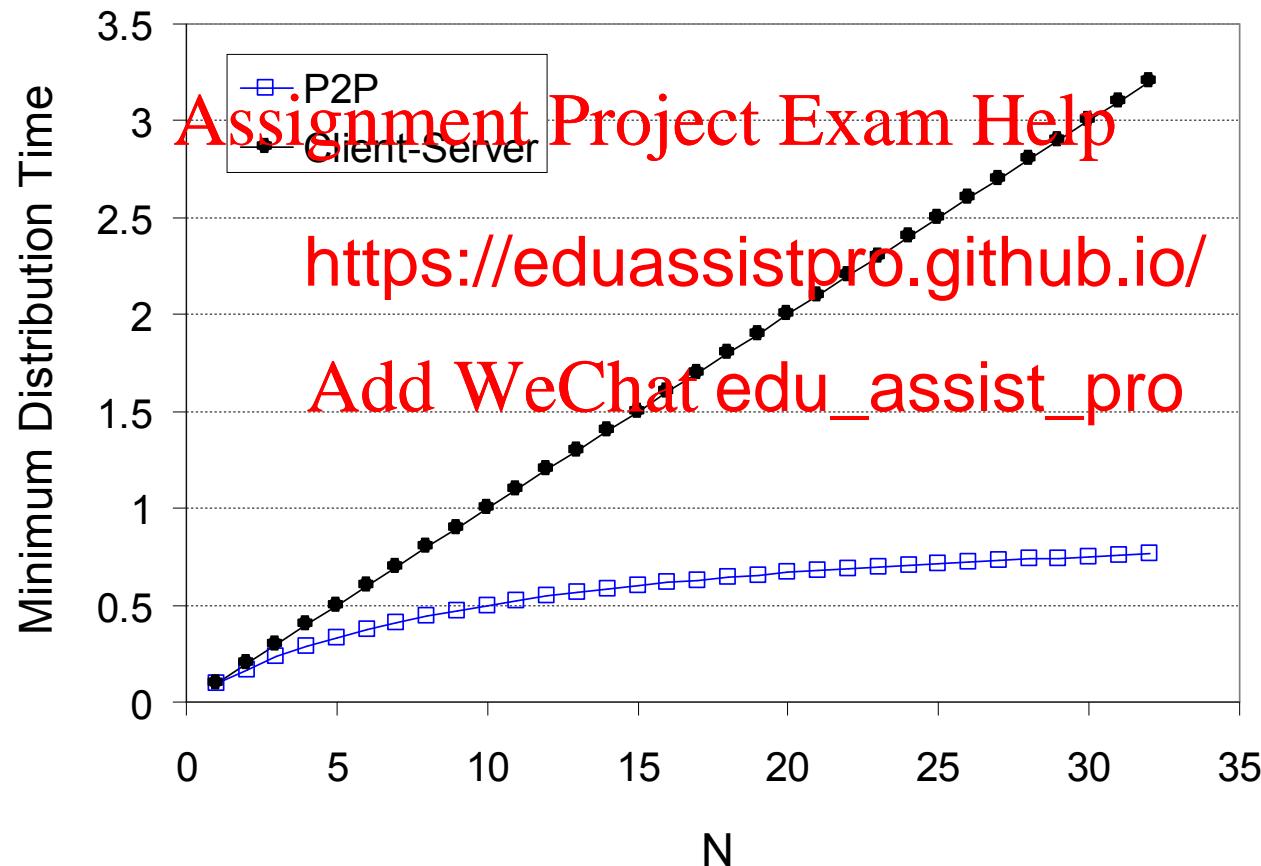
$$D_{P2P} > \max\{F/u_s, F/d_{min}, NF/(u_s + \sum u_i)\}$$

increases linearly in N ...

... but so does this, as each peer brings service capacity



client upload rate = u , $F/u = 1$ hour, $u_s = 10u$, $d_{min} \geq u_s$



BitTorrent, a file sharing application

- › 20% of European internet traffic in 2012.
- › Used for Linux distribution, software patches, distributing movies
- › Goal: quickly replicate large files to large number of clients



<https://eduassistpro.github.io/>

- › Web server hosts a .torrent file (w/ file I tracker's URL...)
- › A tracker tracks downloaders/owners of **Add WeChat edu_assist_pro**
- › Files are divided into chunks (256kB-1MB)
- › Downloaders download chunks from themselves (and owners)
- › Tit-for-tat: the more one shares (**server**), the faster it can download (**client**)

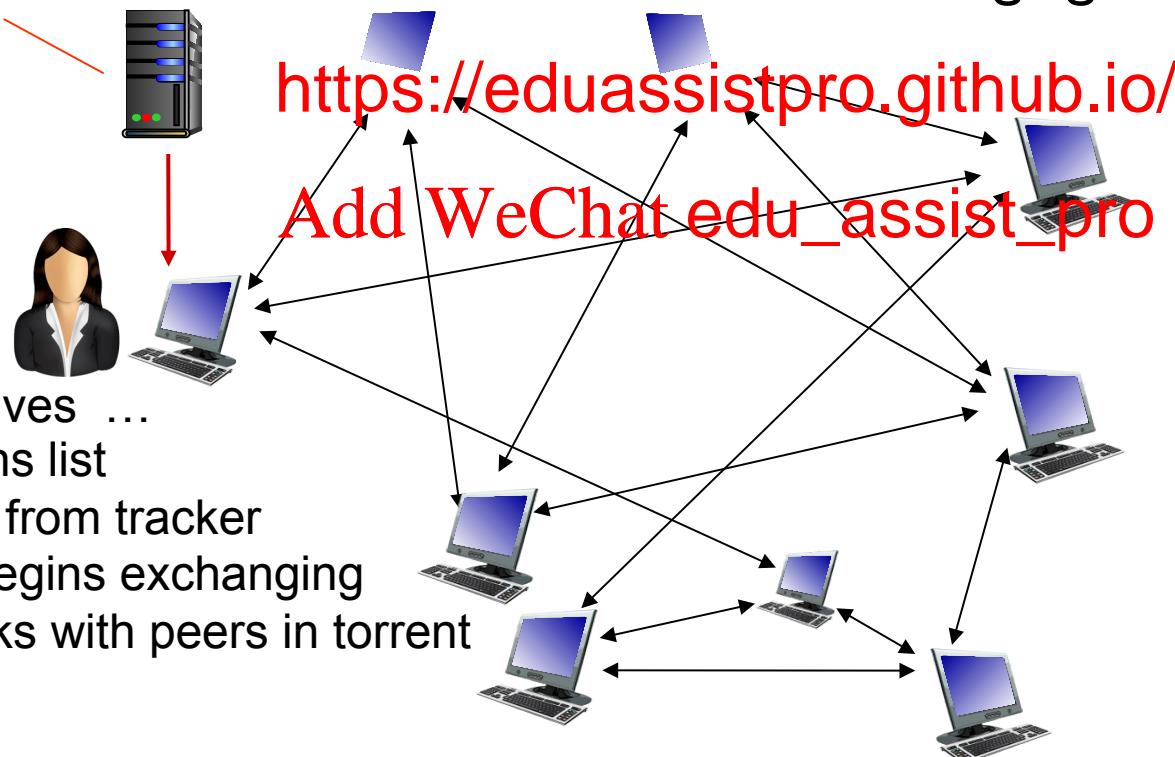
- › file divided into 256KB chunks
- › peers in torrent send/receive file chunks



tracker: tracks peers
participating in torrent

torrent: group of peers
exchanging chunks of a

Alice arrives ...
... obtains list
of peers from tracker
... and begins exchanging
file chunks with peers in torrent



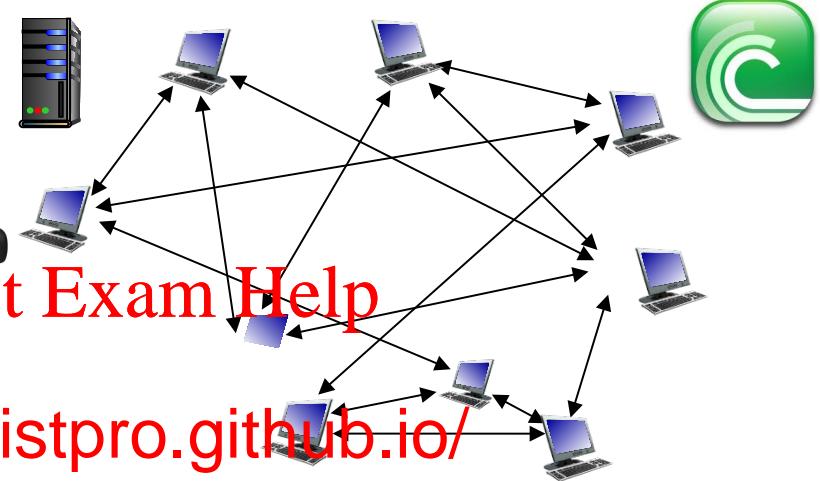
› peer joining torrent:

- has no chunks, but will accumulate them over time from other peers

Assignment Project Exam Help

- registers with tracker
peers, connects to <https://eduassistpro.github.io/>
("neighbors")

Add WeChat edu_assist_pro



- › while downloading, peer uploads chunks to other peers
- › peer may change peers with whom it exchanges chunks
- › *churn*: peers may come and go
- › once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent

requesting chunks:

- › at any given time, different peers have different subsets of file chunks
- › periodically, Alice a peer for list of chunks that they have
- › Alice requests missing chunks from peers, rarest first

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

top 4 every 10 secs

sending chunks: tit-for-tat



- › Alice sends chunks to those four peers currently sending her chunks *at highest rate*

peers are choked by Alice (do one chunk

- › every 30 secs: randomly select another peer, starts sending chunks

› “optimistically unchoke” this peer

› newly chosen peer may join top 4

- (1) Alice “optimistically unchoke” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



Assignment Project Exam Help

