

# Advanced Network Technologies

Application layer

Transport layer

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

Dr. Wei Bao | Lecturer  
School of Computer Science



Drag picture to placeholder or click





# Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Pure peer-to-peer model architecture

- › no always-on server
- › arbitrary end systems directly communicate
- › peers are intermittently connected and change addresses

examples:

- file distribution (BitTorrent)
- Streaming (Zattoo, KanKan)
- VoIP (Skype)

Assignment Project Exam Help

<https://eduassistpro.github.io/>

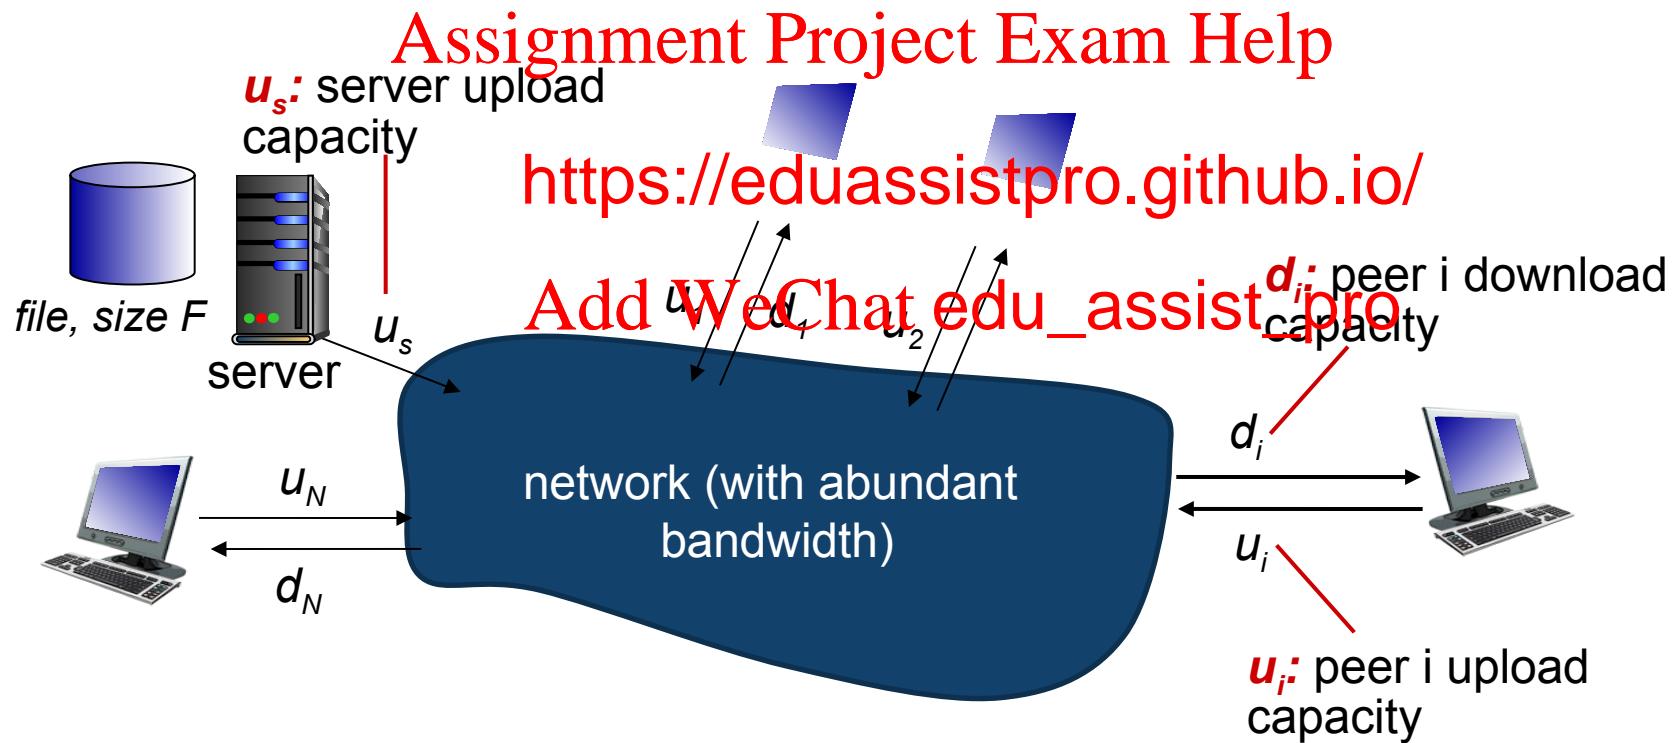
Add WeChat edu\_assist\_pro



# File distribution: client-server vs. p2p

Question: how much time to distribute file (size  $F$ ) from one server to  $N$  peers?

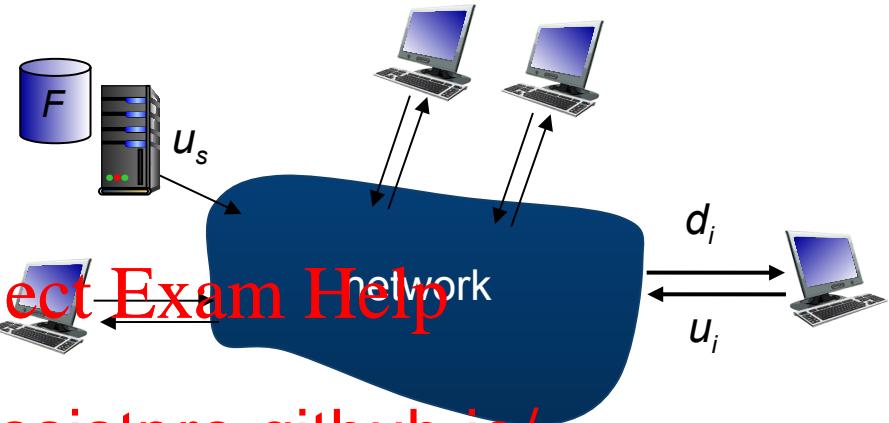
- peer upload/download capacity is limited resource



## File distribution time: client-server

› **server transmission:** must sequentially send (upload)  $N$  file copies:

- time to send one copy:  $F/u_s$
- time to send  $N$  copies:  
❖ **client:** each client <https://eduassistpro.github.io/> download file co
  - $d_{min} = \text{min client download rate}$
  - (worst case) client download time:  $F/d_{min}$



*time to distribute  $F$   
to  $N$  clients using  
client-server approach*

$$D_s > \max\{NF/u_s, F/d_{min}\}$$

increases linearly in  $N$

› **server transmission:** must upload at least one copy

- time to send one copy:  $F/u_s$

- ❖ **client:** each client must download file copy

- client download time

- ❖ **clients:** as aggregate

- Max upload rate  $u_s + \sum u_i$

- $NF/(u_s + \sum u_i)$

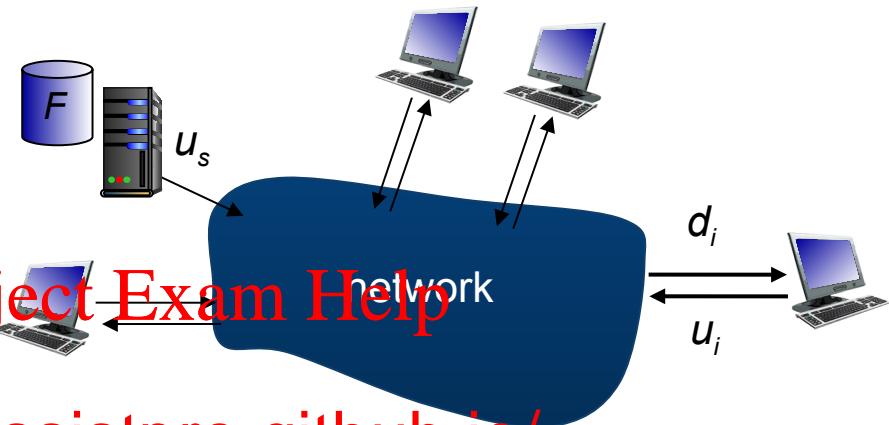
*time to distribute  $F$*

*to  $N$  clients using  
P2P approach*

$$D_{P2P} > \max\{F/u_s, F/d_{min}, NF/(u_s + \sum u_i)\}$$

increases linearly in  $N$  ...

... but so does this, as each peer brings service capacity

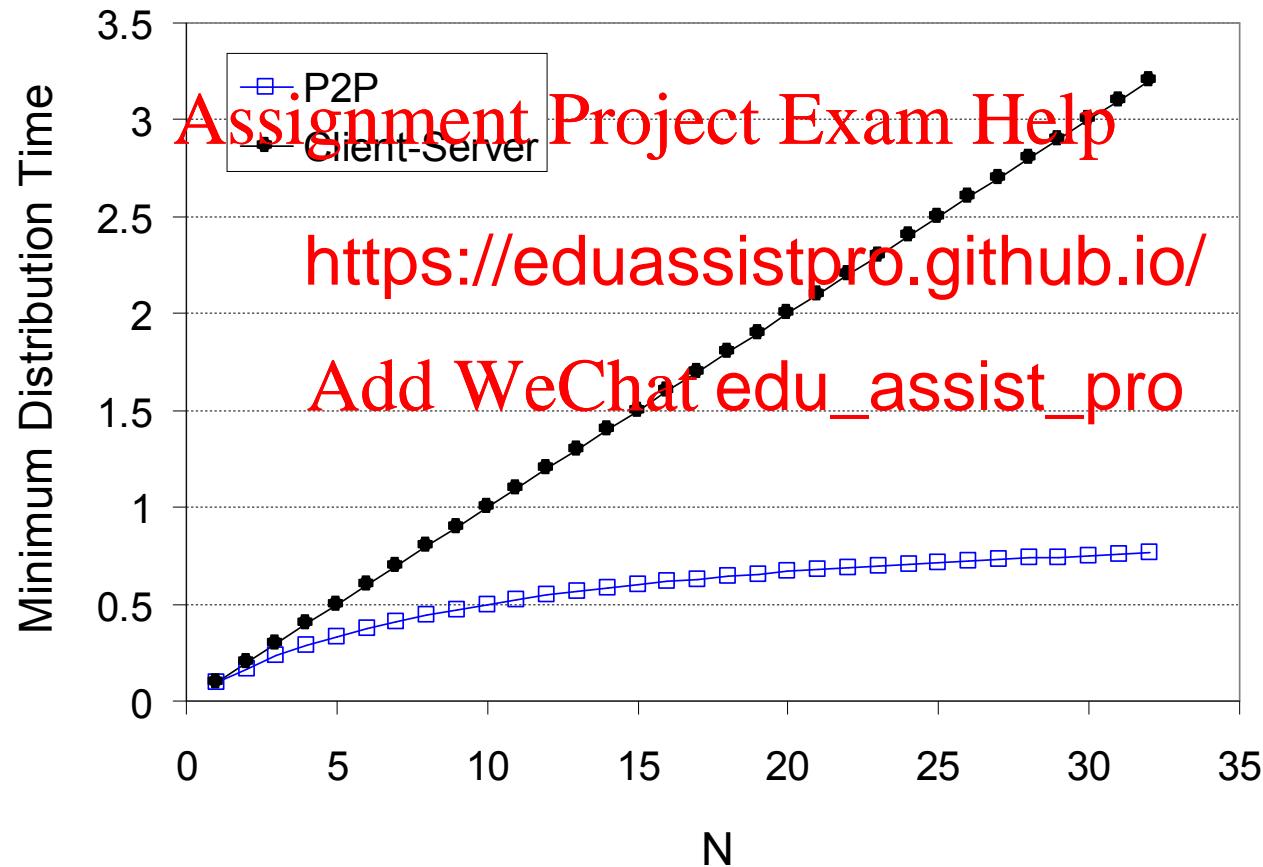


<https://eduassistpro.github.io/>

= upload  $NF$  bits

Add WeChat edu\_assist\_pro

client upload rate =  $u$ ,  $F/u = 1$  hour,  $u_s = 10u$ ,  $d_{min} \geq u_s$



## BitTorrent, a file sharing application

- › 20% of European internet traffic in 2012.
- › Used for Linux distribution, software patches, distributing movies
- › Goal: quickly replicate large files to large number of clients



<https://eduassistpro.github.io/>

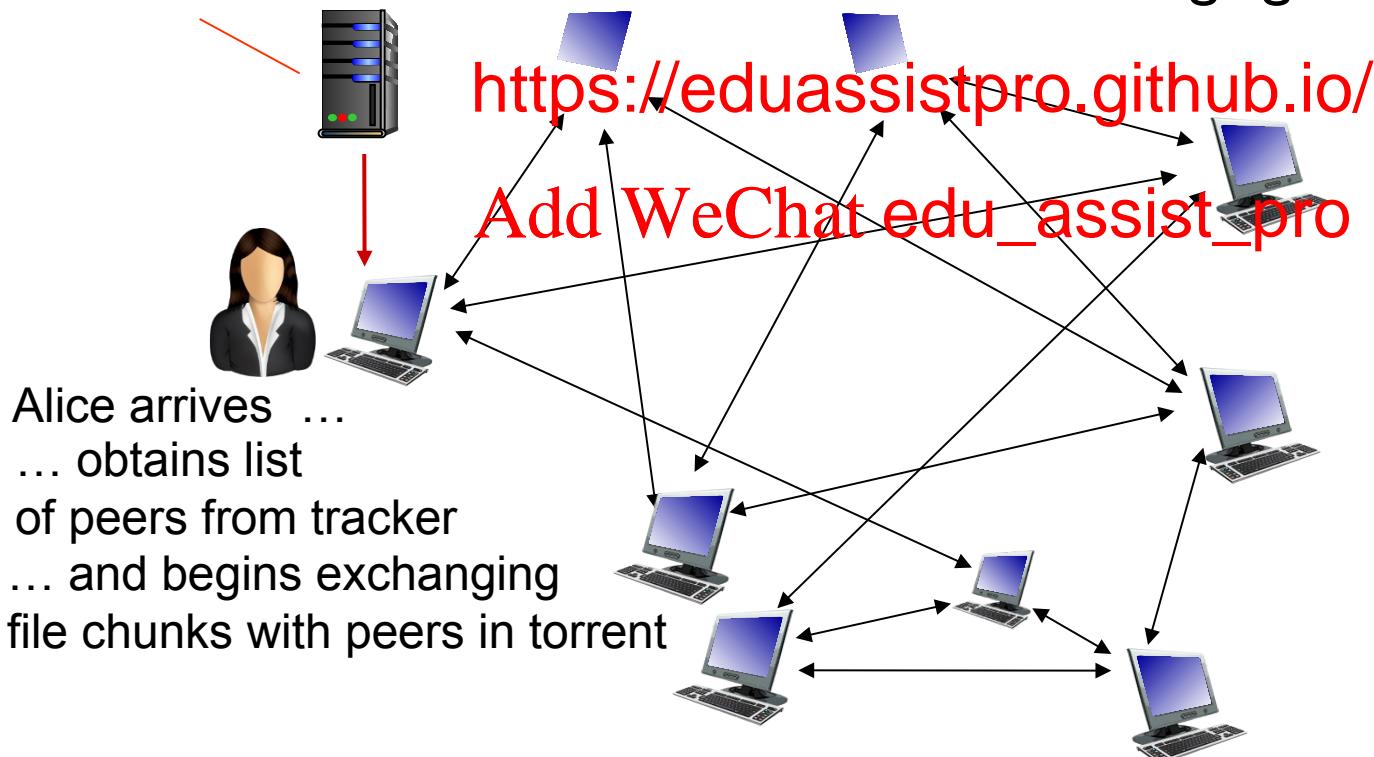
- › Web server hosts a .torrent file (w/ file I tracker's URL...)
- › A tracker tracks downloaders/owners of **Add WeChat edu\_assist\_pro**
- › Files are divided into chunks (256kB-1MB)
- › Downloaders download chunks from themselves (and owners)
- › Tit-for-tat: the more one shares (**server**), the faster it can download (**client**)

- › file divided into 256KB chunks
- › peers in torrent send/receive file chunks



*tracker*: tracks peers  
participating in torrent

*torrent*: group of peers  
exchanging chunks of a



## › peer joining torrent:

- has no chunks, but will accumulate them over time from other peers

Assignment Project Exam Help

- registers with tracker

peers, connects to <https://eduassistpro.github.io/> (“neighbors”)

Add WeChat edu\_assist\_pro



› while downloading, peer uploads chunks to other peers

› peer may change peers with whom it exchanges chunks

› *churn*: peers may come and go

› once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent

## requesting chunks:

- › at any given time, different peers have different subsets of file chunks
- › periodically, Alice a peer for list of chunks that they have
- › Alice requests missing chunks from peers, rarest first

**Assignment Project Exam Help**

<https://eduassistpro.github.io/>

Add WeChat **edu\_assist\_pro**

## sending chunks: tit-for-tat

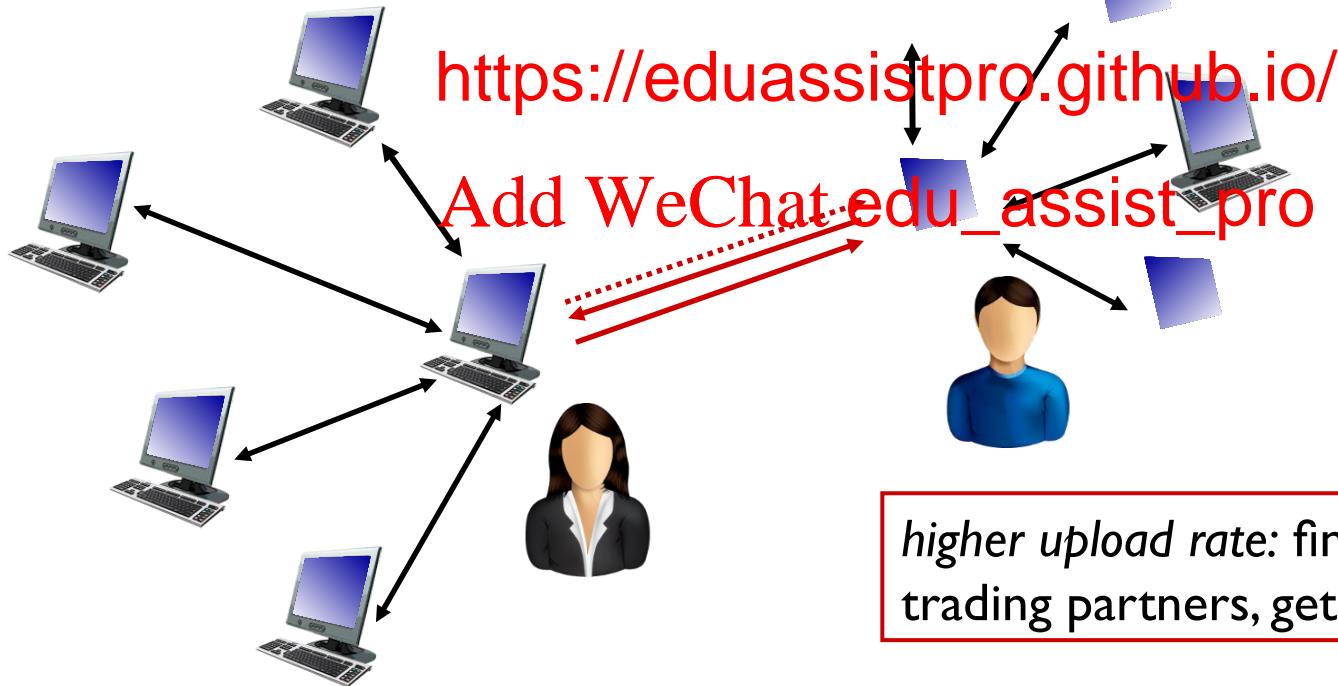
- › Alice sends chunks to those four peers currently sending her chunks *at highest rate*
- › Alice is *choke* (do not receive chunks from her)  
4 every 10 secs
- › every 30 secs: randomly select another peer, starts sending chunks
  - › “optimistically unchoke” this peer
  - › newly chosen peer may join top 4



- (1) Alice “optimistically unchoke” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



## Assignment Project Exam Help





# Distributed Hash Table

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

- › DHT: a *distributed P2P database*
- › database has **(key, value)** pairs; examples:
  - key: social **Assignment Project Exam Help** human name
- › distribute th <https://eduassistpro.github.io/> over the many peers **Add WeChat edu\_assist\_pro**
- › a peer **queries** DHT with key
  - DHT returns values that match the key
- › peers can also **insert** **(key, value)** pairs

- › Assign the keys
- › Lookup the keys

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

## › central issue:

- assigning (key, value) pairs to peers.

## › basic idea:

- Key: generate <https://eduassistpro.github.io/>
- Assign an integer ID to each ~~Add WeChat~~ `edu_assist_pro`
- put (key,value) pair in the peer that is closest to the key

- › **distance**: assign integer identifier to each peer in range  $[0, 2^n - 1]$  for some  $n$ .
  - each identifier represented by  $n$  bits.
- › Each key to be <https://eduassistpro.github.io/>  $[0, 2^n - 1]$
- › to get integer key, hash ori [Add WeChat edu\\_assist\\_pro](#)
  - A hash function is any function that can be used to map data of arbitrary size to data of fixed size (e.g., an integer in  $[0, 2^n - 1]$  ).
  - e.g., 15 = **hash**("Led Zeppelin IV")
  - this is why its is referred to as a *distributed "hash" table*

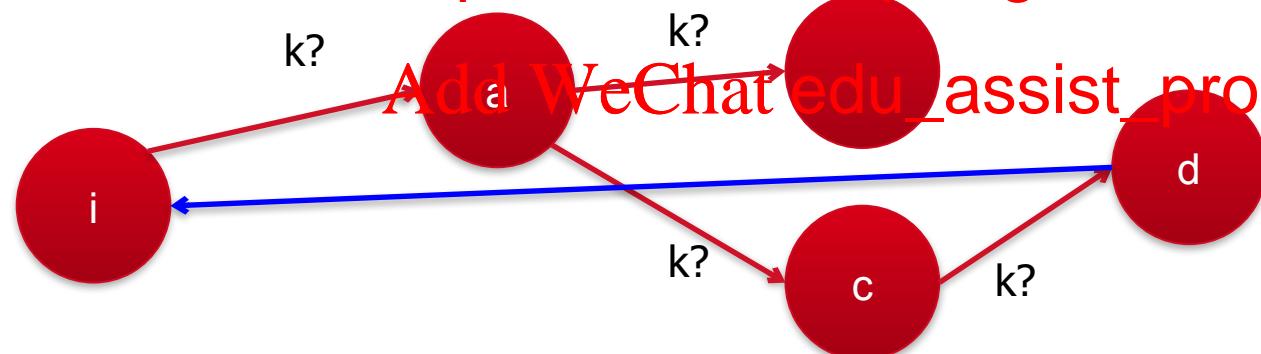
- › rule: assign key to the peer that has the *closest* ID.
- › Here: closest is the ~~Assignment Project Extra Help or~~ predecessor of the key.  
<https://eduassistpro.github.io/>
- › e.g.,  $n = 4$ ; peers: 1, 3, 4, 12, 14:  
~~Add WeChat edu\_assist\_pro~~
- key = 13, then successor peer = 14
- key = 15, then successor peer = 1

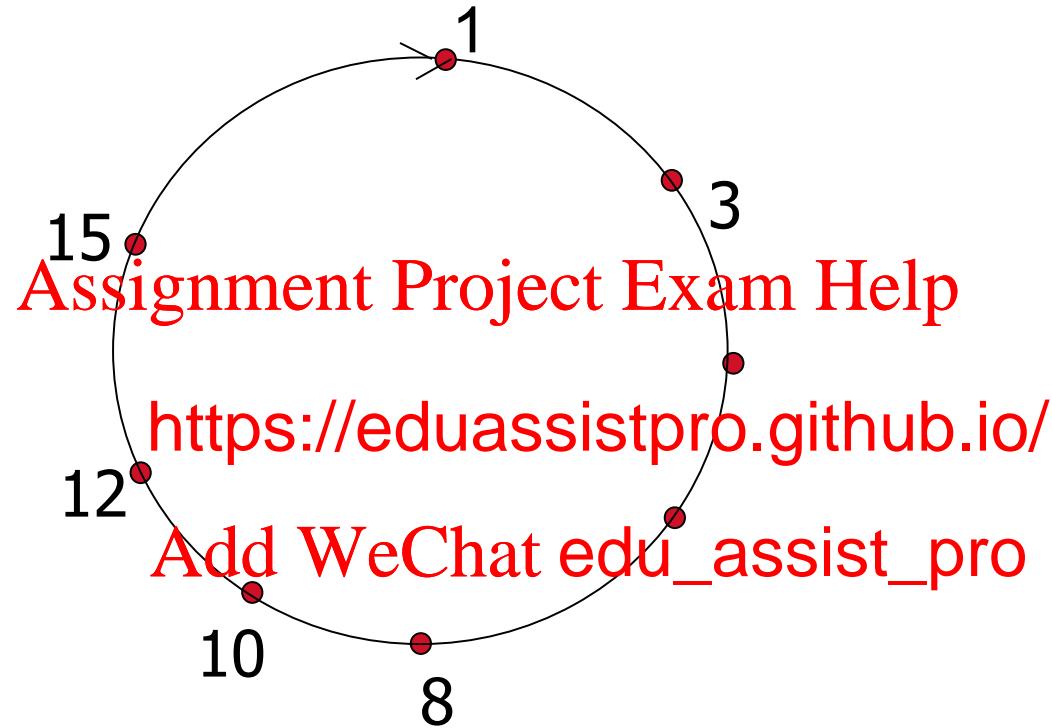
Goal: to provide a distributed lookup service returning the host that owns the key

- Given a key, find the host that owns the key

**Assignment Project Exam Help**

<https://eduassistpro.github.io/>



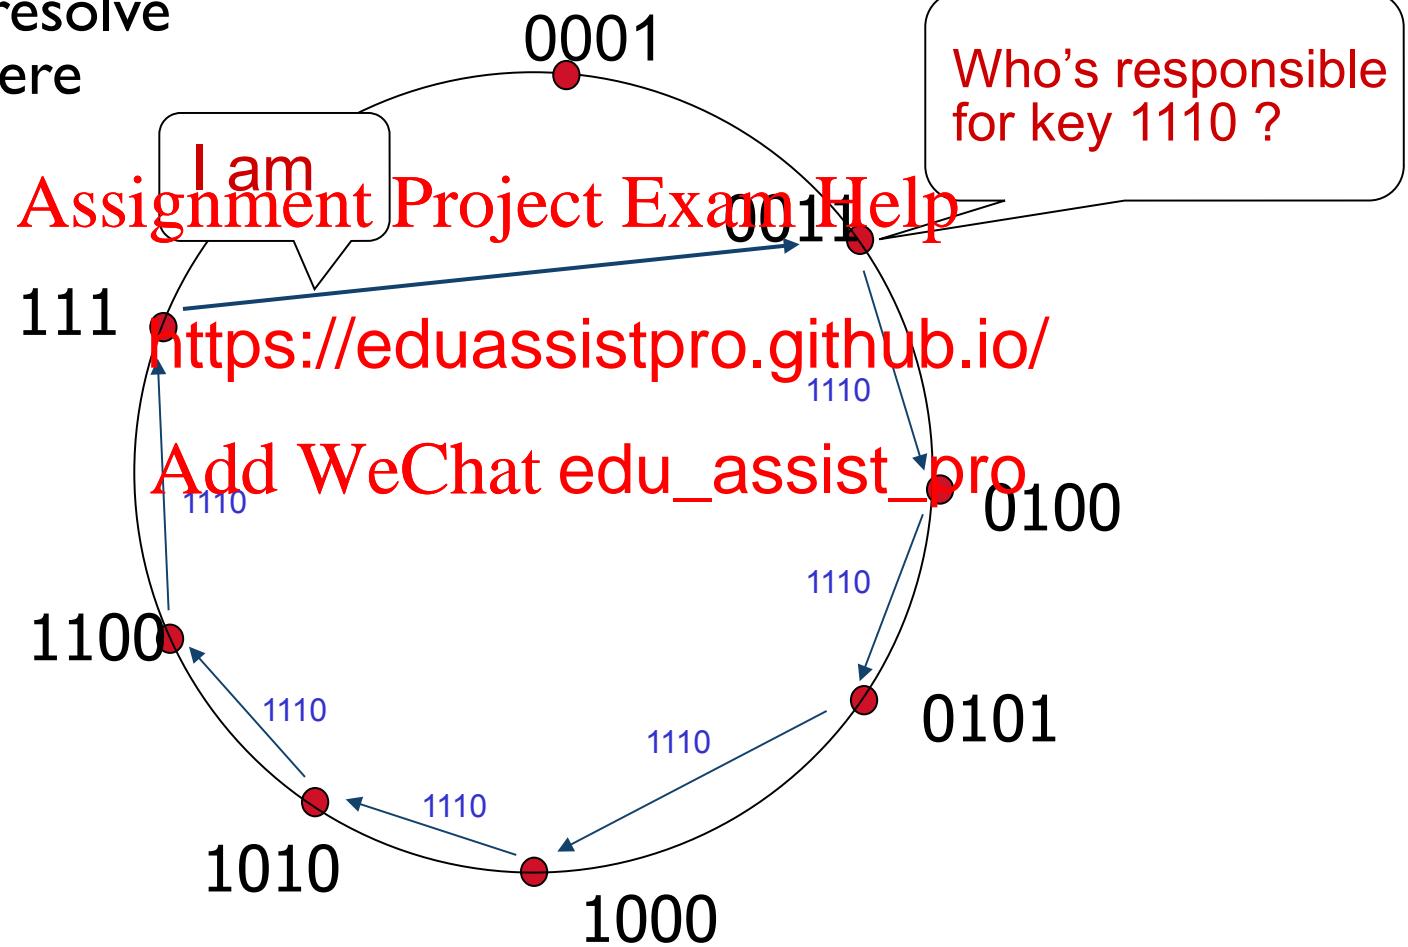


- › each peer *only* aware of immediate successor.

$O(N)$  messages  
on average to resolve  
query, when there  
are  $N$  peers

key 1110 is  
stored at node  
1111

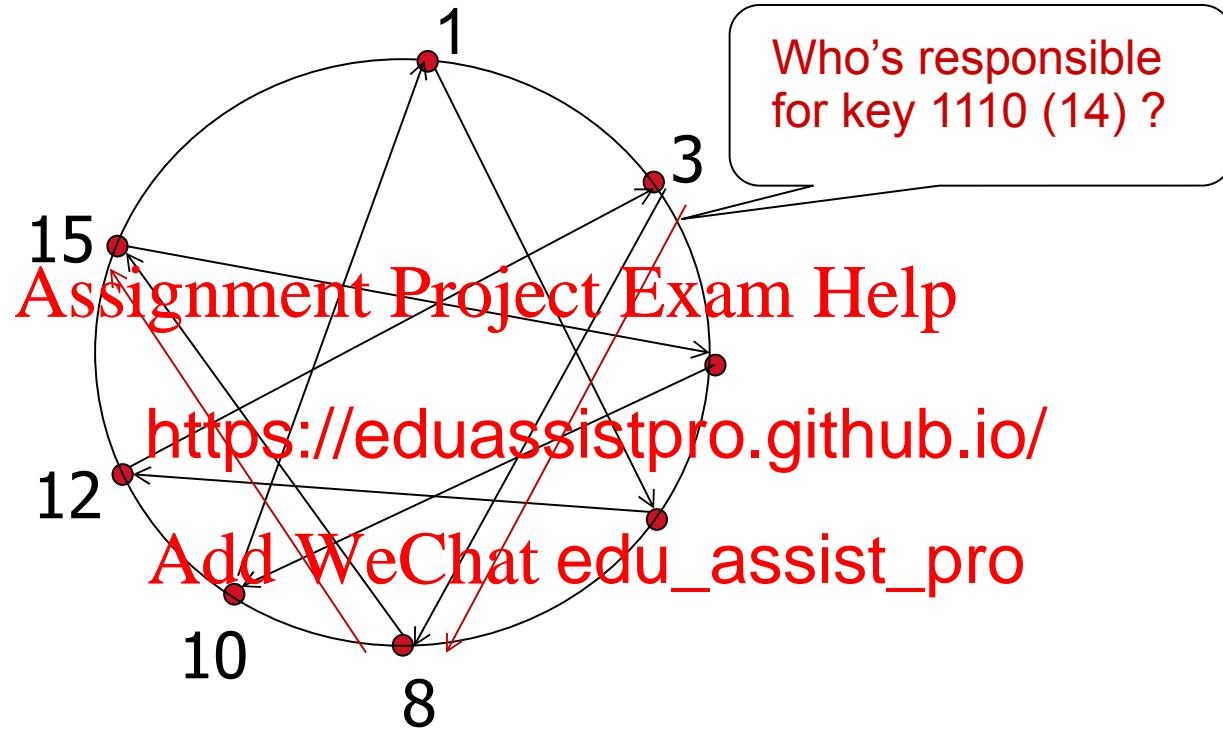
Define closest  
as closest  
successor



### Example: Chord is an example of a Distributed Hash Table (DHT)

As a node:

- › I have a successor peer **Assignment Project Exam Help**
- › I have a predecessor p
- › I have some shortcuts **<https://eduassistpro.github.io/>** to speedup delivery of requests **Add WeChat edu\_assist\_pro**
  
- › Chord: A scalable peer-to-peer lookup service for internet applications. Stoica et al. SIGCOMM 2001.



- › each peer keeps track of predecessor, successor, short cuts.
- › reduced from 6 to 2 messages.
- › possible to design shortcuts so  $O(\log N)$  neighbors,  $O(\log N)$  messages in query



# Assignment Project Exam Help **Transport Layer**

<https://eduassistpro.github.io/>

Add WeChat **edu\_assist\_pro**

## our goals:

- › understand principles behind transport layer services:
  - multiplexing, demultiplexing
  - reliable data transfer
  - flow control
  - congestion control
- › learn about Internet transport layer protocols:
  - <https://eduassistpro.github.io/>
  - Add WeChat
  - Transport layer
  - TCP congestion control

- › Transport-layer services

Assignment Project Exam Help

- › Multiplexing/demultiplexing

<https://eduassistpro.github.io/>

- › Connectionless transport (UDP)

Add WeChat edu\_assist\_pro

- › Principles of reliable data transfer

- › TCP protocol



# Assignment Project Exam Help **Transport Services**

<https://eduassistpro.github.io/>

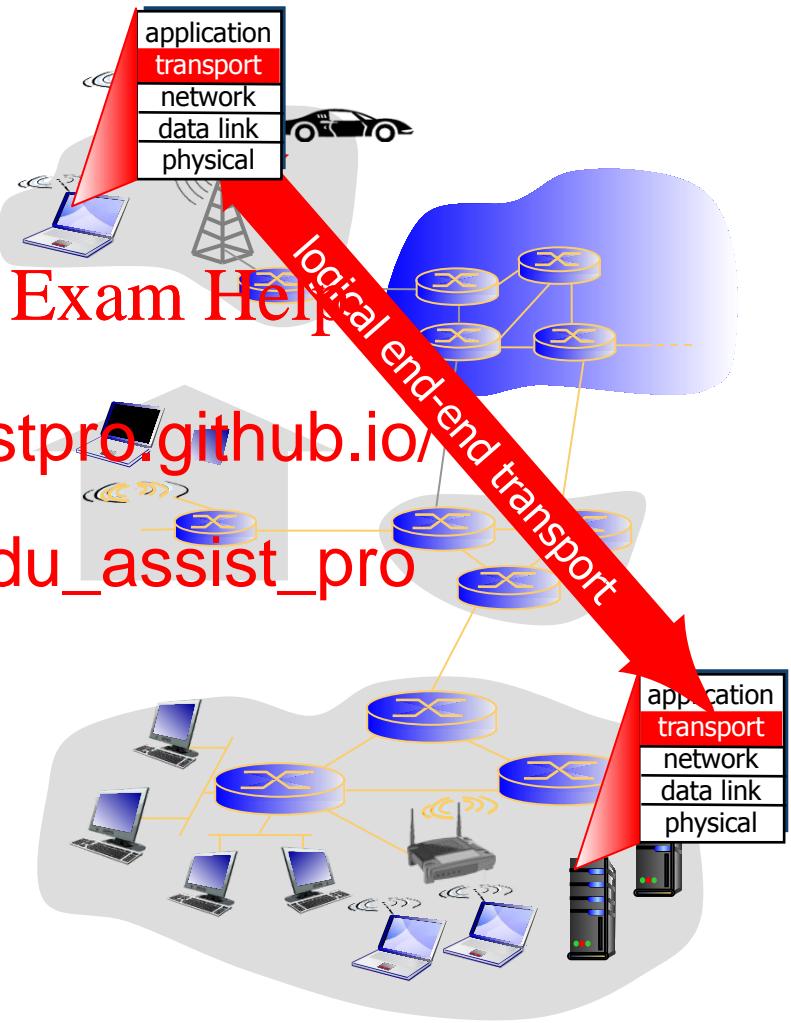
Add WeChat **edu\_assist\_pro**

- › provide *logical communication* between app processes running on different hosts
- › transport protocols run in end systems
  - send side: breaks app into *segments*, passes to network layer
  - rcv side: reassembles segments into messages, passes to app layer
- › more than one transport protocol available to apps
  - Internet: TCP and UDP

**Assignment Project Exam Help**

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro



- › **network layer:** host-to-host communication
  - best-effort, unreliable

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

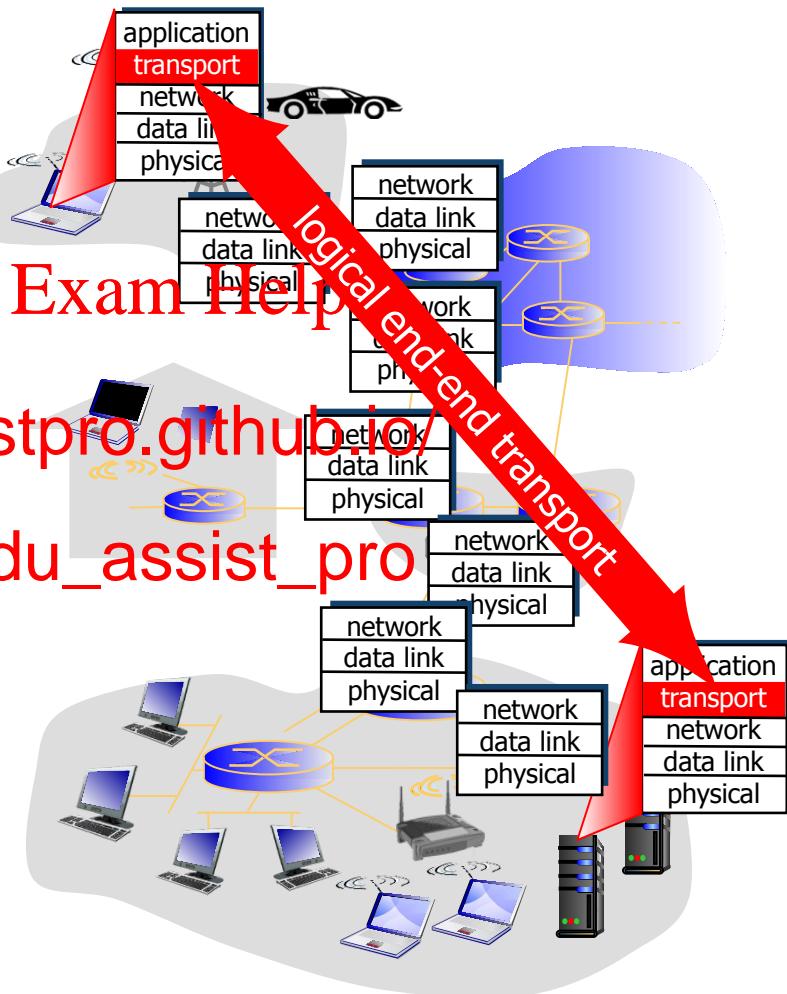
- › **transport layer:** process-to-process communication
  - relies on, enhances, network layer services

- › IP: best effort service
- › reliable, in-order delivery (TCP)
  - congestion control
  - flow control
  - connection setup
- › unreliable, unordered delivery: UDP
  - no-frills extension of “best-effort” IP
- › services not available:
  - delay guarantees
  - bandwidth guarantees

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro





# Assignment Project Exam Help **Transport Services**

<https://eduassistpro.github.io/>

Add WeChat **edu\_assist\_pro**

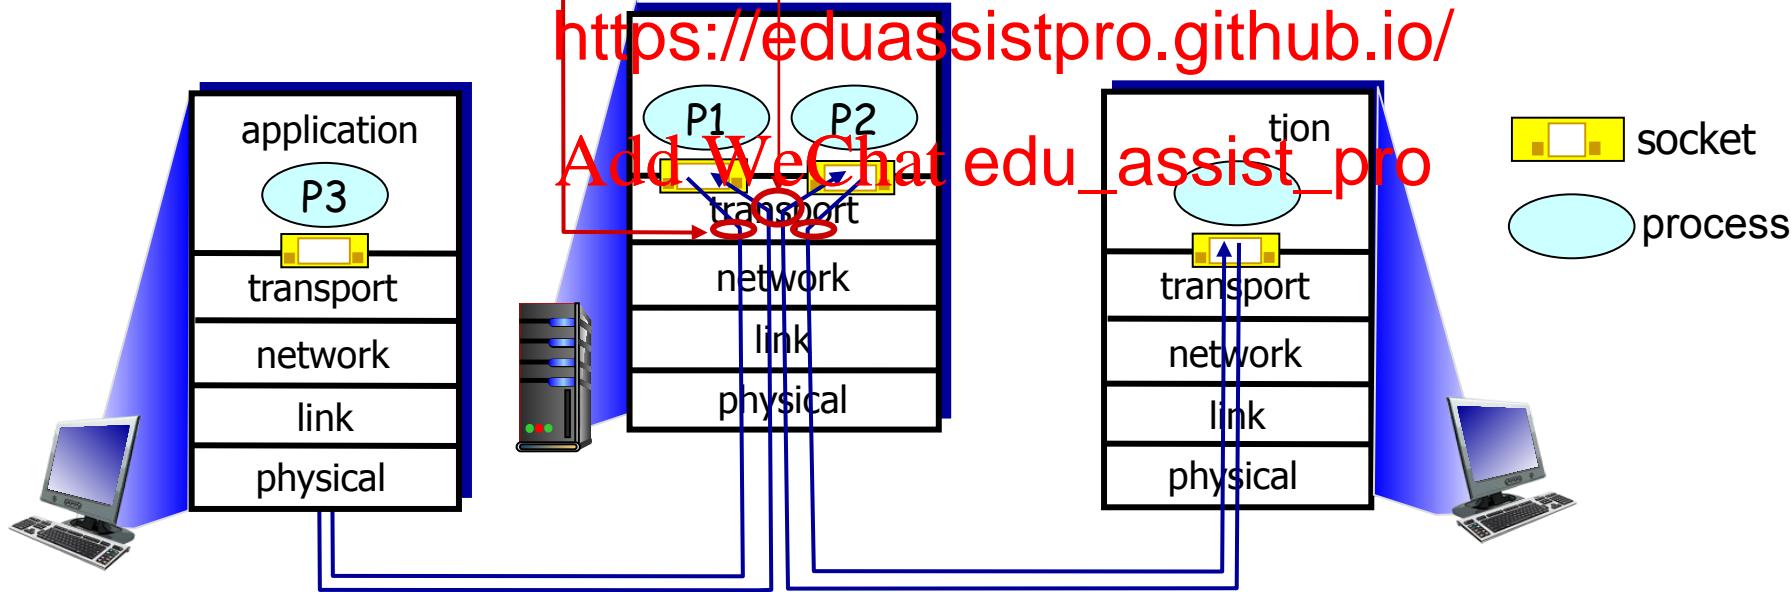
*- multiplexing at sender:*

handle data from multiple sockets, add transport header (later used for demultiplexing)

*- demultiplexing at receiver:*

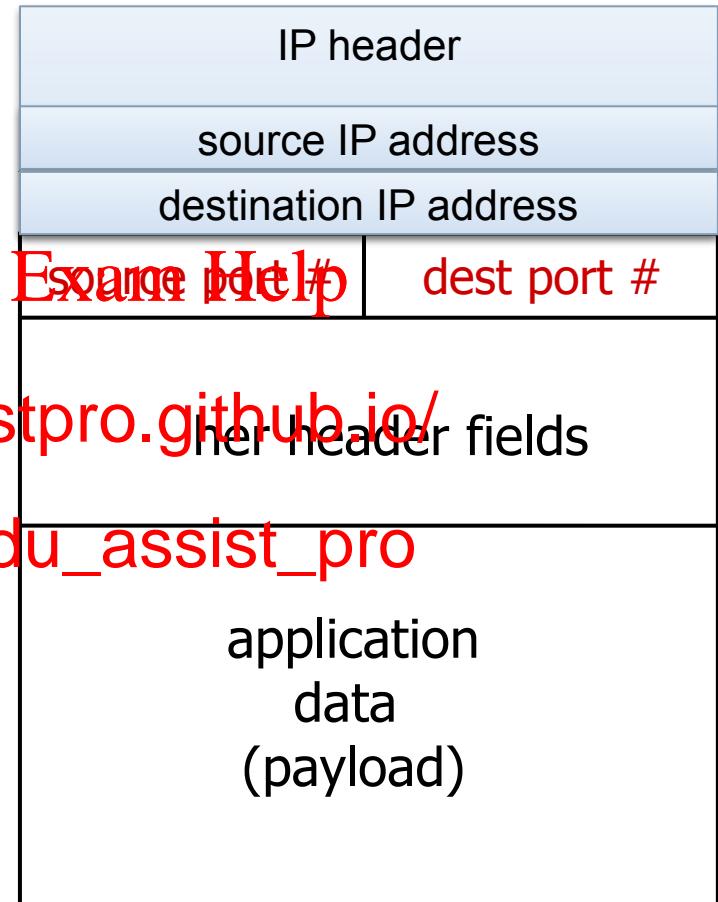
use header info to deliver received segments to correct socket

Assignment Project Exam Help



# How demultiplexing works

- › host receives IP datagrams
  - each datagram has source IP address, destination IP address
  - each datagram carries one transport layer segment
  - each segment has source, number
- › host uses IP addresses & port numbers to direct segment to appropriate socket



TCP/UDP segment format

› Receiver

- › *recall:* created socket has host-local port #:

› Sender

- ❖ *recall:* when creating datagram to send into UDP socket, must specify
  - destination IP address

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat `edu_assist_pro` with same dest. IP

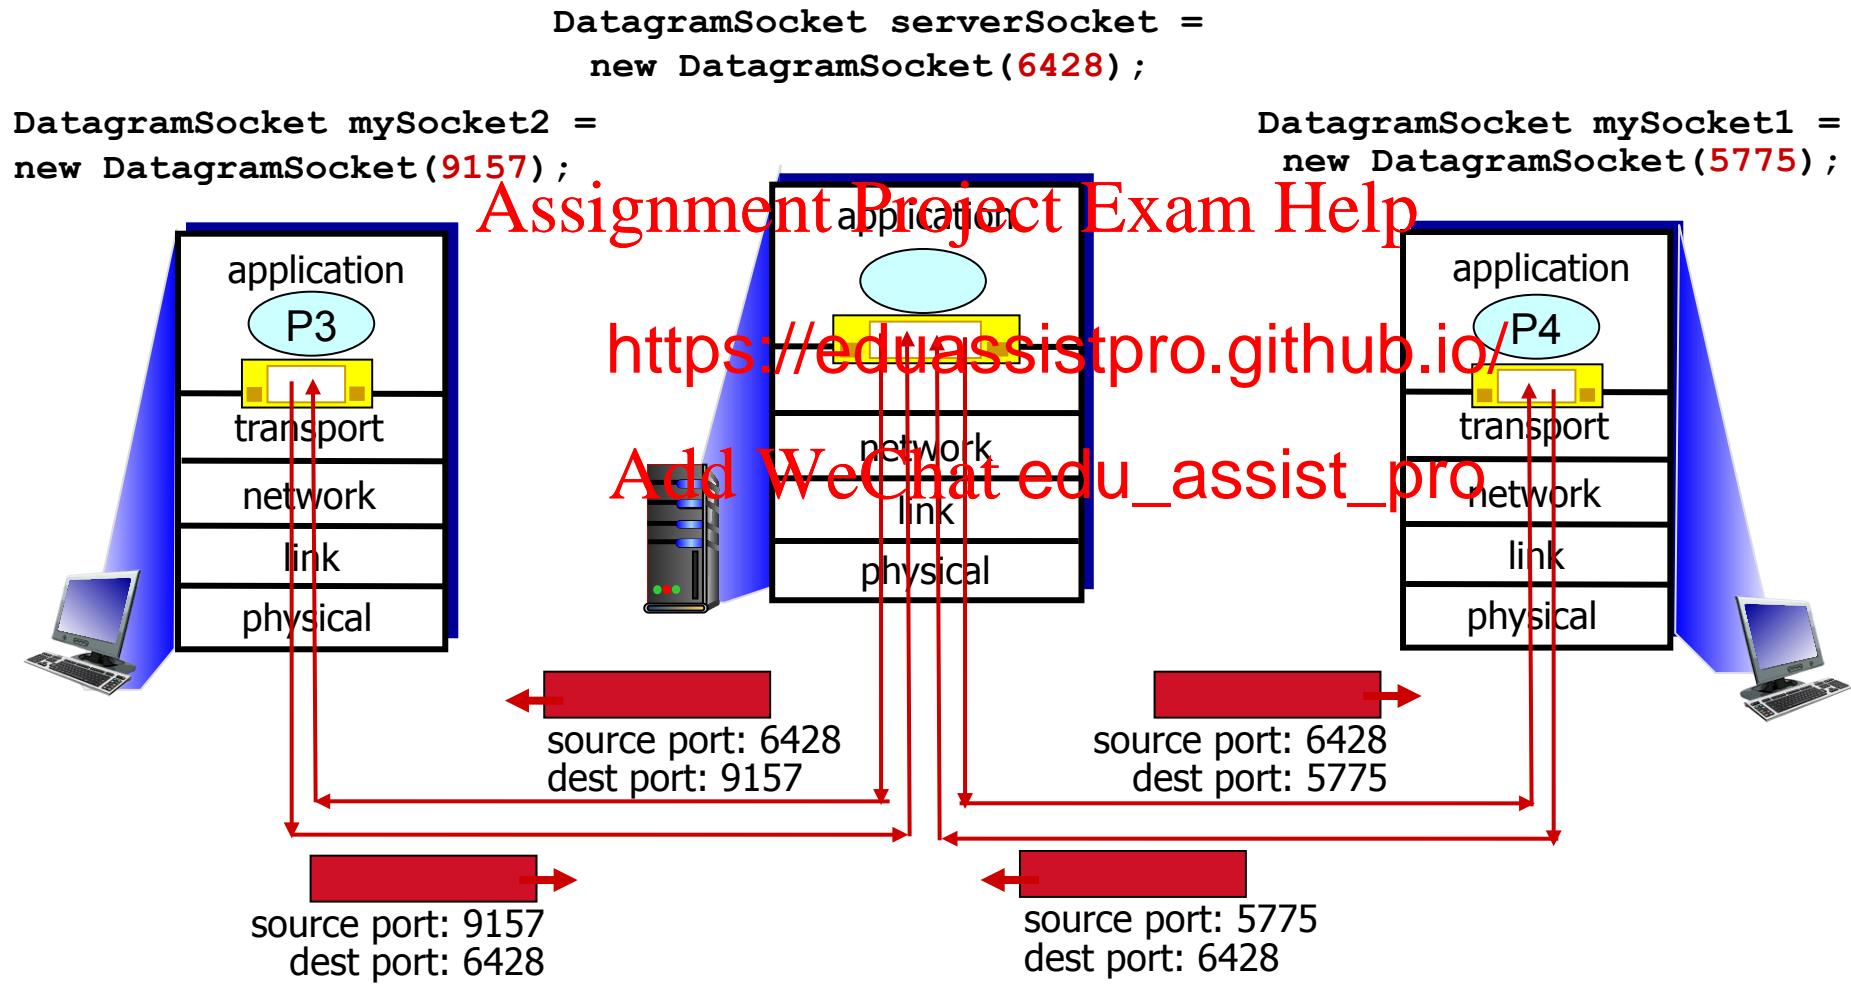
- › when host receives UDP segment:

- Checks destination port # in segment
- directs UDP segment to socket with that port #



address, dest. port #, but different source IP addresses and/or source port numbers will be directed to *same socket* at dest

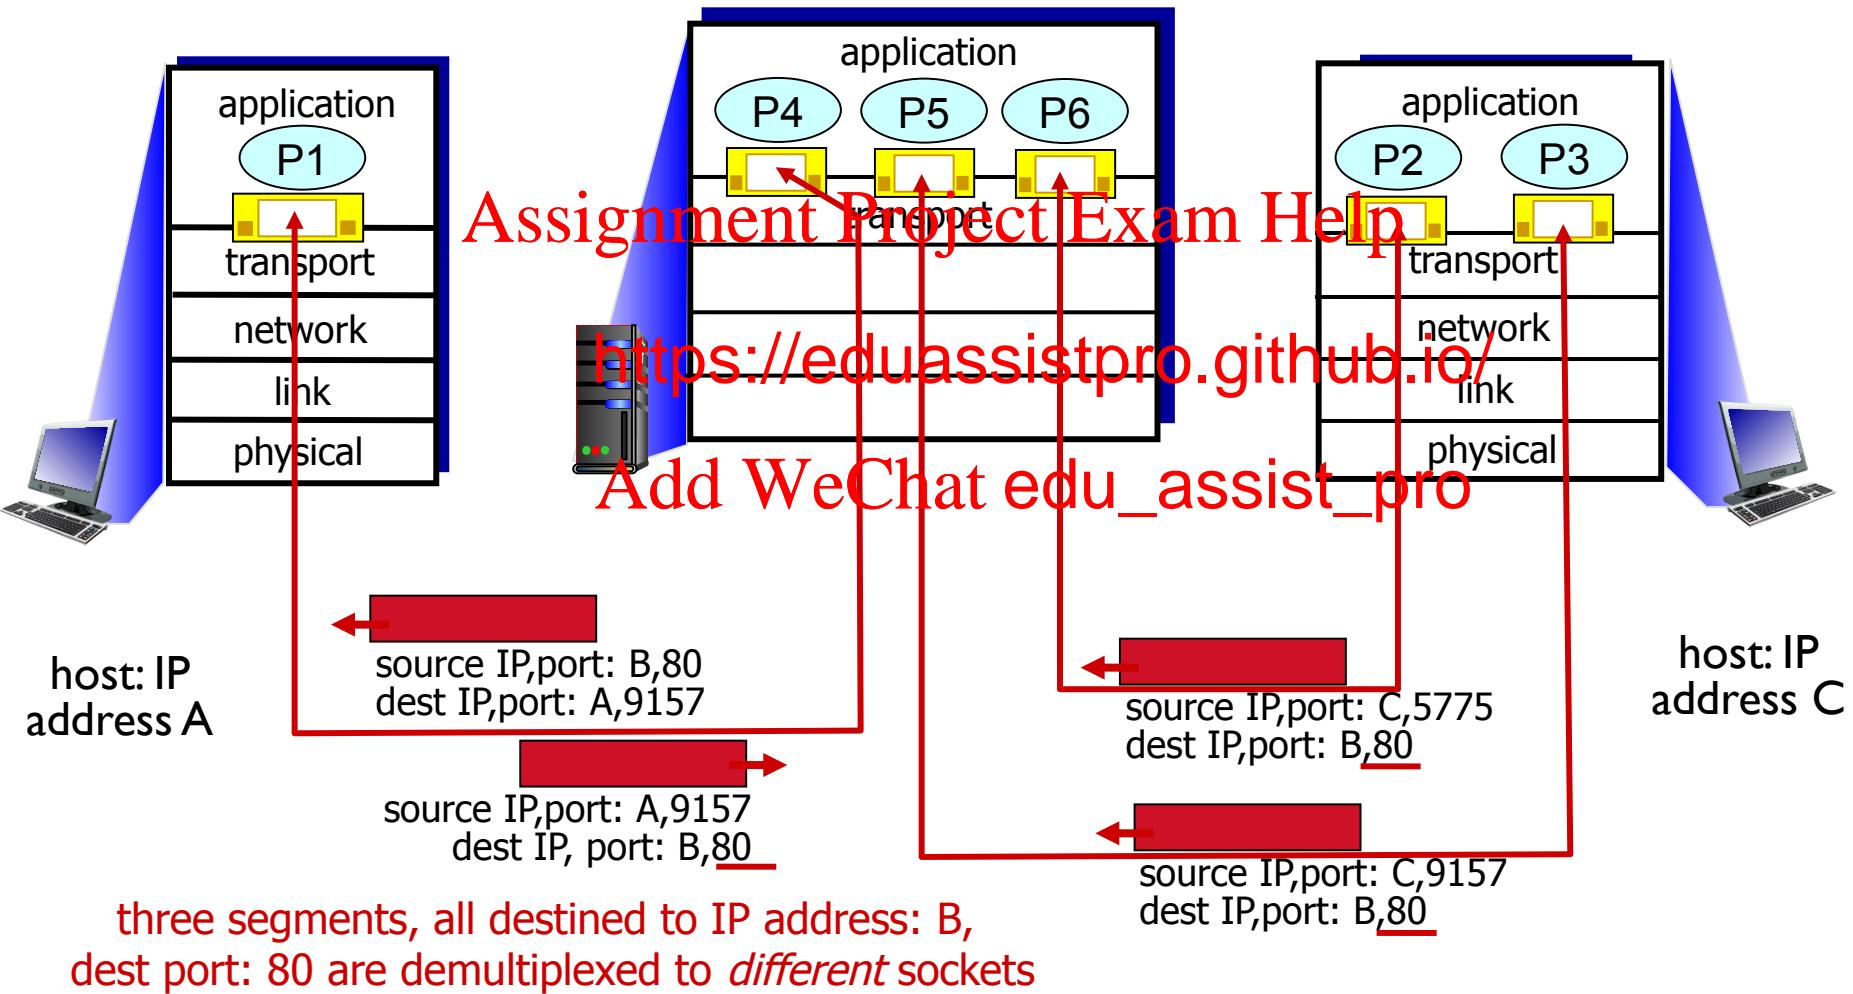
# Connectionless demux: example



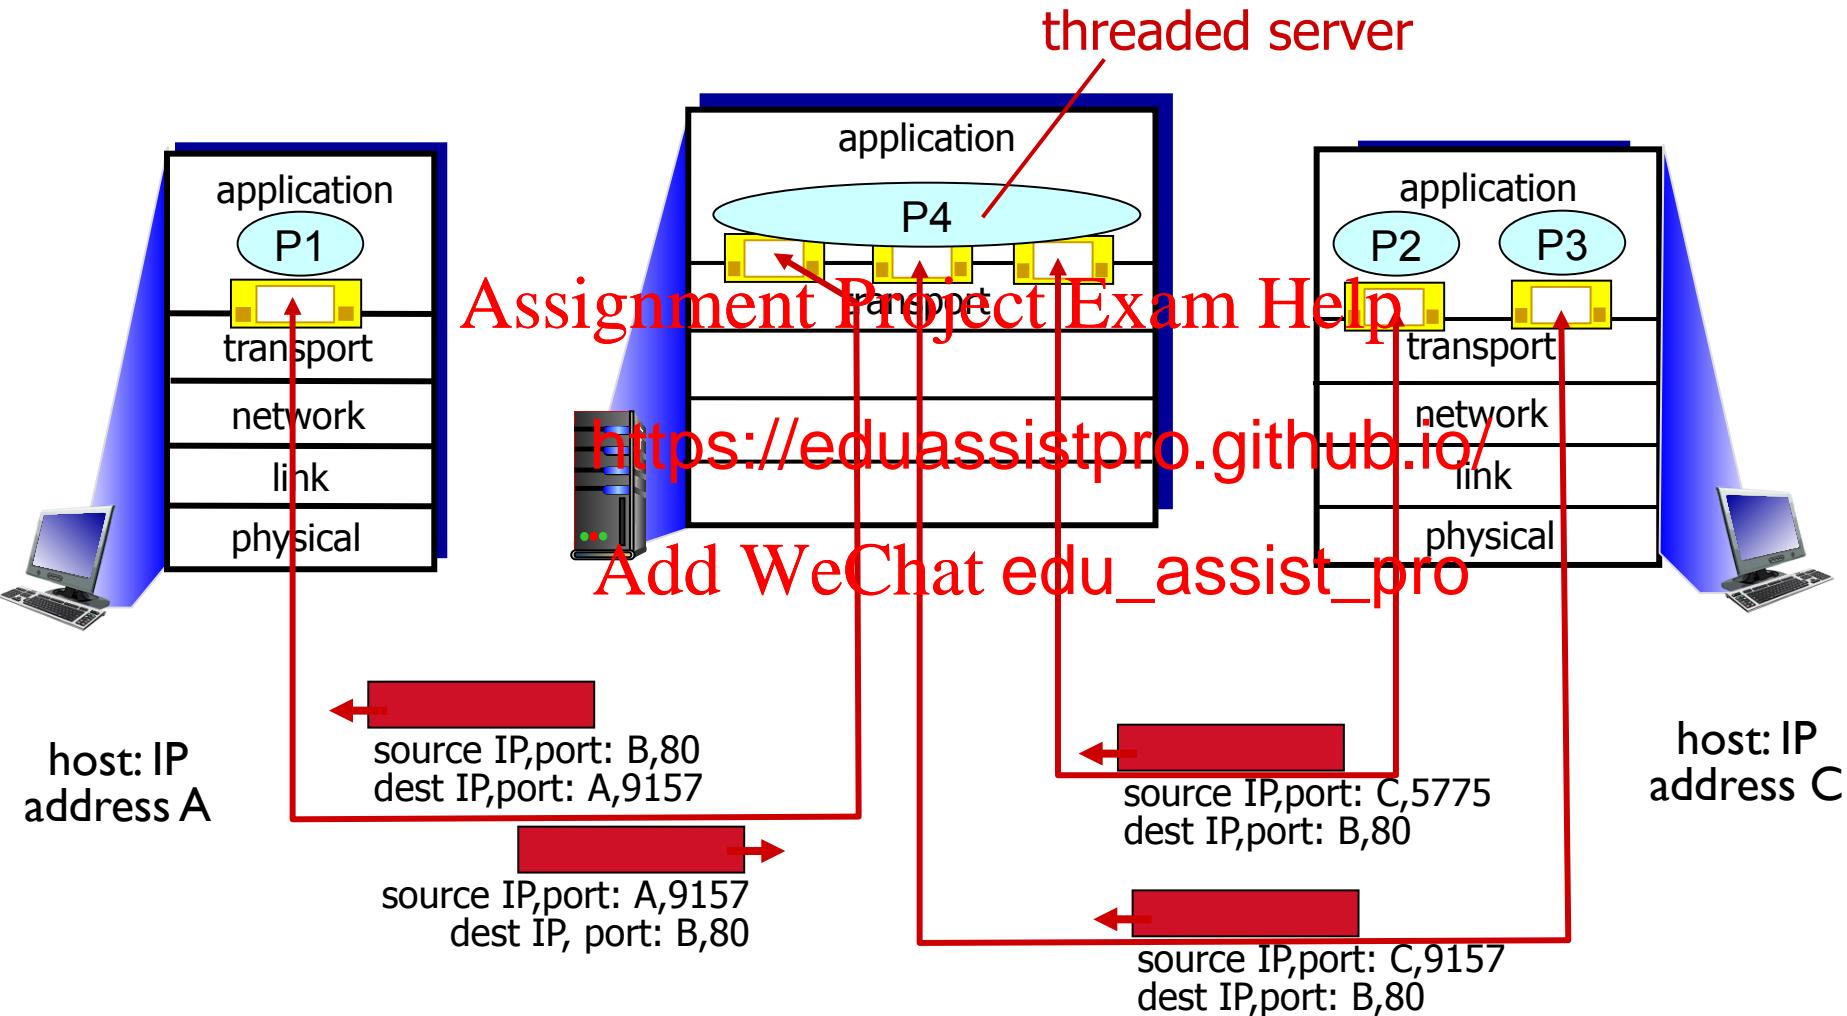


- › TCP socket identified by 4-tuple:
  - source IP address
  - source port number
  - dest IP address
  - dest port number
- › demux: receiver uses all four values to direct segment to appropriate socket
  - › server host may support many simultaneous TCP sockets:
    - et identified by its socket, ie
    - Add WeChat `edu_assist_pro` have different sockets for each connecting client
    - non-persistent HTTP will have different socket for each request

# Connection-oriented demux: example



# Connection-oriented demux: example





# Connectionless Transport UDP

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# UDP: User Datagram Protocol [RFC 768]

› “no frills,” Internet transport protocol

› “best effort” service, UDP segments may be:

- lost
- delivered out-of-order

› **connectionless:**

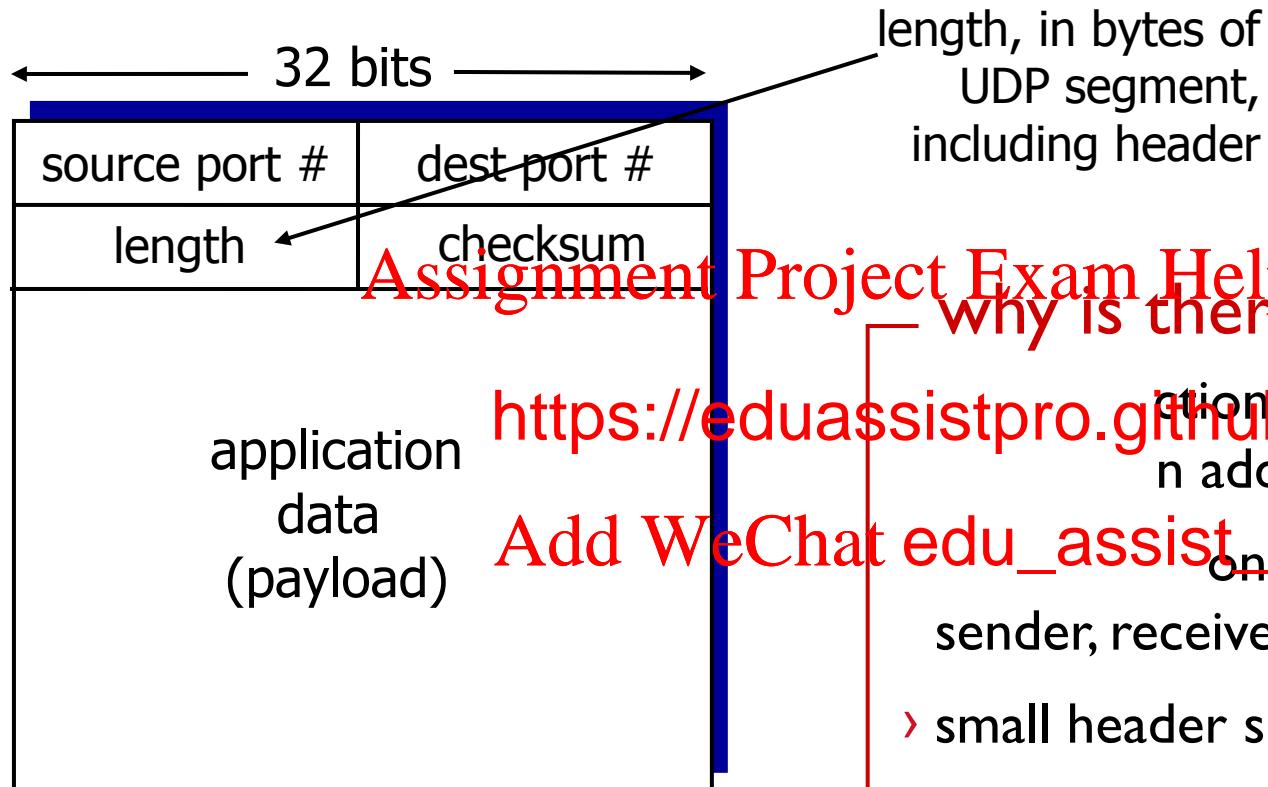
- no handshaking between UDP sender, receiver
- each UDP segment handled independently of others

❖ UDP use:

- streaming multimedia apps (loss tolerant, rate sensitive)
- DNS

<https://eduassistpro.github.io/>

❖ Add WeChat [edu\\_assist\\_pro](https://edu_assist_pro) at application layer  
application-specific error recovery!



UDP segment format

length, in bytes of  
UDP segment,  
including header

Assignment Project Exam Help

why is there a UDP?

connection establishment  
in add d

<https://eduassistpro.github.io/>

Add WeChat **edu\_assist\_pro**

connection state at

sender, receiver

› small header size

› no congestion control: UDP can blast away as fast as desired

**Goal:** detect “errors” (e.g., flipped bits) in transmitted segment

sender: Assignment Project Exam Help

- › treat segment contents including header fields as sequence of 16-bit words
  - › sum: addition (one's complement sum) of segment contents
  - › checksum: complement of sum
  - › sender puts checksum value into UDP checksum field
- receiver: https://eduassistpro.github.io/
- checksum of received segment
- Add WeChat edu\_assist\_pro
- computed checksum field value:
- NO - error detected
  - YES - no error detected.

example: add two 16-bit integers

	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
	A	s	i	g	o	n	m	P	r	o	j	E	x	a	m	0
<hr/>																
carryout	1	0	https://eduassistpro.github.io/0	1	1	1	1	0	1	1	1	0	1	1	1	1
wraparound	0	0							0	0	0	0	0	0	1	
<hr/>																
sum	1	0	1	1	1	0	1			1	1	1	0	0		
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	1	1	

Note: when adding numbers, a carryout from the most significant bit needs to be added to the result



# Principles of Reliable Data Transfer

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

- › important in application, transport, link layers
  - top-10 list of important networking topics!

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

- › characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

- › important in application, transport, link layers
  - top-10 list of important networking topics!

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

- › characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

- › important in application, transport, link layers
  - top-10 list of important networking topics!

## Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat `edu_assist_pro`

Network  
layer

- › characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

**rdt\_send()** : called from above,  
(e.g., by app.). Passed data to  
deliver to receiver upper layer

**deliver\_data()** : called by  
**rdt** to deliver data to upper

## Assignment Project Exam Help

send  
side

<https://eduassistpro.github.io/> receive  
side

Add WeChat **edu\_assist\_pro**

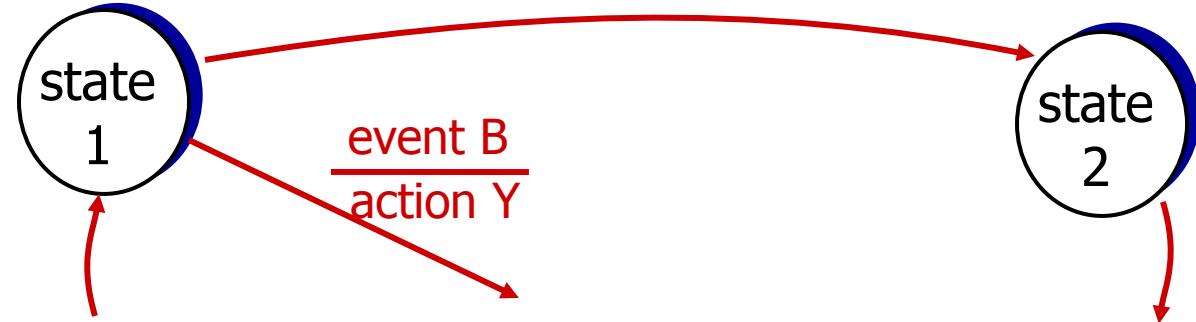
**udt\_send()** : called by rdt,  
to transfer packet over  
unreliable channel to receiver

**rdt\_rcv()** : called when packet  
arrives on rcv-side of channel

We will:

- › incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- › consider only u
  - but control info <https://eduassistpro.github.io/>
- › use finite state machines (FSM) ~~Add WeChat edu\_assist\_pro~~, ~~sender, receiver~~

**state:** when in this "state", next state and action uniquely determined by next event



## › underlying channel perfectly reliable

- no bit errors

- no loss of packets

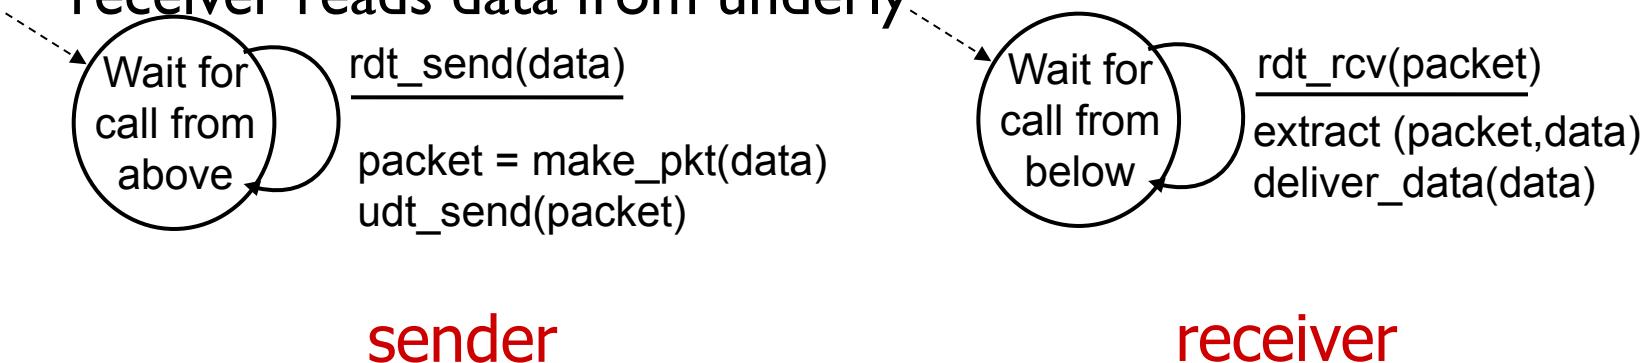
Assignment Project Exam Help

## › separate FSMs for

<https://eduassistpro.github.io/>

- sender sends data

- receiver reads data from underlying



- › underlying channel may flip bits in packet
  - checksum to detect bit errors
- › the question: how to recover from errors:  
**Assignment Project Exam Help**

<https://eduassistpro.github.io/>

*How do humans handle “bit errors”  
during conversation?*

› underlying channel may flip bits in packet

- checksum to detect bit errors

› the question: how to recover from errors:

- acknowledgements (ACKs): receiver explicitly tells sender that pkt received OK

<https://eduassistpro.github.io/>

- negative acknowledgement: explicitly tell that pkt had errors

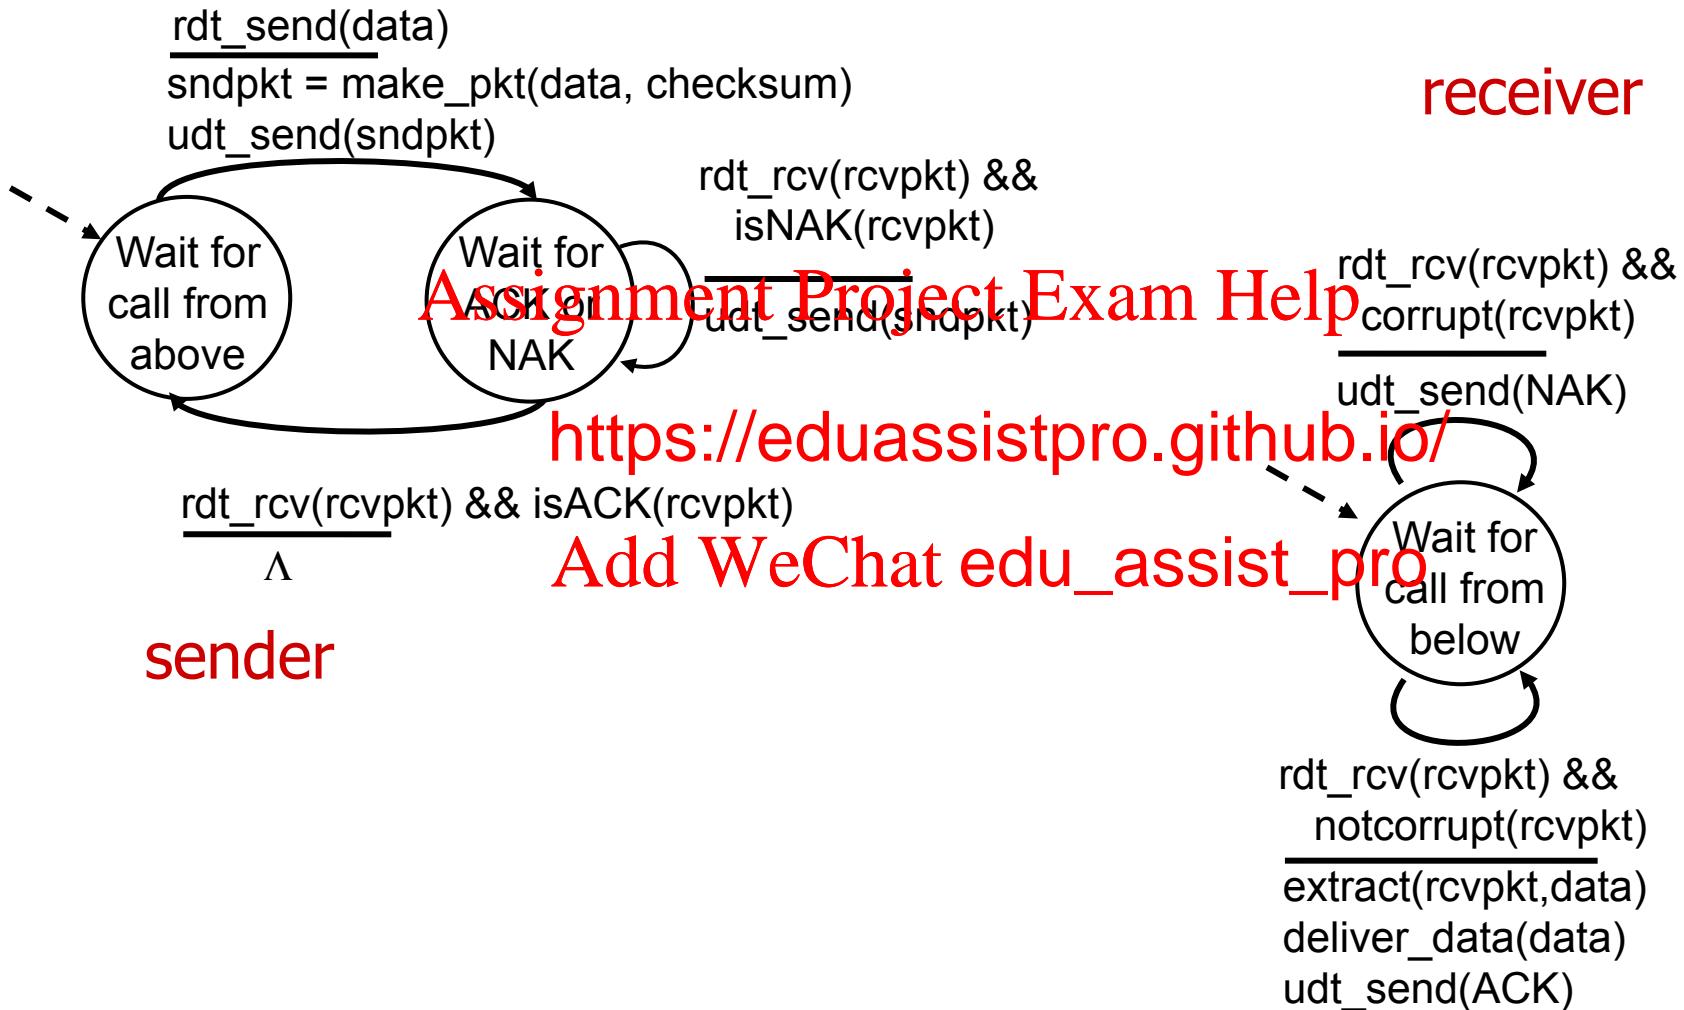
Add WeChat edu\_assist\_pro

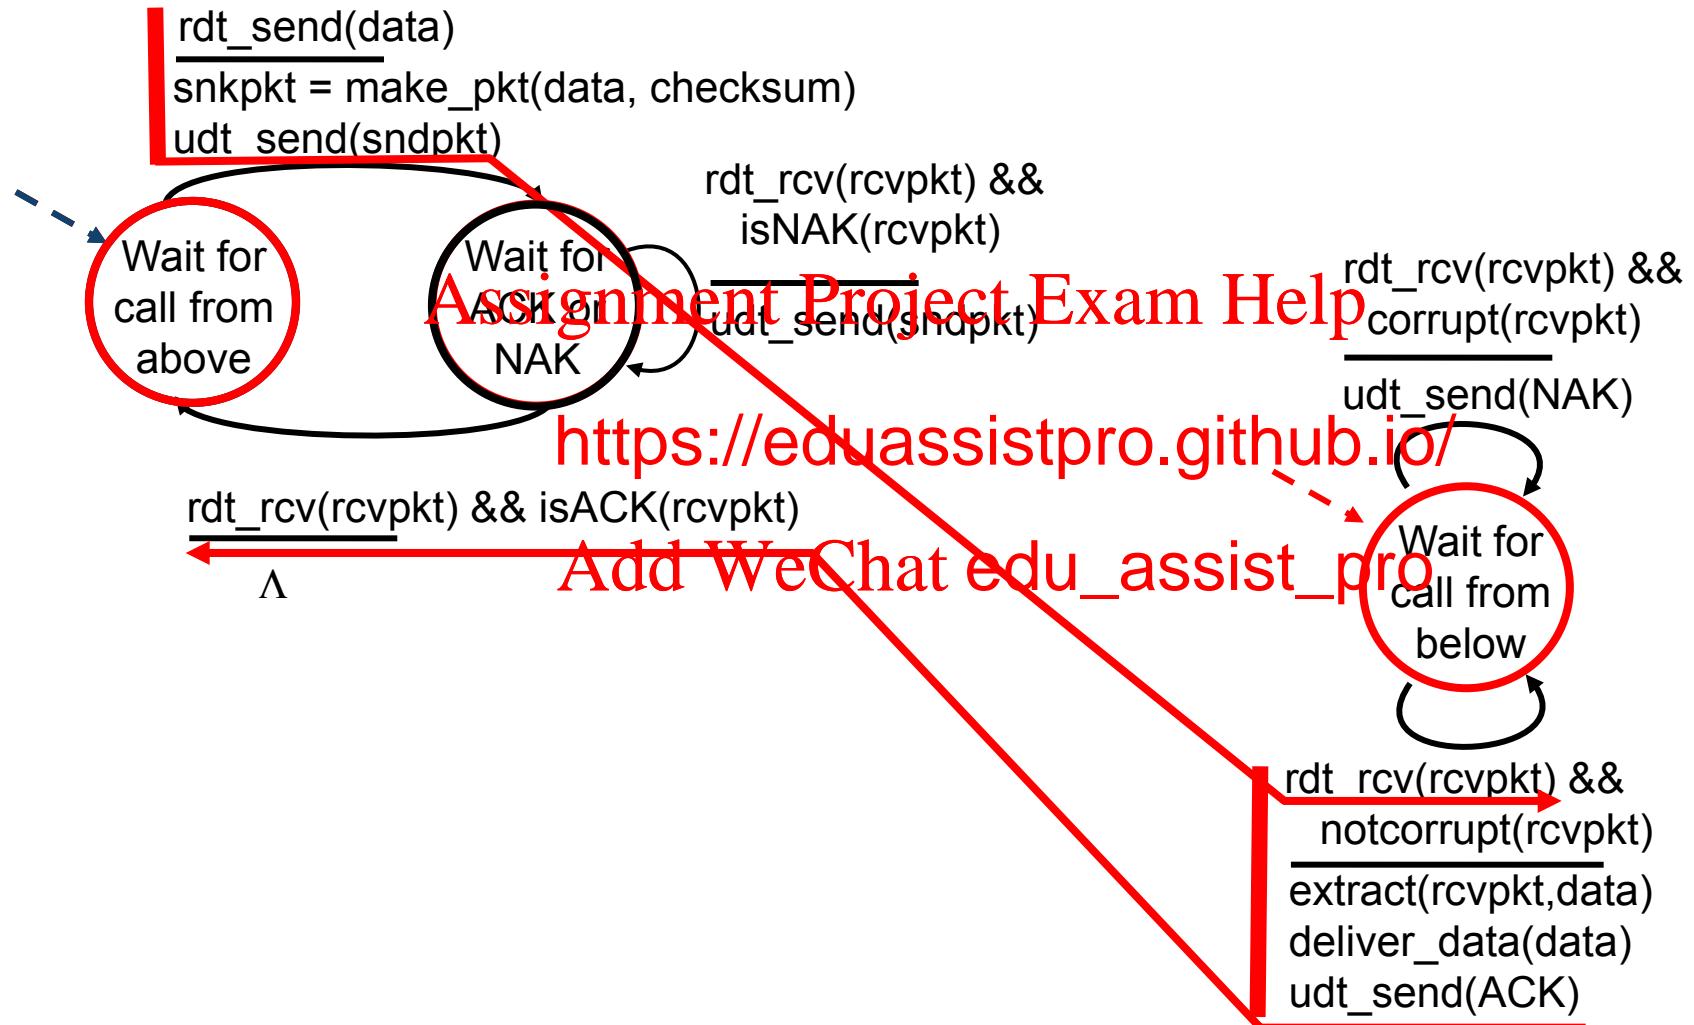
- sender retransmits pkt on receipt of NAK

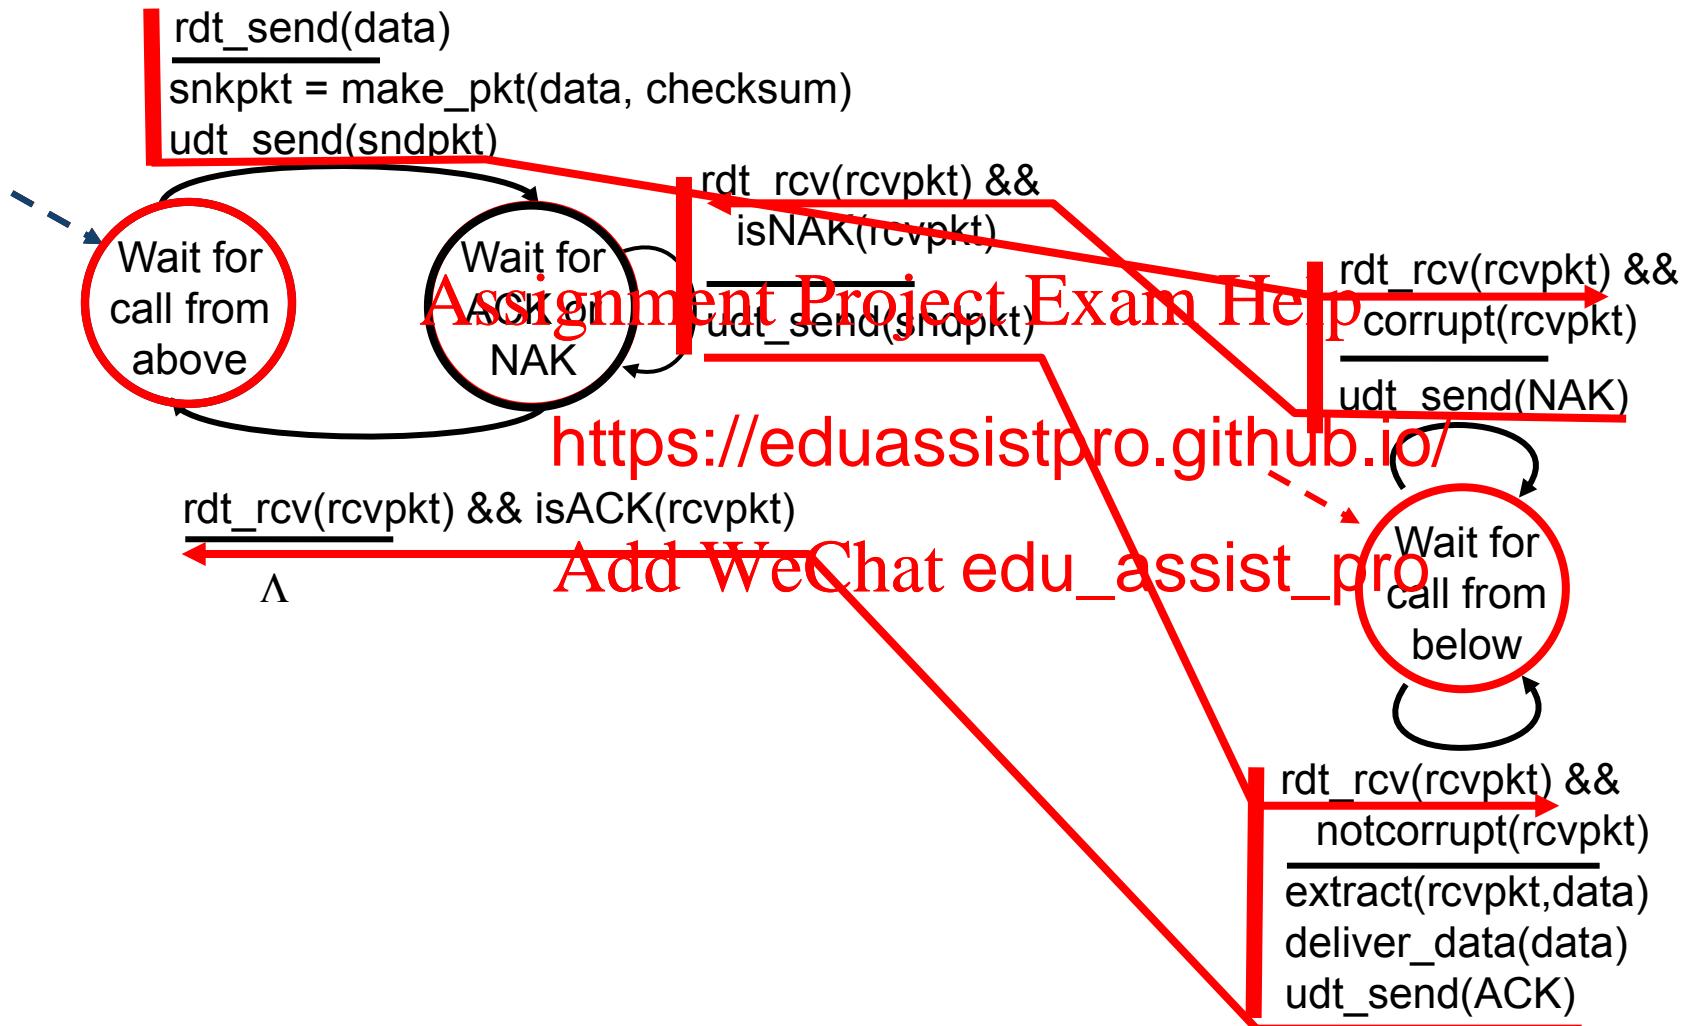
› new mechanisms in rdt2.0 (beyond rdt1.0):

- error detection

- feedback: control msgs (ACK,NAK) from receiver to sender







## what happens if ACK/NAK corrupted?

- › sender does not know what happened at receiver!
- › cannot just retransmit, possible duplicate

## handling duplicates:

- › sender retransmits current pkt if ACK/NAK corrupted

*sequence number* to

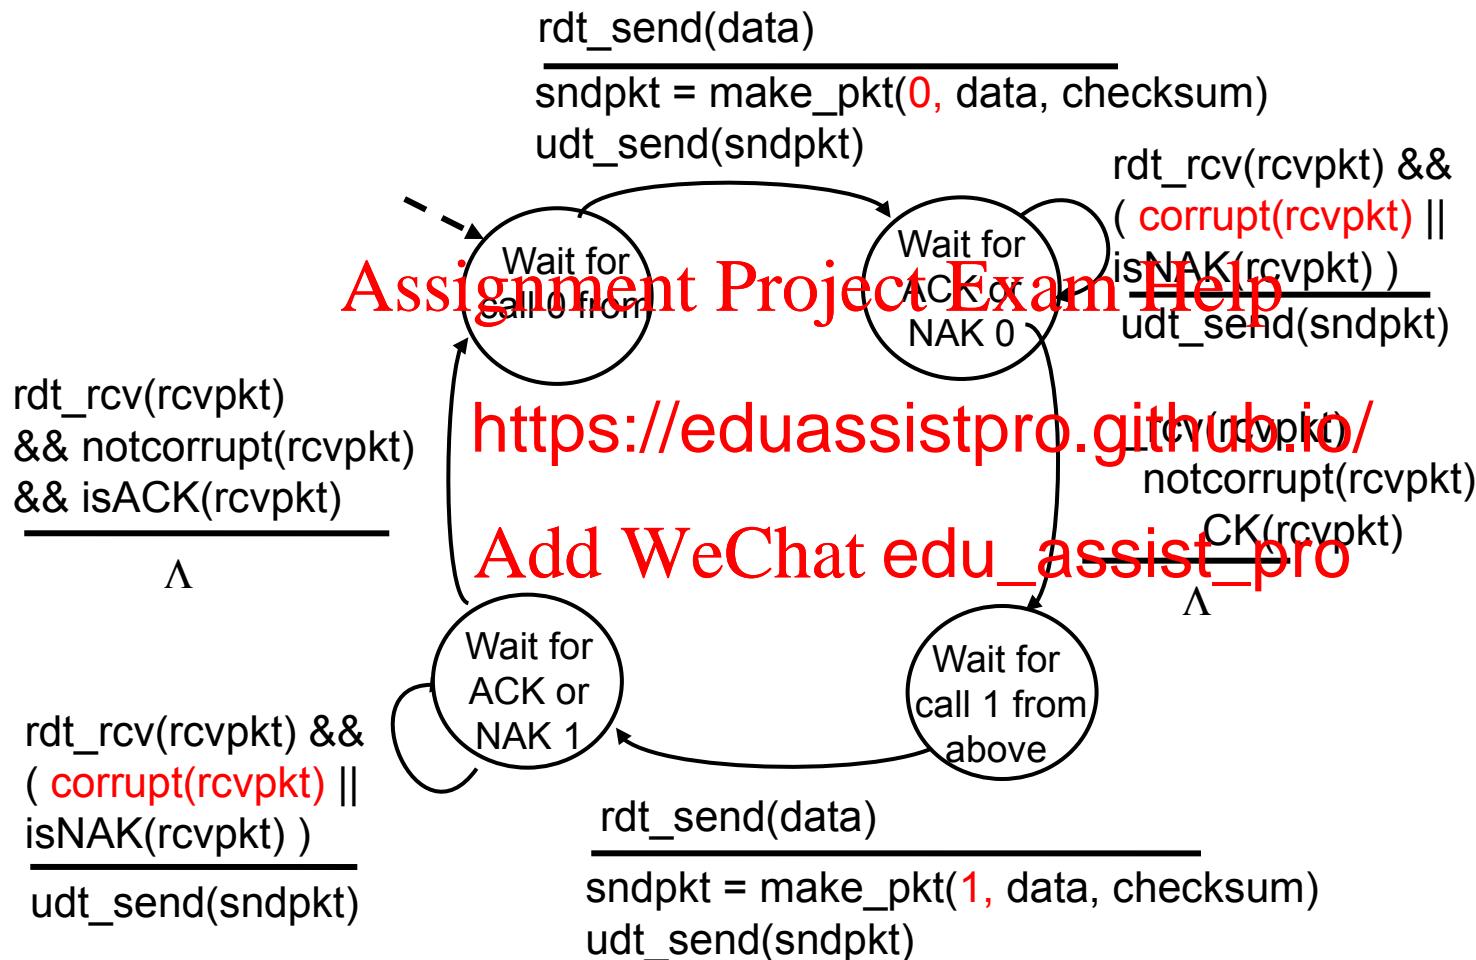
<https://eduassistpro.github.io/>

- › retransmit (does not deliver)

Add WeChat [edu\\_assist\\_pro](#)

stop and wait  
sender sends one packet,  
then waits for receiver  
response

# rdt2.1: sender, handles garbled ACK/NAKs



# rdt2.1: receiver, handles garbled ACK/NAKs

```
rdt_rcv(rcvpkt) &&
(corrupt(rcvpkt))
sndpkt = make_pkt(NAK, chks)
udt_send(sndpkt)
```

```
rdt_rcv(rcvpkt) &&
not corrupt(rcvpkt) &&
has_seq1(rcvpkt)
```

```
sndpkt = make_pkt(ACK, chks)
udt_send(sndpkt)
```

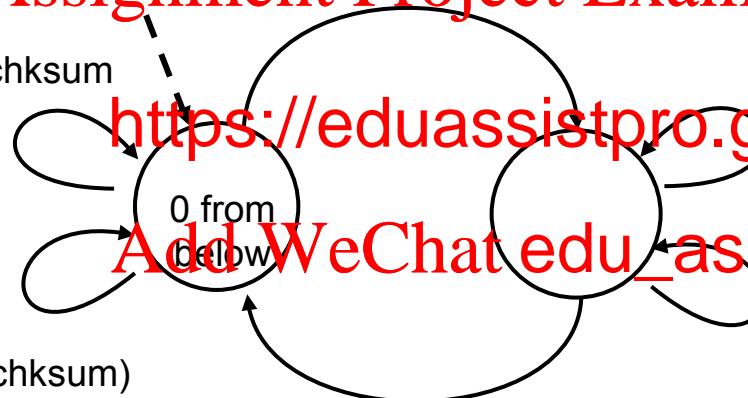
```
rdt_rcv(rcvpkt) && notcorrupt(rcvpkt)
&& has_seq0(rcvpkt)
```

```
extract(rcvpkt,data)
deliver_data(data)
sndpkt = make_pkt(ACK, chks)
udt_send(sndpkt)
```

```
rdt_rcv(rcvpkt) && notcorrupt(rcvpkt)
&& has_seq1(rcvpkt)
```

```
extract(rcvpkt,data)
deliver_data(data)
sndpkt = make_pkt(ACK, chks)
udt_send(sndpkt)
```

Assignment Project Exam Help



```
rat_rcv(rcvpkt) &&
(corrupt(rcvpkt))
sndpkt = make_pkt(NAK, chks)
udt_send(sndpkt)
```

```
rat_rcv(rcvpkt) &&
notcorrupt(rcvpkt) &&
has_seq0(rcvpkt)
```

```
sndpkt = make_pkt(ACK, chks)
udt_send(sndpkt)
```

sender:

- › seq # added to pkt
- › two seq. #'s (0,1)

suffice.

- › must check if received ACK/NAK corrupted

- › twice as many states

- state must “remember” whether “expected” pkt should have seq # of 0 or 1

receiver:

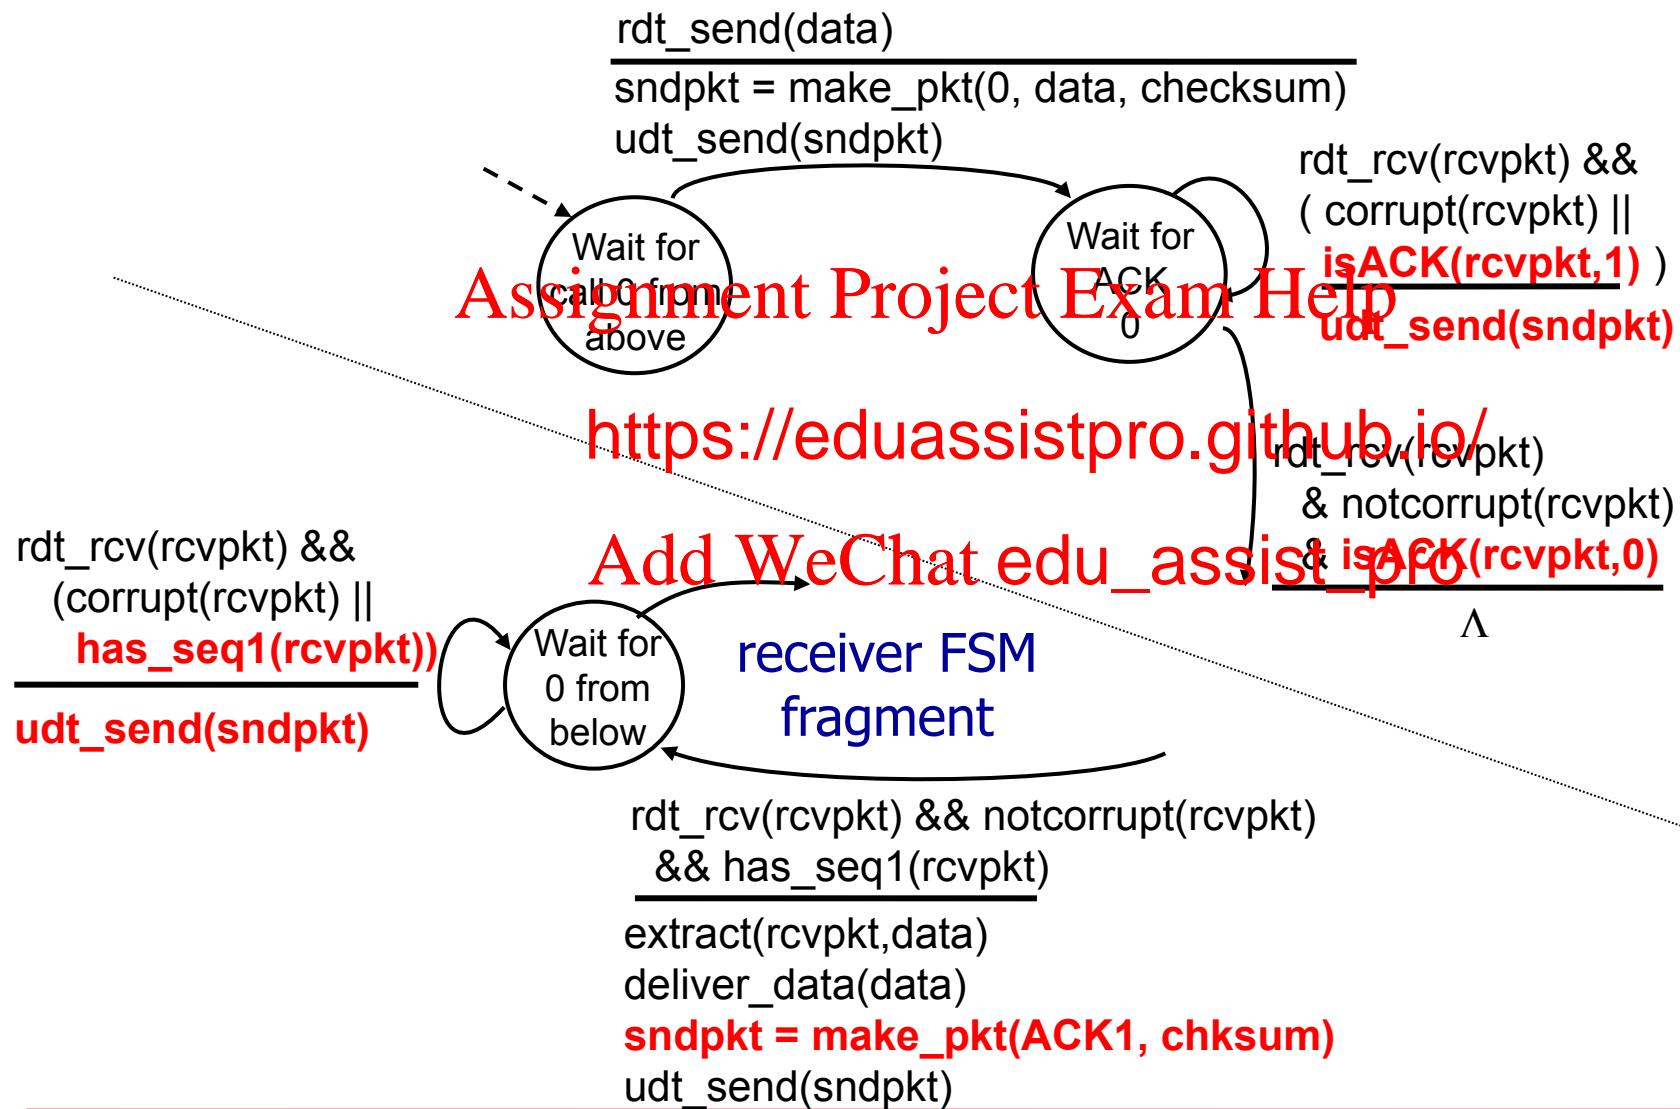
- › must check if received packet is duplicate

<https://eduassistpro.github.io/>

ates whether 0 or  
ed pkt seq

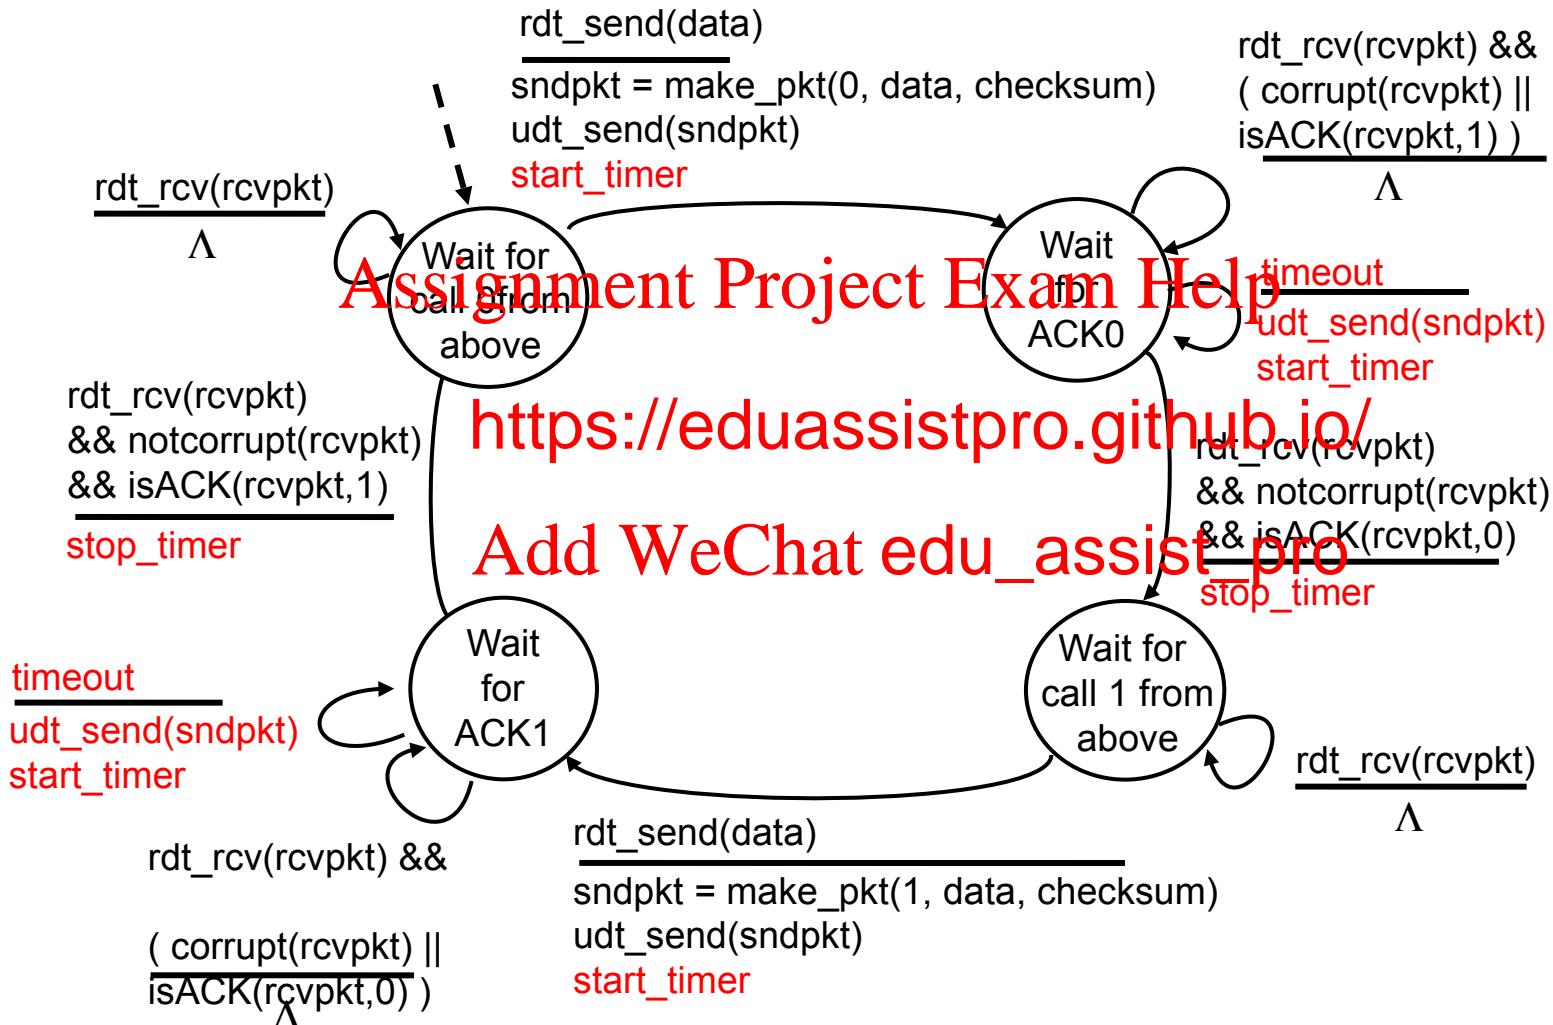
Add WeChat edu\_assist\_pro

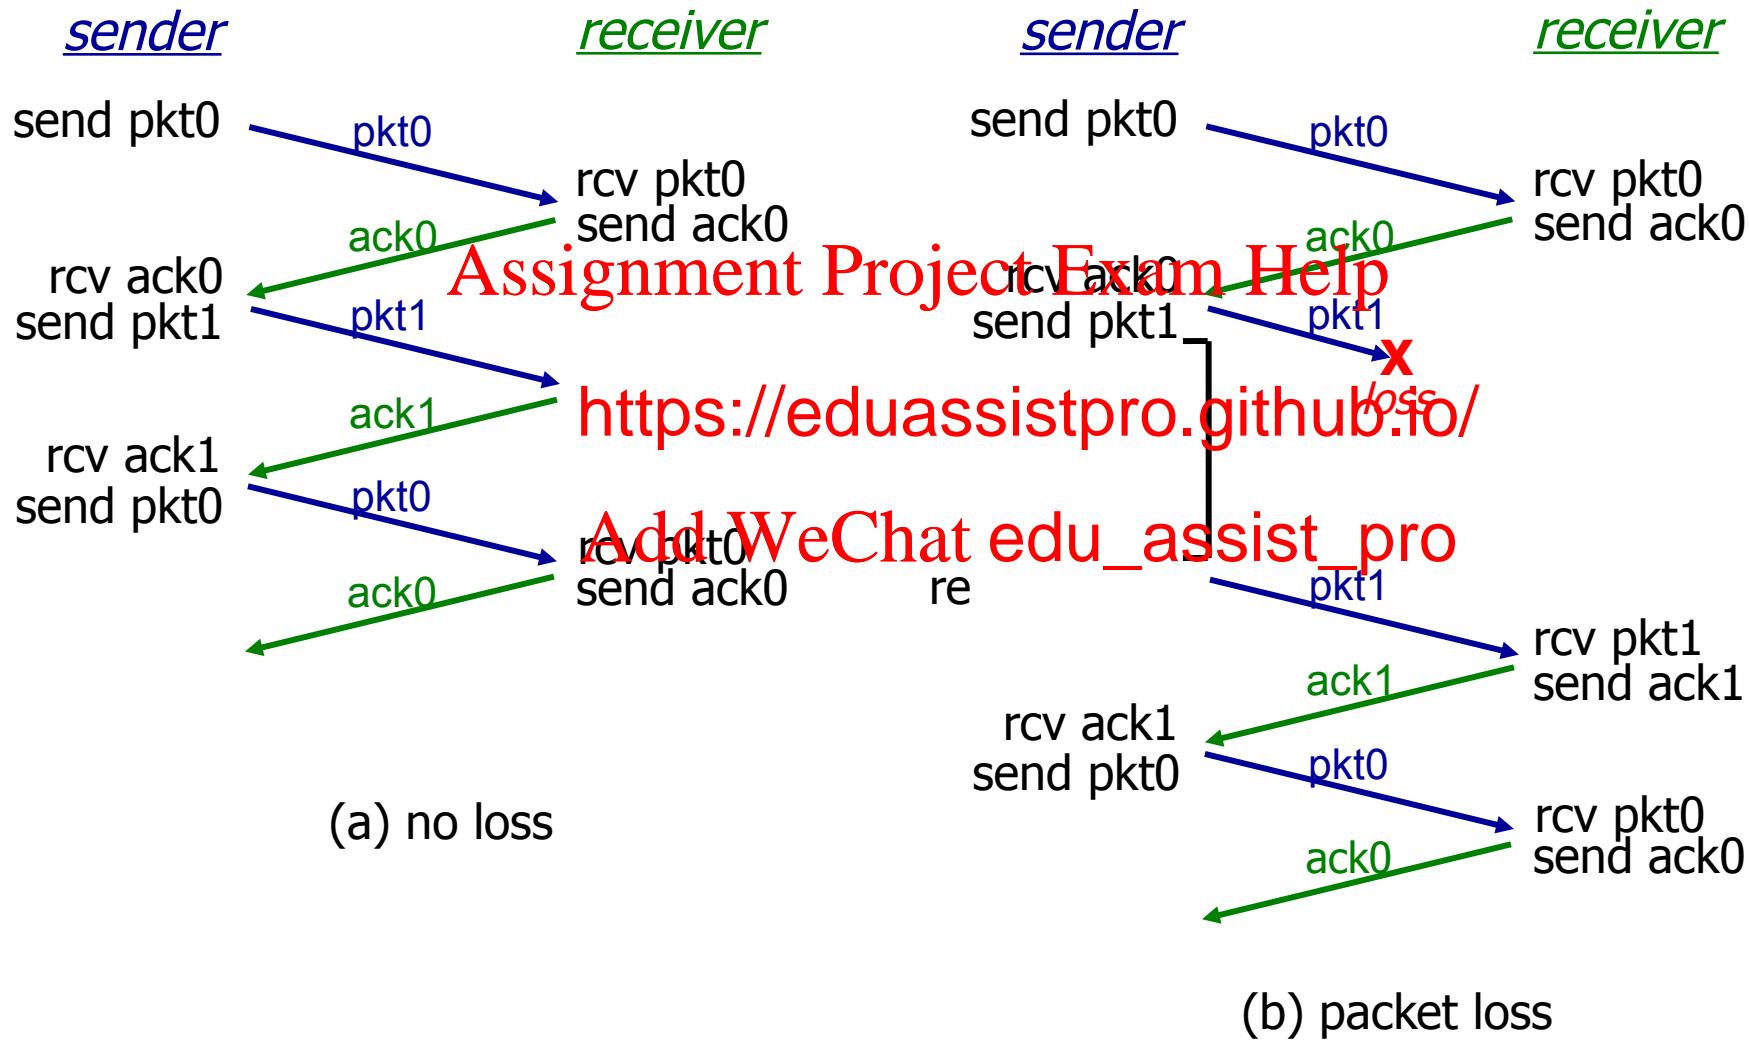
- › same functionality as rdt2.1, using ACKs only
- › instead of NAK, receiver sends ACK for last pkt received **OKAssignment Project Exam Help**
  - receiver must ex <https://eduassistpro.github.io/> being ACKed
- › "unexpected" A n same action as NAK: *retransmit current pkt* **Add WeChat edu\_assist\_pro**

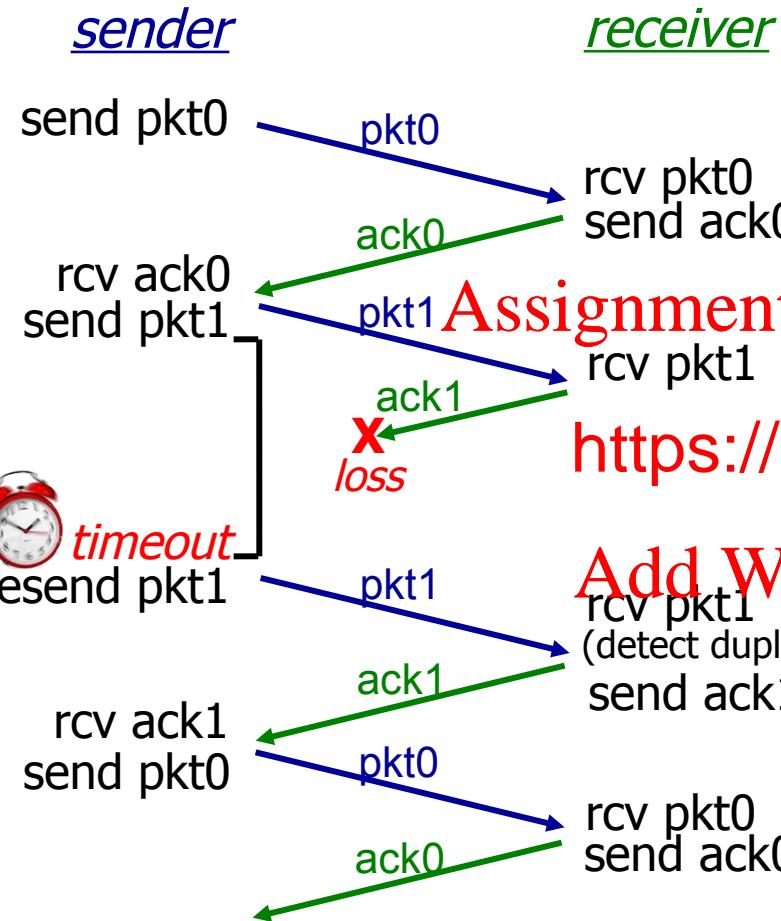


new assumption: underlying approach: sender waits  
channel can also lose “reasonable” amount of  
packets (data, ACKs) time for ACK

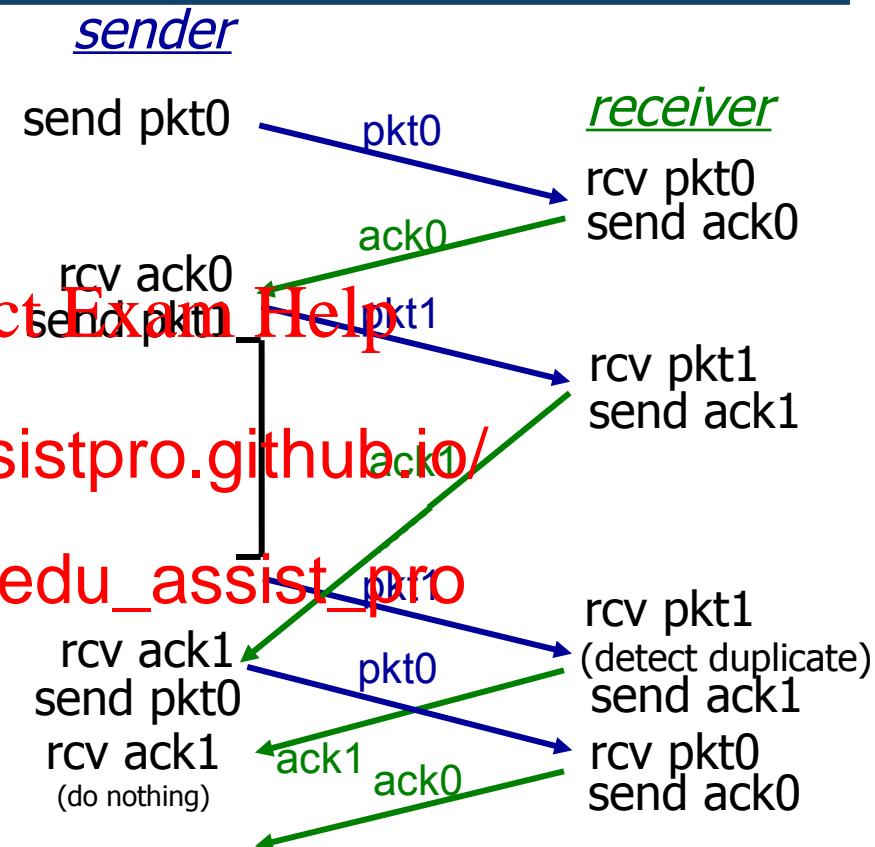
- Assignment Project Exam Help
- checksum, seq. #  
retransmissions  
help ... but not e
  - › if no ACK received in  
<https://eduassistpro.github.io/>  
K) just delayed (not  
lost)
  - Add WeChat `edu_assist_pro`  
n will be  
duplicate, but seq. #'s already  
handles this
    - receiver must specify seq # of  
pkt being ACKed
    - › requires countdown timer







(c) ACK loss



(d) premature timeout/ delayed ACK

- › rdt3.0 is correct, but performance stinks
- › e.g.: 1 Gbps link, 15 ms prop. delay, 8000 bit packet:

$$D_{trans} = \frac{L}{R} = \frac{8000 \text{ bits}}{8 \text{ microsecs}} = 1 \text{ msec}$$

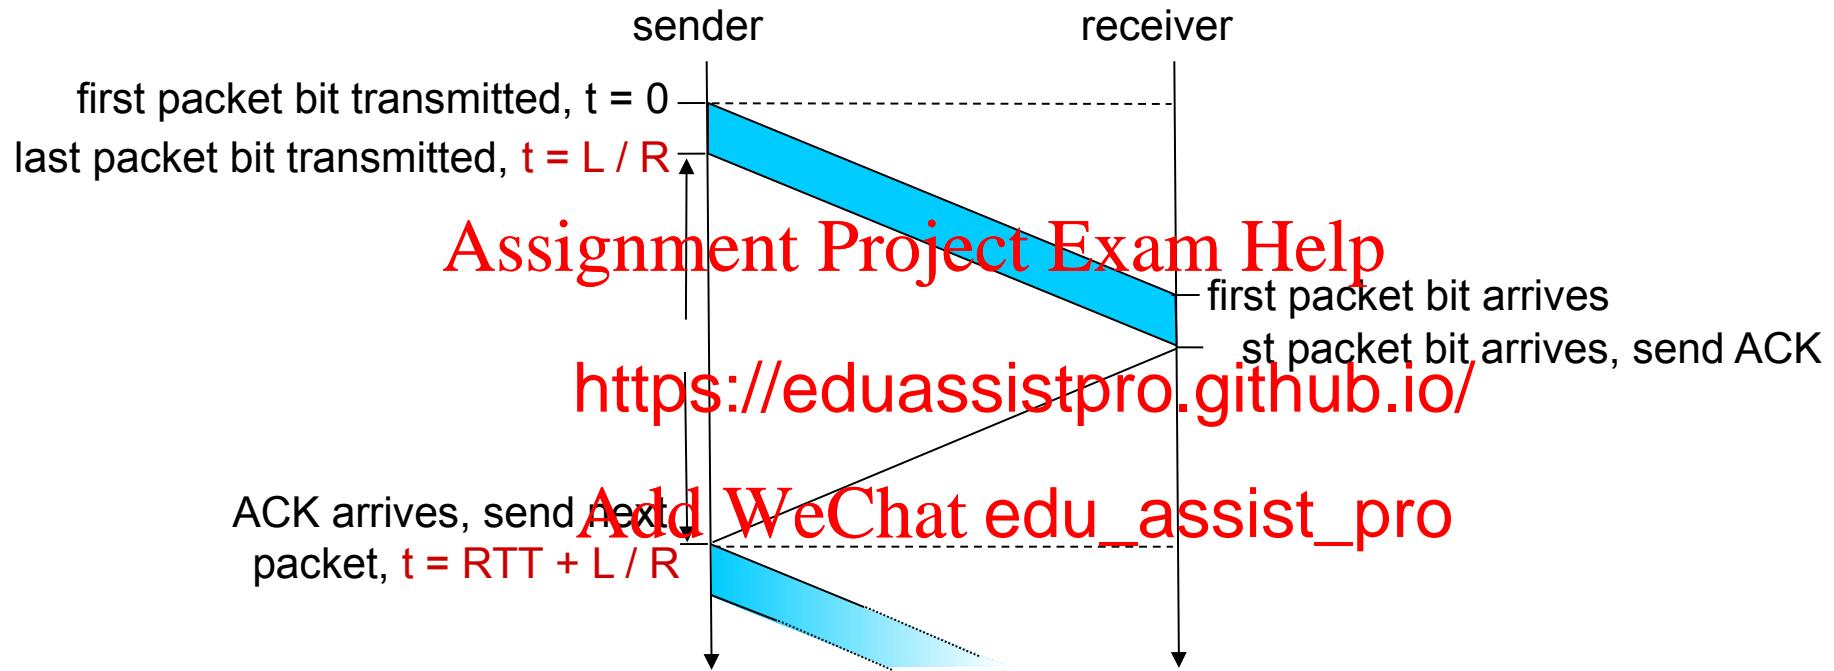
<https://eduassistpro.github.io/>

- $U_{\text{sender}}$ : utilization – fraction of time busy sending

Add WeChat [edu\\_assist\\_pro](https://edu_assist_pro)

$$U_{\text{sender}} = \frac{L/R}{RTT + L/R} = \frac{30.008}{30.008} = 0.00027$$

- if RTT=30 msec, 1KB pkt every 30 msec: 33kB/sec thruput over 1 Gbps link
- ❖ network protocol limits use of physical resources!



$$U_{\text{sender}} = \frac{L/R}{RTT + L/R} = \frac{.008}{30.008} = 0.00027$$

**pipelining:** sender allows multiple, “in-flight”, yet-to-be-acknowledged pkts

- range of sequence numbers must be increased
- buffering at send

**Assignment Project Exam Help**

<https://eduassistpro.github.io/>

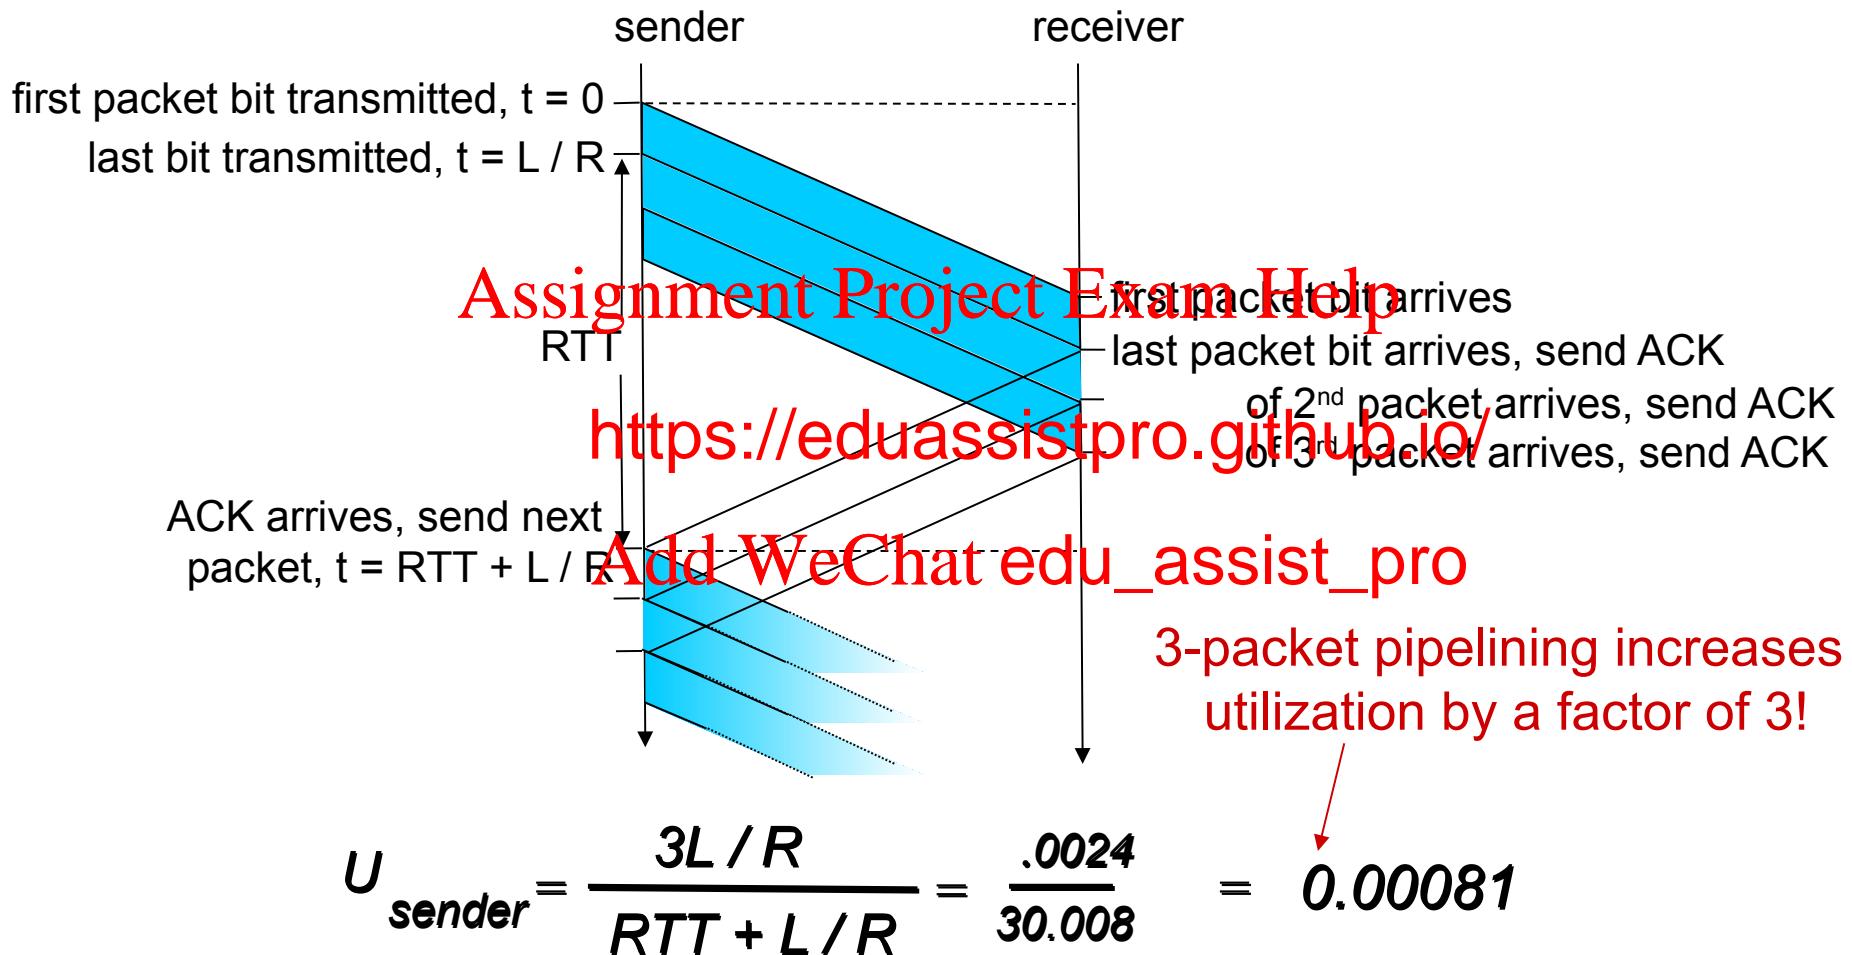


Add WeChat **edu\_assist\_pro**



- › two generic forms of pipelined protocols: **go-Back-N**, **selective repeat**

# Pipelining: increased utilization



### Go-back-N:

- › sender can have up to N unacked packets in pipeline
- › receiver only sees *cumulative ack*
  - does not ack packet if there is a gap
- › sender has timer for oldest unacked packet
  - when timer expires, retransmit *all* unacked packets

### Selective Repeat:

- › sender can have up to N unacked packets in pipeline
- › receiver sends *individual ack* packet
- › Add WeChat `edu_assist_pro` contains timer for unacked packet
  - when timer expires, retransmit only that unacked packet

- › “window” of up to N, consecutive unacked pkts allowed

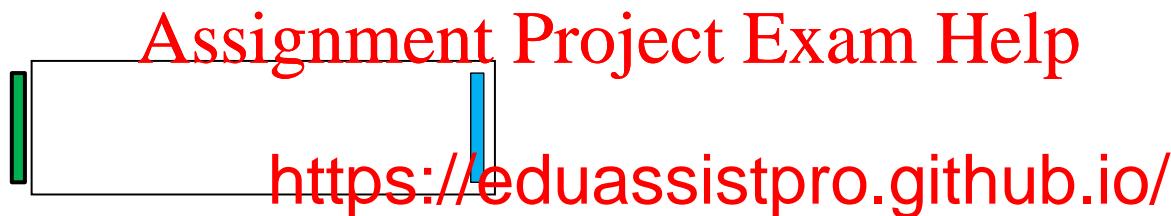
Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

- ❖ ACK(n):ACKs all pkts up to, including seq # n - “*cumulative ACK*”
  - may receive duplicate ACKs (see receiver)
- ❖ timer for oldest in-flight pkt
- ❖ *timeout(n)*: retransmit packet n and all higher seq # pkts in window

- › “window” of up to N, consecutive unacked pkts allowed



Add WeChat `edu_assist_pro`

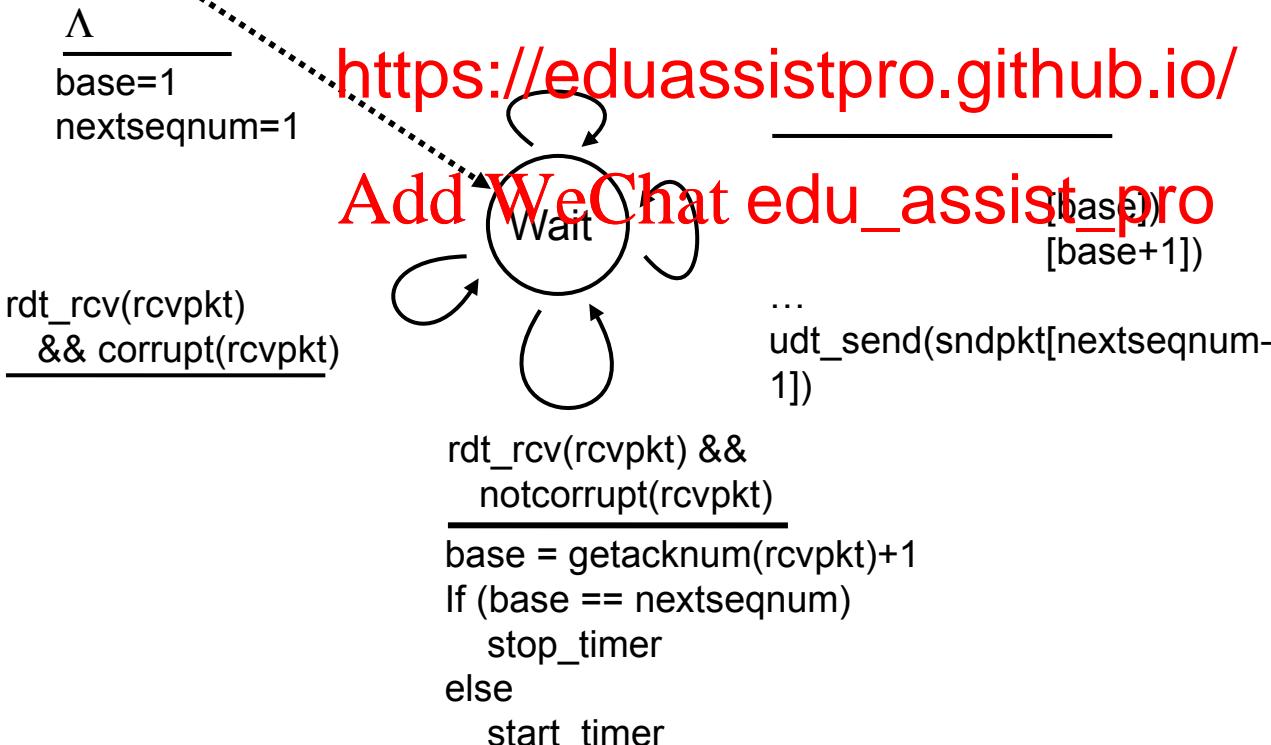
- ❖ ACK(n):ACKs all pkts up to, including seq # n - “*cumulative ACK*”
  - may receive duplicate ACKs (see receiver)
- ❖ timer for oldest in-flight pkt
- ❖ *timeout(n)*: retransmit packet n and all higher seq # pkts in window

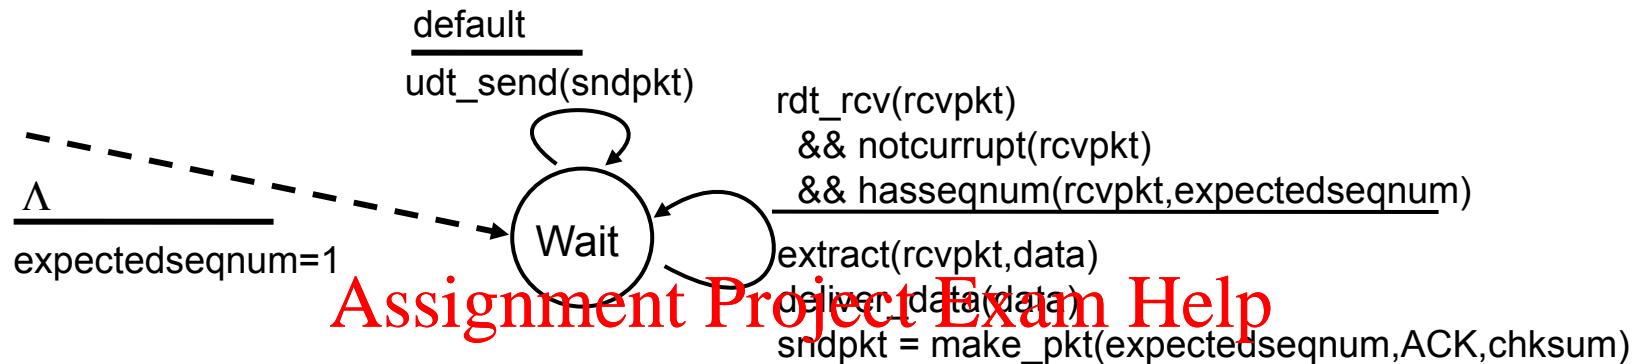
```

rdt_send(data)
if (nextseqnum < base+N) {
    sndpkt[nextseqnum] = make_pkt(nextseqnum,data,chksum)
    udt_send(sndpkt[nextseqnum])
    if (base == nextseqnum)
        start_timer
    nextseqnum++
}

```

Assignment Project Exam Help





Assignment Project Exam Help

<https://eduassistpro.github.io/>

ACK-only: always send ACK  
highest *in-order* seq #

try-received pkt with

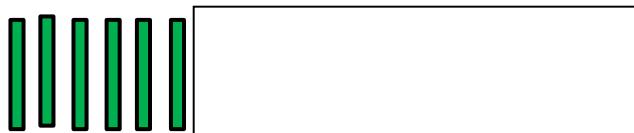
- may generate duplicate ACKs
  - need only remember **expectedseqnum**
- › out-of-order pkt:
- discard (don't buffer): *no receiver buffering!*
  - re-ACK pkt with highest in-order seq #





- › receiver *individually* acknowledges all correctly received pkts
  - buffers pkts as needed, for eventual in-order delivery to upper layer
- › sender only re <https://eduassistpro.github.io/> received
  - sender timer for each unACKed
- › sender window
- › receiver window

## Selective repeat: sender, receiver windows



Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro



## sender

**data from above:**

- › if next available seq # in window, send pkt

**timeout(n):**

- › resend pkt n, res

**ACK(n) in [sendbase,sendbase+N-1]:**

- › mark pkt n as received

- › if n is smallest unACKed pkt, advance window base to next unACKed seq #

## receiver

**pkt n in [rcvbase, rcvbase+N-1]**

- ❖ send ACK(n)
- ❖ out-of-order: buffer
  - ❖ der: deliver (also en buffered, in-order window to t-yet received pkt [rcvbase-N,rcvbase-1])

- ❖ ACK(n)

**otherwise:**

- ❖ ignore

<https://eduassistpro.github.io/>

Add WeChat **edu\_assist\_pro**



Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

